

FUSE

Filesystem in User space

Chian yu Tseng
2012-6-11

Outline

- Introduction
- The FUSE structure
- 如何運作
- Struct fuse_operations
- Example

Introduction(1/2)

- FUSE is a **loadable kernel module** for Unix-like computer operating systems that lets non-privileged users create their own file systems without editing kernel code.
- This is achieved by running file system code in user space while the FUSE module provides only a "**bridge**" to the actual kernel interfaces.
- FUSE is particularly useful for writing **virtual file systems**. Unlike traditional file systems that essentially save data to and retrieve data from disk, virtual filesystems do not actually store data themselves. They act as a view or translation of an existing file system or storage device.

Introduction(2/2)

- The FUSE system was originally part of A Virtual Filesystem ([AVFS](#)), but has since split off into its own project on [SourceForge.net](#).
- FUSE is available for Linux, FreeBSD, NetBSD, OpenSolaris, and Mac OS X. It was officially merged into the mainstream Linux kernel tree in kernel version 2.6.14.

Examples(1/2)

- **ExpanDrive**: A commercial filesystem implementing SFTP/FTP/FTPS using FUSE.
- **GlusterFS**: Clustered Distributed Filesystem having capability to scale up to several petabytes.
- **SSHFS**: Provides access to a remote filesystem through SSH.
- **GmailFS**: Filesystem which stores data as mail in Gmail
- **EncFS**: Encrypted virtual filesystem

Examples(2/2)

- **NTFS-3G**和**Captive NTFS**: allowing access to NTFS filesystem.
- **WikipediaFS** : View and edit Wikipedia articles as if they were real files.
- Sun Microsystems's **Lustre** cluster filesystem
- Sun Microsystems's **ZFS**
- **HDFS**: FUSE bindings exist for the open source Hadoop distributed filesystem.

FUSE Installation

- <http://fuse.sourceforge.net/>
- ./configure
- make
- make install

FUSE source code

- **./doc**: contains FUSE-related documentation. Ex: **how-fuse-works**
- **./include**: contains the FUSE API headers, which you need to create a file system. The only one you need now is **fuse.h**.
- **./lib**: holds the source code to create the FUSE libraries that you will be linking with your binaries to create a file system.
- **./util**: has the source code for the FUSE utility library.
- **./example**: contains samples for your reference.

FUSE structure

- FUSE kernel module (fuse.ko)
 - inode.c, dev.c, control.c, dir.c, file.c
- LibFUSE module (libfuse.*)
 - helper.c, fuse_kern_chan.c, fuse_mt.c, fuse.c, fuse_lowlevel.c, fuse_loop.c, fuse_loop_mt.c, fuse_session.c
- Mount utility(fusermount)
 - fusermount, mount.fuse.c, mount_util.c, mount.c, mount_bsd.c,

FUSE Library

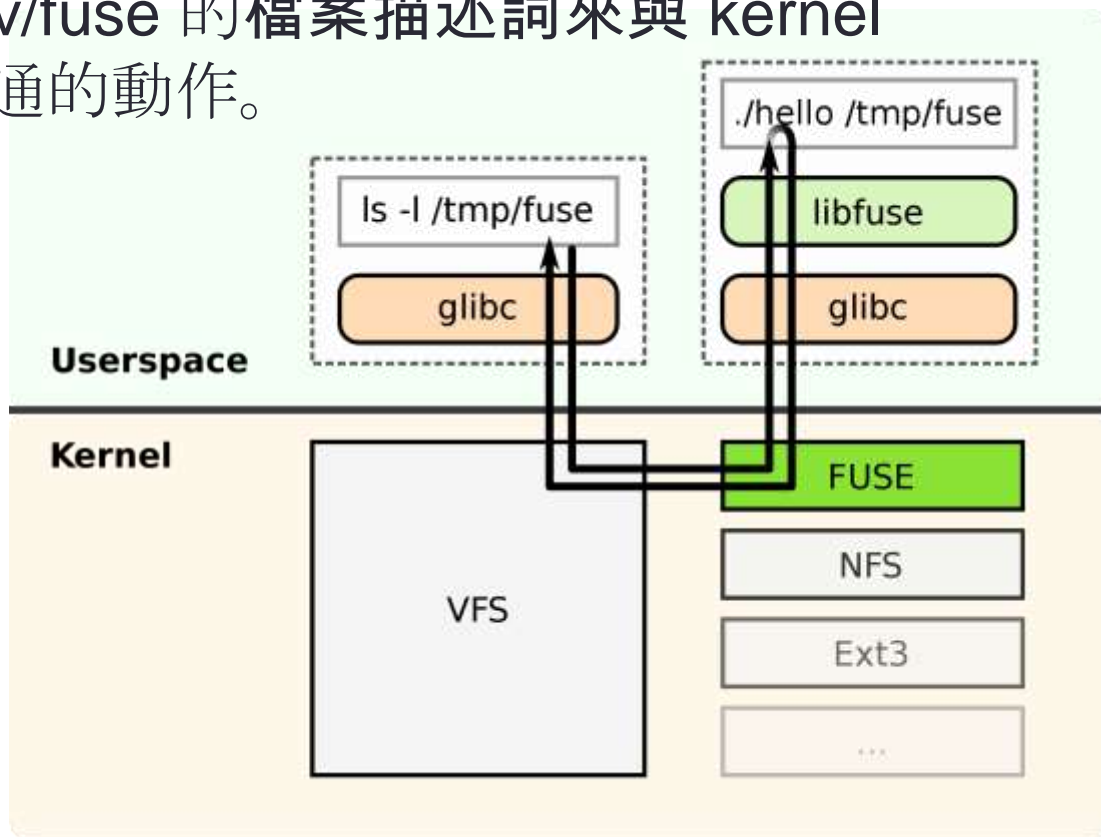
- `include/fuse.h` → the library interface of FUSE (HighLevel)
- `include/fuse_common.h` → common
- `include/fuse_lowlevel.h` → Lowlevel API
- `include/fuse_opt.h` → option parsing interface of FUSE

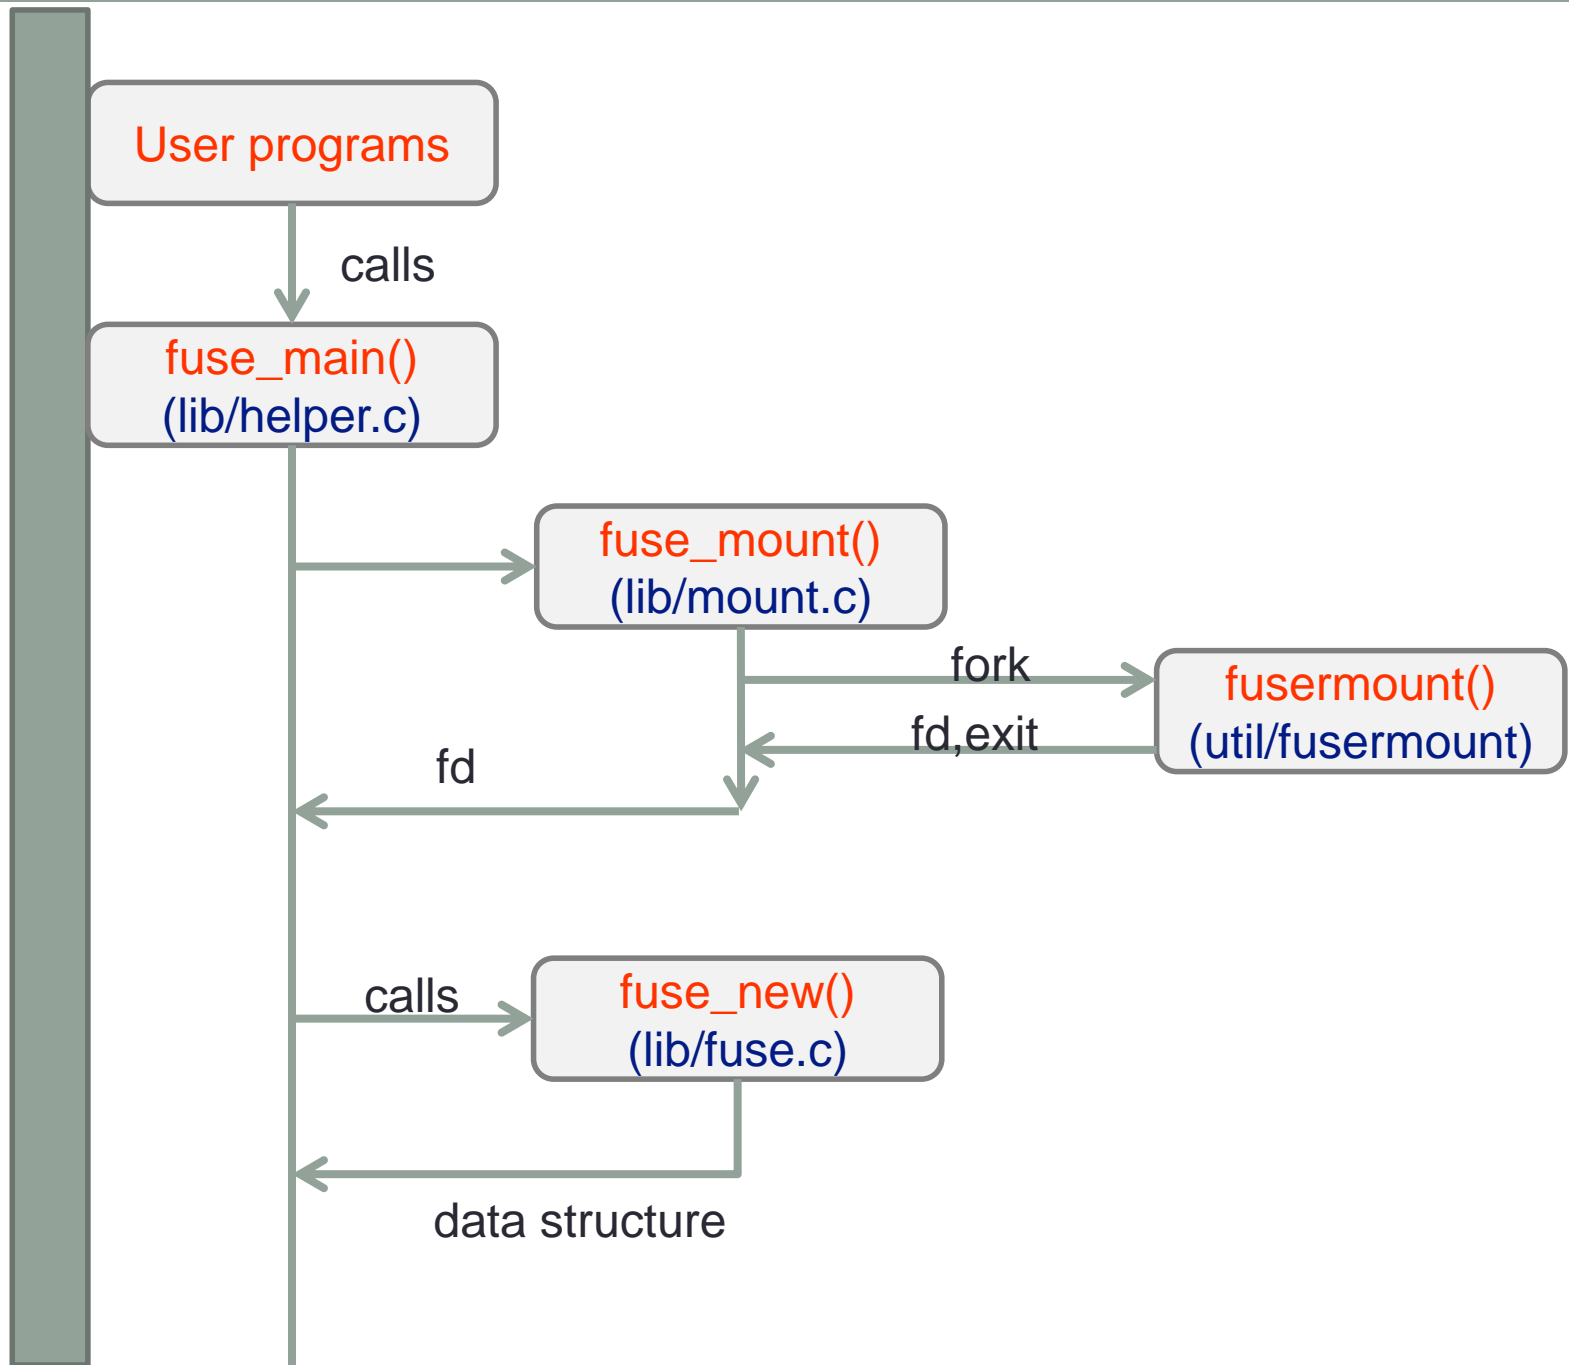
如何運作

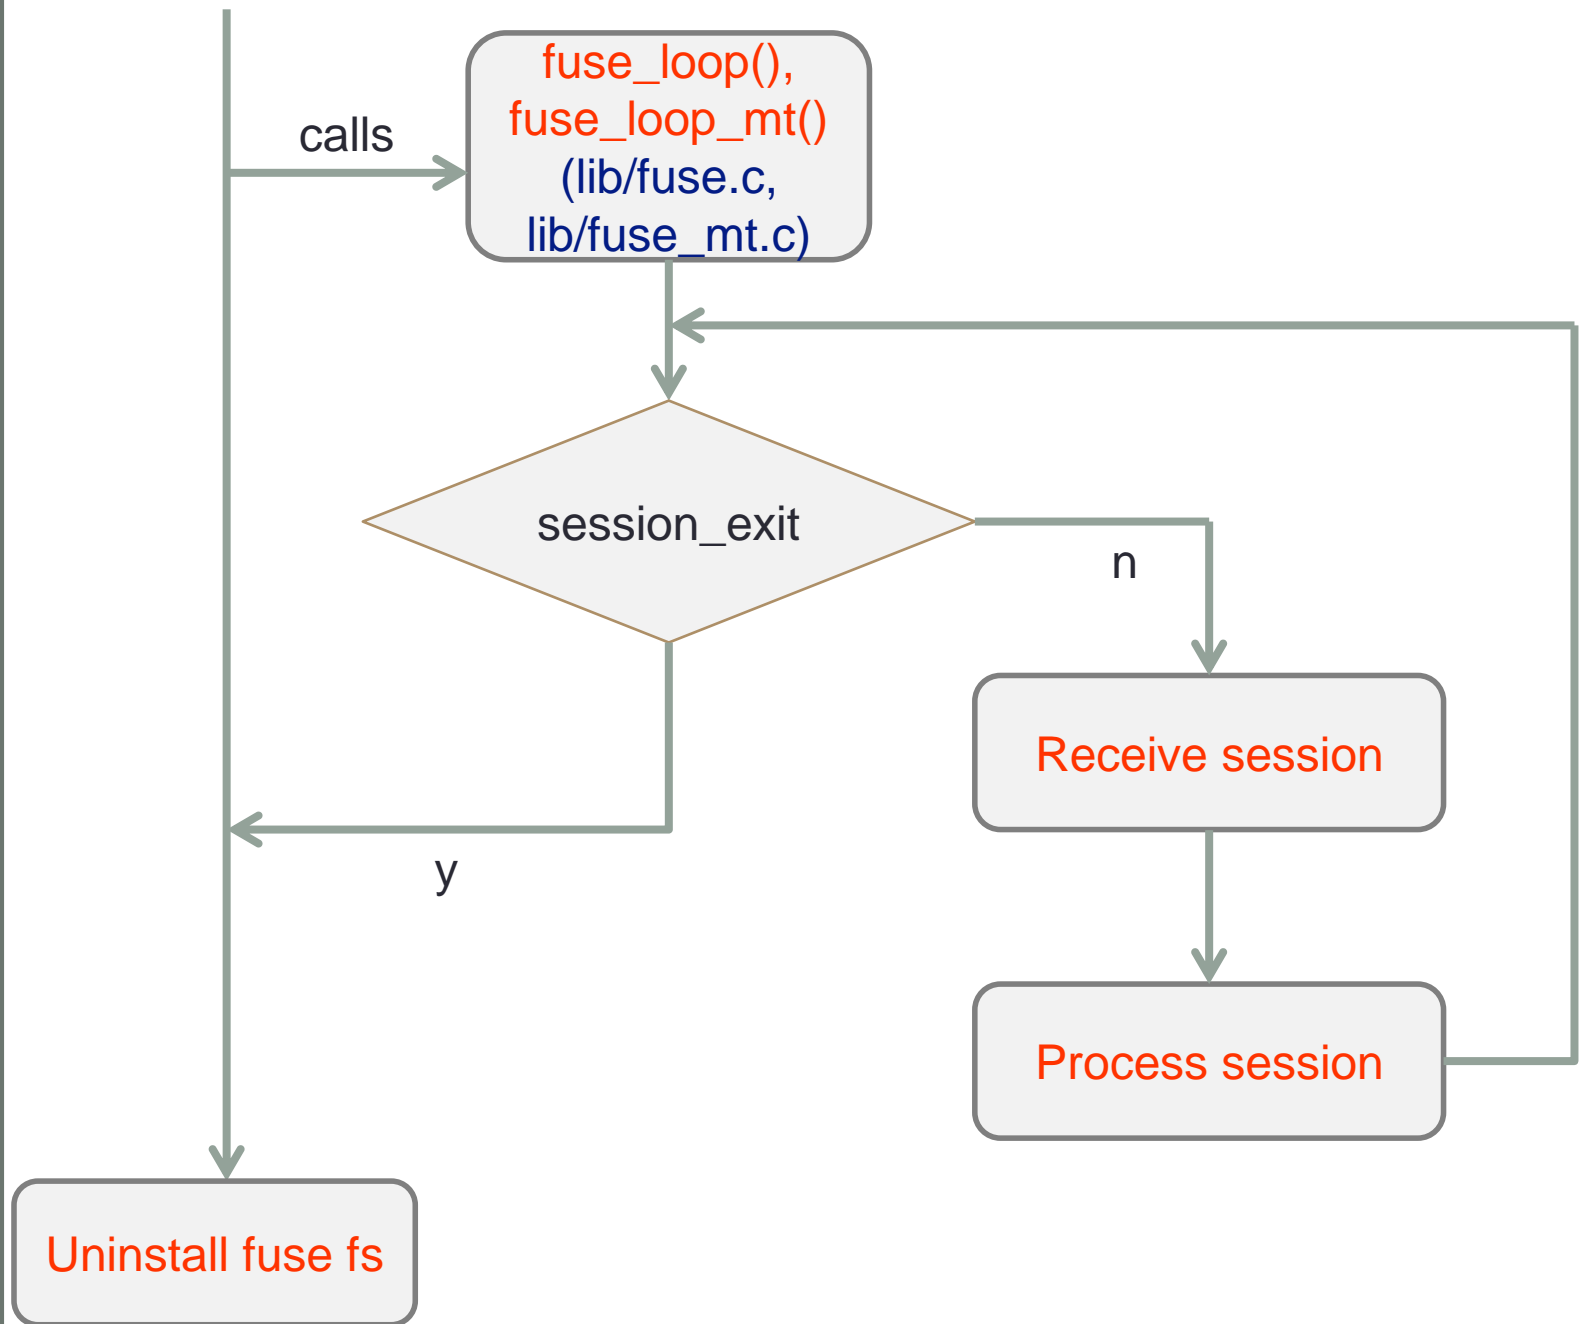
- 在 FUSE daemon 啟動的時候，會先進行掛載的動作，將 /dev/fuse 掛載到指定的目錄底下，並回傳 /dev/fuse 的檔案描述詞(file descriptor)，而 FUSE daemon 在預設上會使用 multi-thread 的方式，透過 /dev/fuse 的檔案描述詞來接收 requests，再根據 requests 的類別來進行處理，最後透過 replies，將結果傳回去。

如何運作

- `ls` : FUSE daemon 會接收到 `OPENDIR`、`REaddir` 等 requests, 並採用 userspace library(`libfuse.*`)的函式, 讀取 file 目錄的資訊, 並將此資訊傳回去, 其中 FUSE daemon 就是透過 `/dev/fuse` 的檔案描述詞來與 kernel module(`fuse.ko`)作溝通的動作。







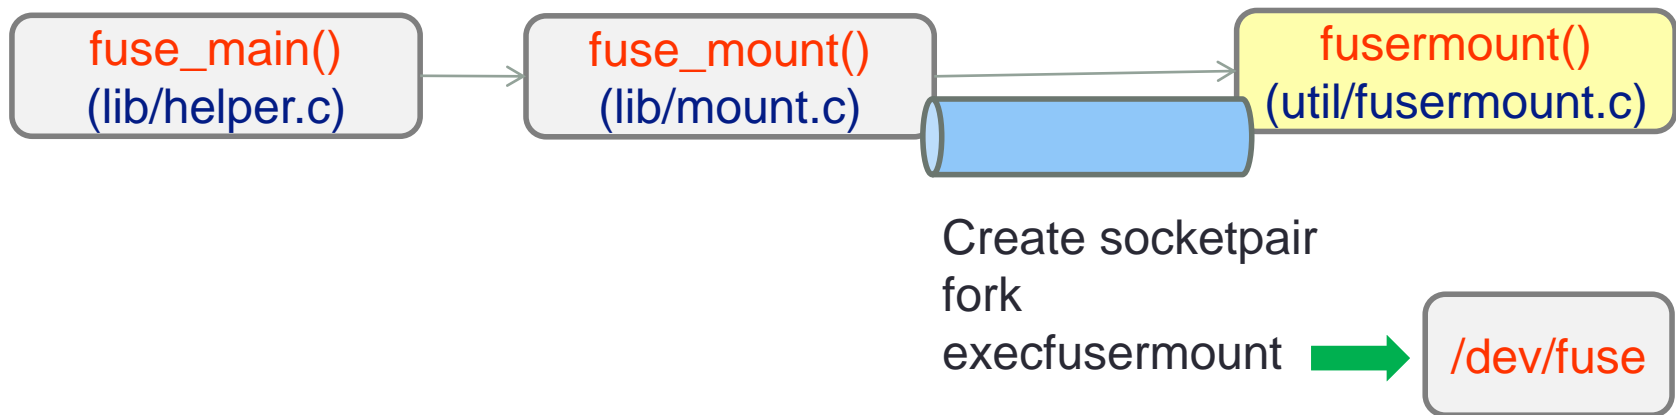
The fuse library(1/5)

- When your user mode program calls `fuse_main()` (`lib/helper.c`), `fuse_main()` parses the arguments passed to your user mode program, then calls `fuse_mount()` (`lib/mount.c`).



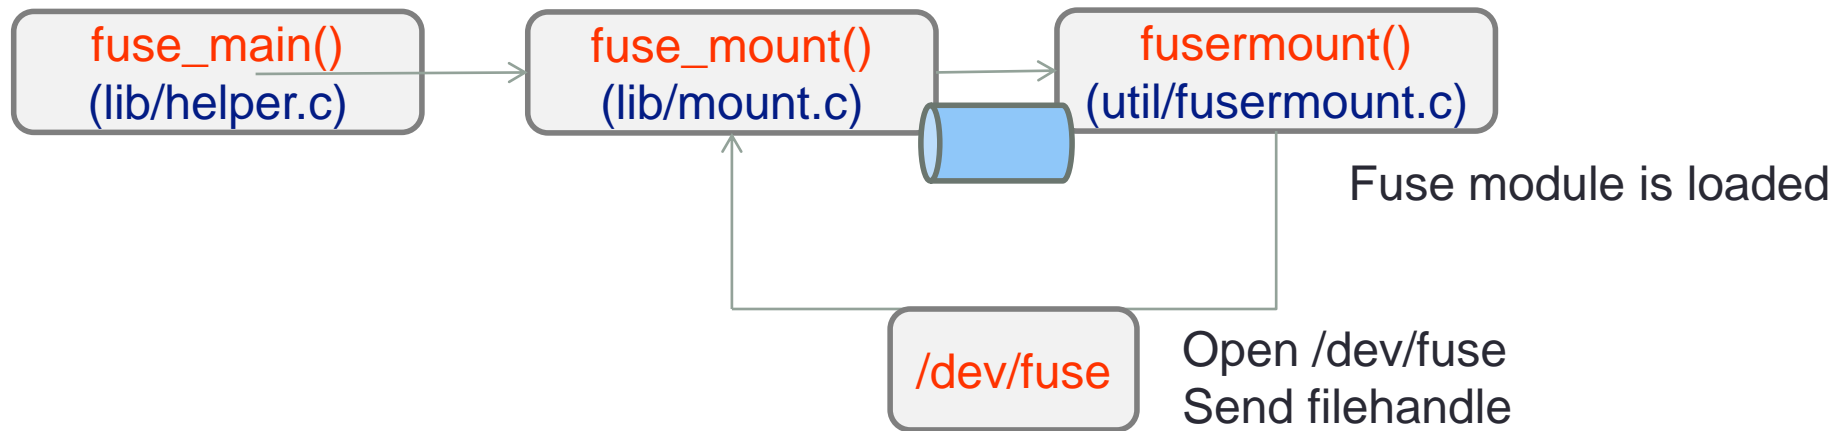
The fuse library(2/5)

- `fuse_mount()` creates a UNIX domain socket pair, then forks and `execsfusermount` ([util/fusermount.c](#)) passing it one end of the socket in the `FUSE_COMMFD_ENV` environment variable.



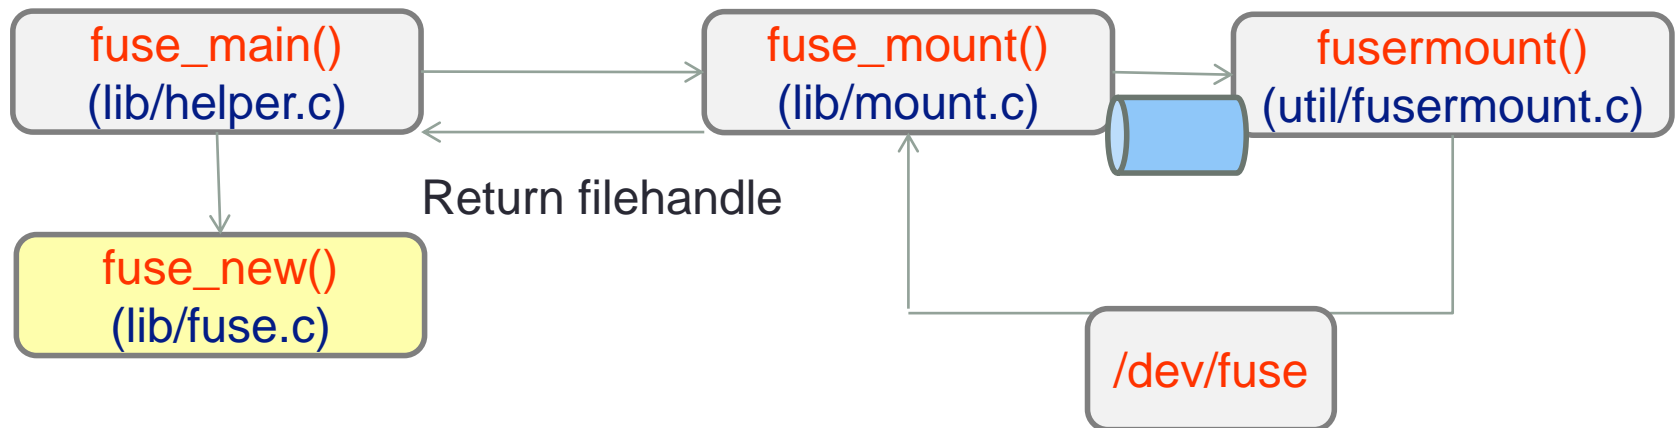
The fuse library(3/5)

- `fusermount` (`util/fusermount.c`) makes sure that the fuse module is loaded. `fusermount` then open `/dev/fuse` and send the file handle over a UNIX domain socket back to `fuse_mount()`.



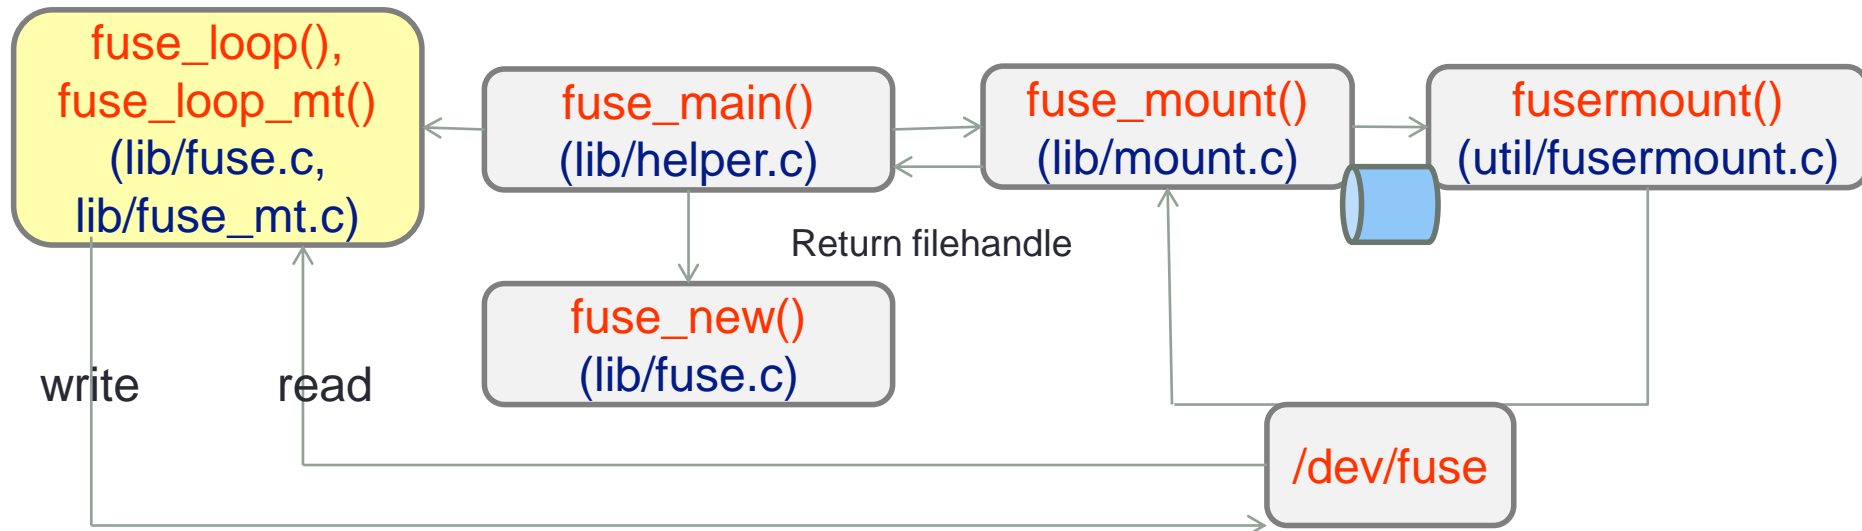
The fuse library(4/5)

- `fuse_mount()` returns the file handle for `/dev/fuse` to `fuse_main()`.
- `fuse_main()` calls `fuse_new()` (`lib/fuse.c`) which allocates the struct fuse data structure that stores and maintains a cached image of the filesystem data.



The fuse library (5/5)

- Lastly, `fuse_main()` calls either `fuse_loop()` (`lib/fuse.c`) or `fuse_loop_mt()` (`lib/fuse_mt.c`) which both start to read the file system system calls from the `/dev/fuse`, call the user mode functions stored in **`struct fuse_operations`** data structure before calling `fuse_main()`.
- The results of those calls are then written back to the `/dev/fuse` file where they can be forwarded back to the system calls.



Struct fuse_operations (1/9)

- `int (*getattr) (const char *, struct stat *);`
 - Get file attributes.
- `int (*readlink) (const char *, char *, size_t);`
 - Read the target of a symbolic link
- `int (*mknod) (const char *, mode_t, dev_t);`
 - Create a file node.
- `int (*mkdir) (const char *, mode_t);`
 - Create a directory. Note that the mode argument may not have the type specification bits set, i.e. `S_ISDIR(mode)` can be false. To obtain the correct directory type bits use `mode | S_IFDIR`

Struct fuse_operations (2/9)

- `int (*unlink) (const char *);`
 - Remove a file
- `int (*rmdir) (const char *);`
 - Remove a directory
- `int (*symlink) (const char *, const char *);`
 - Create a symbolic link
- `int (*rename) (const char *, const char *);`
 - Rename a file
- `int (*link) (const char *, const char *);`
 - Create a hard link to a file

Struct fuse_operations (3/9)

- `int (*chmod) (const char *, mode_t);`
 - Change the permission bits of a file
- `int (*chown) (const char *, uid_t, gid_t);`
 - Change the owner and group of a file
- `int (*truncate) (const char *, off_t);`
 - Change the size of a file
- `int (*open) (const char *, struct fuse_file_info *);`
 - File open operation.

Struct fuse_operations (4/9)

- `int (*read) (const char *, char *, size_t, off_t, struct fuse_file_info *);`
 - Read data from an open file.
- `int (*write) (const char *, const char *, size_t, off_t, struct fuse_file_info *);`
 - Write data to an open file
- `int (*statfs) (const char *, struct statvfs *);`
 - Get file system statistics
- `int (*flush) (const char *, struct fuse_file_info *);`
 - Possibly flush cached data

Struct fuse_operations (5/9)

- `int (*release) (const char *, struct fuse_file_info *);`
 - Release an open file. Release is called when there are no more references to an open file: all file descriptors are closed and all memory mappings are unmapped.
- `int (*fsync) (const char *, int, struct fuse_file_info *);`
 - Synchronize file contents
- `int (*setxattr) (const char *, const char *, const char *, size_t, int);`
 - Set extended attributes
- `int (*getxattr) (const char *, const char *, char *, size_t);`
 - Get extended attributes

Struct fuse_operations (6/9)

- `int (*listxattr) (const char *, char *, size_t);`
 - List extended attributes
- `int (*removexattr) (const char *, const char *);`
 - Remove extended attributes
- `int (*opendir) (const char *, struct fuse_file_info *);`
 - Open directory. Unless the 'default_permissions' mount option is given, this method should check if opendir is permitted for this directory. Optionally opendir may also return an arbitrary **filehandle** in the [fuse_file_info](#) structure, which will be passed to readdir, closedir and fsyncdir.

Struct fuse_operations (7/9)

- `int (*readdir) (const char *, void *, fuse_fill_dir_t, off_t, struct fuse_file_info *);`
 - Read directory
- `int (*releasedir) (const char *, struct fuse_file_info *);`
 - Release directory
- `int (*fsyncdir) (const char *, int, struct fuse_file_info *);`
 - Synchronize directory contents
- `void *(*init) (struct fuse_conn_info *conn);`
 - Initialize file system.

Struct fuse_operations (8/9)

- `void (*destroy) (void *);`
 - Clean up filesystem
- `int (*access) (const char *, int);`
 - Check file access permissions
- `int (*create) (const char *, mode_t, struct fuse_file_info *);`
 - Create and open a file. If the file does not exist, first create it with the specified mode, and then open it.
- `int (*ftruncate) (const char *, off_t, struct fuse_file_info *);`
 - Change the size of an open file
- `int (*fgetattr) (const char *, struct stat *, struct fuse_file_info *);`
 - Get attributes from an open file

Struct fuse_operations(9/9)

- `int (*lock) (const char *, struct fuse_file_info *, int cmd, struct flock *);`
 - Perform POSIX file locking operation
- `int (*utimens) (const char *, const struct timespec tv[2]);`
 - Change the access and modification times of a file with nanosecond resolution
- `int (*bmap) (const char *, size_t blocksize, uint64_t *idx);`
 - Map block index within file to block index within device

Example1: Hello.c

```
11  #define FUSE_USE_VERSION 26
12
13  #include <fuse.h>
14  #include <stdio.h>
15  #include <string.h>
16  #include <errno.h>
17  #include <fcntl.h>
18
19  static const char *hello_str = "Hello World!\n";
20  static const char *hello_path = "/hello";
21
```

```
86  static struct fuse_operations hello_oper = {
87      .getattr    = hello_getattr,
88      .readdir    = hello_readdir,
89      .open       = hello_open,
90      .read       = hello_read,
91  };
92
93  int main(int argc, char *argv[])
94  {
95      return fuse_main(argc, argv, &hello_oper, NULL);
96  }
97
```

hello-getattr()

```
22 static int hello_getattr(const char *path, struct stat *stbuf)
23 {
24     int res = 0;
25
26     memset(stbuf, 0, sizeof(struct stat));
27     if (strcmp(path, "/") == 0) {
28         stbuf->st_mode = S_IFDIR | 0755;
29         stbuf->st_nlink = 2;
30     } else if (strcmp(path, hello_path) == 0) {
31         stbuf->st_mode = S_IFREG | 0444;
32         stbuf->st_nlink = 1;
33         stbuf->st_size = strlen(hello_str);
34     } else
35         res = -ENOENT; A component of the path path does not exist
36
37     return res;
38 }
```

```
danny@danny-desktop:/tmp$ ls -l
total 8
drwxr-xr-x 2 root root 0 1970-01-01 08:00 fuse
drwx----- 2 gdm gdm 4096 2012-05-22 11:41 orbit-gdm
drwx----- 2 gdm gdm 4096 2012-05-22 11:11 pulse-PKdhtXMmr18n
danny@danny-desktop:/tmp$ ls -l fuse/
total 0
-r--r--r-- 1 root root 13 1970-01-01 08:00 hello
```

hello_readdir()

```
40 static int hello_readdir(const char *path, void *buf, fuse_fill_dir_t filler,  
41                          off_t offset, struct fuse_file_info *fi)  
42 {  
43     (void) offset;  
44     (void) fi;  
45  
46     if (strcmp(path, "/") != 0)  
47         return -ENOENT;  
48  
49     filler(buf, ".", NULL, 0);  
50     filler(buf, "..", NULL, 0);  
51     filler(buf, hello_path + 1, NULL, 0);  
52  
53     return 0;  
54 }
```

```
typedef int(* fuse_fill_dir_t)(void *buf, const char *name, const struct stat *stbuf, off_t off)
```

Function to add an entry in a readdir() operation

Parameters:

buf the buffer passed to the readdir() operation
name the file name of the directory entry
stat file attributes, can be NULL
off offset of the next entry or zero

Returns:

1 if buffer is full, zero otherwise

hello_open()

- This function checks whatever user is ***permitted*** to open the /hello file with flags given in the [fuse_file_info](#) structure.

```
56 static int hello_open(const char *path, struct fuse_file_info *fi)
57 {
58     if (strcmp(path, hello_path) != 0)
59         return -ENOENT;
60
61     if ((fi->flags & 3) != O_RDONLY)
62         return -EACCES;
63
64     return 0;
65 }
```


hello_read()

```
67 static int hello_read(const char *path, char *buf, size_t size, off_t offset,
68                      struct fuse_file_info *fi)
69 {
70     size_t len;
71     (void) fi;
72     if(strcmp(path, hello_path) != 0)
73         return -ENOENT;
74
75     len = strlen(hello_str);
76     if (offset < len) {
77         if (offset + size > len)
78             size = len - offset;
79         memcpy(buf, hello_str + offset, size);
80     } else
81         size = 0;
82
83     return size;
84 }
```

Example1: Hello.c 執行

- `./hello /tmp/fuse -d`

```
danny@danny-desktop: ~/fuse-2.9.0/example [106x35]
連線(C) 編輯(E) 檢視(V) 視窗(W) 選項(O) 說明(H)
danny@danny-desktop:~/fuse-2.9.0/example$ ./hello /tmp/fuse -d
FUSE library version: 2.9.0
nullpath_ok: 0
nopath: 0
utime_omit_ok: 0
unique: 1, opcode: INIT (26), nodeid: 0, insize: 56, pid: 0
INIT: 7.17
flags=0x00000047b
max_readahead=0x00020000
INIT: 7.18
flags=0x000000011
max_readahead=0x00020000
max_write=0x00020000
max_background=0
congestion_threshold=0
unique: 1, success, outsize: 40
```

```
danny@danny-desktop:/tmp$ ls -l
total 8
drwxr-xr-x 2 root root 0 1970-01-01 08:00 fuse
drwx----- 2 gdm gdm 4096 2012-05-22 11:41 orbit-gdm
drwx----- 2 gdm gdm 4096 2012-05-22 11:11 pulse-PKdhtXMmr18n
danny@danny-desktop:/tmp$ ls -l fuse/
total 0
-r--r--r-- 1 root root 13 1970-01-01 08:00 hello
danny@danny-desktop:/tmp$ cat fuse/hello
Hello World!
danny@danny-desktop:/tmp$
```

Example2: fusexmp_fh.c

```
500 static struct fuse_operations xmp_oper = {
501     .getattr      = xmp_getattr,
502     .fgetattr     = xmp_fgetattr,
503     .access       = xmp_access,
504     .readlink     = xmp_readlink,
505     .opendir      = xmp_opendir,
506     .readdir      = xmp_readdir,
507     .releasedir   = xmp_releasedir,
508     .mknod        = xmp_mknod,
509     .mkdir        = xmp_mkdir,
510     .symlink      = xmp_symlink,
511     .unlink       = xmp_unlink,
512     .rmdir        = xmp_rmdir,
513     .rename       = xmp_rename,
514     .link         = xmp_link,
515     .chmod        = xmp_chmod,
516     .chown        = xmp_chown,
517     .truncate     = xmp_truncate,
518     .ftruncate    = xmp_ftruncate,
519 #ifdef HAVE_UTIMENSAT
520     .utimens      = xmp_utimens,
521 #endif
522     .create       = xmp_create,
523     .open         = xmp_open,
524     .read         = xmp_read,
525     .read_buf     = xmp_read_buf,
526     .write        = xmp_write,
```

```
526     .write        = xmp_write,
527     .write_buf    = xmp_write_buf,
528     .statfs       = xmp_statfs,
529     .flush        = xmp_flush,
530     .release      = xmp_release,
531     .fsync        = xmp_fsync,
532 #ifdef HAVE_SETXATTR
533     .setxattr     = xmp_setxattr,
534     .getxattr     = xmp_getxattr,
535     .listxattr    = xmp_listxattr,
536     .removexattr  = xmp_removexattr,
537 #endif
538     .lock         = xmp_lock,
539     .flock        = xmp_flock,
540
541     .flag_nullpath_ok = 1,
542 #if HAVE_UTIMENSAT
543     .flag_utime_omit_ok = 1,
544 #endif
545 };
546
547 int main(int argc, char *argv[])
548 {
549     umask(0);
550     return fuse_main(argc, argv, &xmp_oper, NULL);
551 }
552
```

xmp_getattr(), xmp_fgetattr()

```
36 static int xmp_getattr(const char *path, struct stat *stbuf)
37 {
38     int res;
39
40     res = lstat(path, stbuf);
41     if (res == -1)
42         return -errno;
43
44     return 0;
45 }
46
47 static int xmp_fgetattr(const char *path, struct stat *stbuf,
48                        struct fuse_file_info *fi)
49 {
50     int res;
51
52     (void) path;
53
54     res = fstat(fi->fh, stbuf);
55     if (res == -1)
56         return -errno;
57
58     return 0;
59 }
```

xmp_access(), xmp_readlink()

```
61 static int xmp_access(const char *path, int mask)
62 {
63     int res;
64
65     res = access(path, mask);
66     if (res == -1)
67         return -errno;
68
69     return 0;
70 }
71
72 static int xmp_readlink(const char *path, char *buf, size_t size)
73 {
74     int res;
75
76     res = readlink(path, buf, size - 1);
77     if (res == -1)
78         return -errno;
79
80     buf[res] = '\0';
81     return 0;
82 }
```

Struct xmp_dirp, xmp_opendir()

```
84 struct xmp_dirp {
85     DIR *dp;
86     struct dirent *entry;
87     off_t offset;
88 };
89
90 static int xmp_opendir(const char *path, struct fuse_file_info *fi)
91 {
92     int res;
93     struct xmp_dirp *d = malloc(sizeof(struct xmp_dirp));
94     if (d == NULL)
95         return -ENOMEM;
96
97     d->dp = opendir(path);
98     if (d->dp == NULL) {
99         res = -errno;
100         free(d);
101         return res;
102     }
103     d->offset = 0;
104     d->entry = NULL;
105
106     fi->fh = (unsigned long) d;
107     return 0;
108 }
```

xmp_readdir() (1/2)

```
114
115 static int xmp_readdir(const char *path, void *buf, fuse_fill_dir_t filler,
116                       off_t offset, struct fuse_file_info *fi)
117 {
118     struct xmp_dirp *d = get_dirp(fi);
119
120     (void) path;
121     if (offset != d->offset) {
122         seekdir(d->dp, offset);
123         d->entry = NULL;
124         d->offset = offset;
125     }
126     while (1) {
127         struct stat st;
128         off_t nextoff;
129
130         if (!d->entry) {
131             d->entry = readdir(d->dp);
132             if (!d->entry)
133                 break;
134         }
135
136         memset(&st, 0, sizeof(st));
137         st.st_ino = d->entry->d_ino;
138         st.st_mode = d->entry->d_type << 12;
139         nextoff = telldir(d->dp);
```

xmp_readdir() (2/2)

```
140         if (filler(buf, d->entry->d_name, &st, nextoff))
141             break;
142
143         d->entry = NULL;
144         d->offset = nextoff;
145     }
146
147     return 0;
148 }
```


xmp_releasedir(), xmp_mknod()

```
150 static int xmp_releasedir(const char *path, struct fuse_file_info *fi)
151 {
152     struct xmp_dirp *d = get_dirp(fi);
153     (void) path;
154     closedir(d->dp);
155     free(d);
156     return 0;
157 }
158
159 static int xmp_mknod(const char *path, mode_t mode, dev_t rdev)
160 {
161     int res;
162
163     if (S_ISFIFO(mode))
164         res = mkfifo(path, mode);
165     else
166         res = mknod(path, mode, rdev);
167     if (res == -1)
168         return -errno;
169
170     return 0;
171 }
```

xmp_mkdir(), xmp_unlink()

```
173 static int xmp_mkdir(const char *path, mode_t mode)
174 {
175     int res;
176
177     res = mkdir(path, mode);
178     if (res == -1)
179         return -errno;
180
181     return 0;
182 }
183
184 static int xmp_unlink(const char *path)
185 {
186     int res;
187
188     res = unlink(path);
189     if (res == -1)
190         return -errno;
191
192     return 0;
193 }
```

xmp_rmdir(), xmp_symlink()

```
195 static int xmp_rmdir(const char *path)
196 {
197     int res;
198
199     res = rmdir(path);
200     if (res == -1)
201         return -errno;
202
203     return 0;
204 }
205
206 static int xmp_symlink(const char *from, const char *to)
207 {
208     int res;
209
210     res = symlink(from, to);
211     if (res == -1)
212         return -errno;
213
214     return 0;
215 }
```

xmp_rename(), xmp_link()

```
217 static int xmp_rename(const char *from, const char *to)
218 {
219     int res;
220
221     res = rename(from, to);
222     if (res == -1)
223         return -errno;
224
225     return 0;
226 }
227
228 static int xmp_link(const char *from, const char *to)
229 {
230     int res;
231
232     res = link(from, to);
233     if (res == -1)
234         return -errno;
235
236     return 0;
237 }
```

xmp_chmod(), xmp_chown()

```
239 static int xmp_chmod(const char *path, mode_t mode)
240 {
241     int res;
242
243     res = chmod(path, mode);
244     if (res == -1)
245         return -errno;
246
247     return 0;
248 }
249
250 static int xmp_chown(const char *path, uid_t uid, gid_t gid)
251 {
252     int res;
253
254     res = lchown(path, uid, gid);
255     if (res == -1)
256         return -errno;
257
258     return 0;
259 }
```

xmp_truncate(), xmp_ftruncate()

```
261 static int xmp_truncate(const char *path, off_t size)
262 {
263     int res;
264
265     res = truncate(path, size);
266     if (res == -1)
267         return -errno;
268
269     return 0;
270 }
271
272 static int xmp_ftruncate(const char *path, off_t size,
273                          struct fuse_file_info *fi)
274 {
275     int res;
276
277     (void) path;
278
279     res = ftruncate(fi->fh, size);
280     if (res == -1)
281         return -errno;
282
283     return 0;
284 }
```

xmp_utimens(), xmp_create()

```
286 #ifdef HAVE_UTIMENSAT
287 static int xmp_utimens(const char *path, const struct timespec ts[2])
288 {
289     int res;
290
291     /* don't use utime/utimes since they follow symlinks */
292     res = utimensat(0, path, ts, AT_SYMLINK_NOFOLLOW);
293     if (res == -1)
294         return -errno;
295
296     return 0;
297 }
298 #endif
299
300 static int xmp_create(const char *path, mode_t mode, struct fuse_file_info *fi)
301 {
302     int fd;
303
304     fd = open(path, fi->flags, mode);
305     if (fd == -1)
306         return -errno;
307
308     fi->fh = fd;
309     return 0;
310 }
```

xmp_open(), xmp_read()

```
312 static int xmp_open(const char *path, struct fuse_file_info *fi)
313 {
314     int fd;
315
316     fd = open(path, fi->flags);
317     if (fd == -1)
318         return -errno;
319
320     fi->fh = fd;
321     return 0;
322 }
323
324 static int xmp_read(const char *path, char *buf, size_t size, off_t offset,
325                     struct fuse_file_info *fi)
326 {
327     int res;
328
329     (void) path;
330     res = pread(fi->fh, buf, size, offset);
331     if (res == -1)
332         res = -errno;
333
334     return res;
335 }
```


xmp_read_buf()

```
337 static int xmp_read_buf(const char *path, struct fuse_bufvec **bufp,  
338                          size_t size, off_t offset, struct fuse_file_info *fi)  
339 {  
340     struct fuse_bufvec *src;  
341  
342     (void) path;  
343  
344     src = malloc(sizeof(struct fuse_bufvec));  
345     if (src == NULL)  
346         return -ENOMEM;  
347  
348     *src = FUSE_BUFVEC_INIT(size);  
349  
350     src->buf[0].flags = FUSE_BUF_IS_FD | FUSE_BUF_FD_SEEK;  
351     src->buf[0].fd = fi->fh;  
352     src->buf[0].pos = offset;  
353  
354     *bufp = src;  
355  
356     return 0;  
357 }
```

xmp_write(), xmp_write_buf()

```
359 static int xmp_write(const char *path, const char *buf, size_t size,
360                      off_t offset, struct fuse_file_info *fi)
361 {
362     int res;
363
364     (void) path;
365     res = pwrite(fi->fh, buf, size, offset);
366     if (res == -1)
367         res = -errno;
368
369     return res;
370 }
371
372 static int xmp_write_buf(const char *path, struct fuse_bufvec *buf,
373                          off_t offset, struct fuse_file_info *fi)
374 {
375     struct fuse_bufvec dst = FUSE_BUFVEC_INIT(fuse_buf_size(buf));
376
377     (void) path;
378
379     dst.buf[0].flags = FUSE_BUF_IS_FD | FUSE_BUF_FD_SEEK;
380     dst.buf[0].fd = fi->fh;
381     dst.buf[0].pos = offset;
382
383     return fuse_buf_copy(&dst, buf, FUSE_BUF_SPLICE_NONBLOCK);
384 }
```

xmp_statfs(), xmp_flush()

```
386 static int xmp_statfs(const char *path, struct statvfs *stbuf)
387 {
388     int res;
389
390     res = statvfs(path, stbuf);
391     if (res == -1)
392         return -errno;
393
394     return 0;
395 }
396
397 static int xmp_flush(const char *path, struct fuse_file_info *fi)
398 {
399     int res;
400
401     (void) path;
402     /* This is called from every close on an open file, so call the
403        close on the underlying filesystem. But since flush may be
404        called multiple times for an open file, this must not really
405        close the file. This is important if used on a network
406        filesystem like NFS which flush the data/metadata on close() */
407     res = close(dup(fi->fh));
408     if (res == -1)
409         return -errno;
410
411     return 0;
412 }
```

xmp_release(), xmp_fsync()

```
414 static int xmp_release(const char *path, struct fuse_file_info *fi)
415 {
416     (void) path;
417     close(fi->fh);
418
419     return 0;
420 }
421
422 static int xmp_fsync(const char *path, int isdatasync,
423                     struct fuse_file_info *fi)
424 {
425     int res;
426     (void) path;
427
428     #ifndef HAVE_FDATASYNC
429         (void) isdatasync;
430     #else
431         if (isdatasync)
432             res = fdatasync(fi->fh);
433         else
434             res = fsync(fi->fh);
435         if (res == -1)
436             return -errno;
437
438     return 0;
439 }
440
```

xmp_setattr(), xmp_getattr()

```
443  /* xattr operations are optional and can safely be left unimplemented */
444  static int xmp_setxattr(const char *path, const char *name, const char *value,
445                        size_t size, int flags)
446  {
447      int res = lsetxattr(path, name, value, size, flags);
448      if (res == -1)
449          return -errno;
450      return 0;
451  }
452
453  static int xmp_getxattr(const char *path, const char *name, char *value,
454                        size_t size)
455  {
456      int res = lgetxattr(path, name, value, size);
457      if (res == -1)
458          return -errno;
459      return res;
460  }
```

xmp_listattr(), xmp_removexattr()

```
462 static int xmp_listxattr(const char *path, char *list, size_t size)
463 {
464     int res = llistxattr(path, list, size);
465     if (res == -1)
466         return -errno;
467     return res;
468 }
469
470 static int xmp_removexattr(const char *path, const char *name)
471 {
472     int res = lremovexattr(path, name);
473     if (res == -1)
474         return -errno;
475     return 0;
476 }
477 #endif /* HAVE_SETXATTR */
```

xmp_lock(), xmp_flock()

```
479 static int xmp_lock(const char *path, struct fuse_file_info *fi, int cmd,
480                    struct flock *lock)
481 {
482     (void) path;
483
484     return ulockmgr_op(fi->fh, cmd, lock, &fi->lock_owner,
485                      sizeof(fi->lock_owner));
486 }
487
488 static int xmp_flock(const char *path, struct fuse_file_info *fi, int op)
489 {
490     int res;
491     (void) path;
492
493     res = flock(fi->fh, op);
494     if (res == -1)
495         return -errno;
496
497     return 0;
498 }
```

Example2: fusexmp_fh.c 執行

```
danny@danny-desktop: ~/fuse-2.9.0 [106x35]
連線(C) 編輯(E) 檢視(V) 視窗(W) 選項(O) 說明(H)
danny@danny-desktop:~/fuse-2.9.0/example$ ./fusexmp_fh /tmp/fuse -d
FUSE library version: 2.9.0
nullpath_ok: 1
nopath: 0
utime_omit_ok: 1
unique: 1, opcode: INIT (26), nodeid: 0, insize: 56, pid: 0
INIT: 7.17
flags=0x0000047b
max_readahead=0x00020000
INIT: 7.18
flags=0x00000413
max_readahead=0x00020000
max_write=0x00020000
max_background=0
congestion_threshold=0
unique: 1, success, outsize: 40
█
```

```
danny@danny-desktop: ~/fuse-2.9.0 [106x35]
連線(C) 編輯(E) 檢視(V) 視窗(W) 選項(O) 說明(H)
danny@danny-desktop:~/tmp$ cd fuse/
danny@danny-desktop:~/tmp/fuse$ ls
bin      debug  home      lib      mnt      root      selinux  sys      usr      vmlinuz.old
boot    dev    initrd.img  lost+found  opt     /sbin      srv      tmp      var
cdrom   etc    initrd.img.old  media    proc    scratchbox  stuff    tracing  vmlinux
```


The End

Thank you for your listening