



Gluster Technical Overview

Kaleb KEITHLEY
Red Hat
11 April, 2015

Concepts:

- Translator — a shared library that implements storage semantics
 - e.g. distribution, replication, performance, etc.
- Translator Stack — a directed graph of translators
- Brick — a server and a directory
- Volume — a collection of bricks



Gluster Processes

- Server: glusterfsd — the heart of it all
- Clients:
 - FUSE, glusterfs
 - NFS server, glusterfs (gluster client, NFS server)
- Management: glusterd
- Healing: glustershd
- And more——



Let's get real — creating a volume

```
Srv1% gluster peer probe srv1
```

```
...
```

```
srv1% gluster volume create my_vol
```

```
srv1:/bricks/my_vol
```

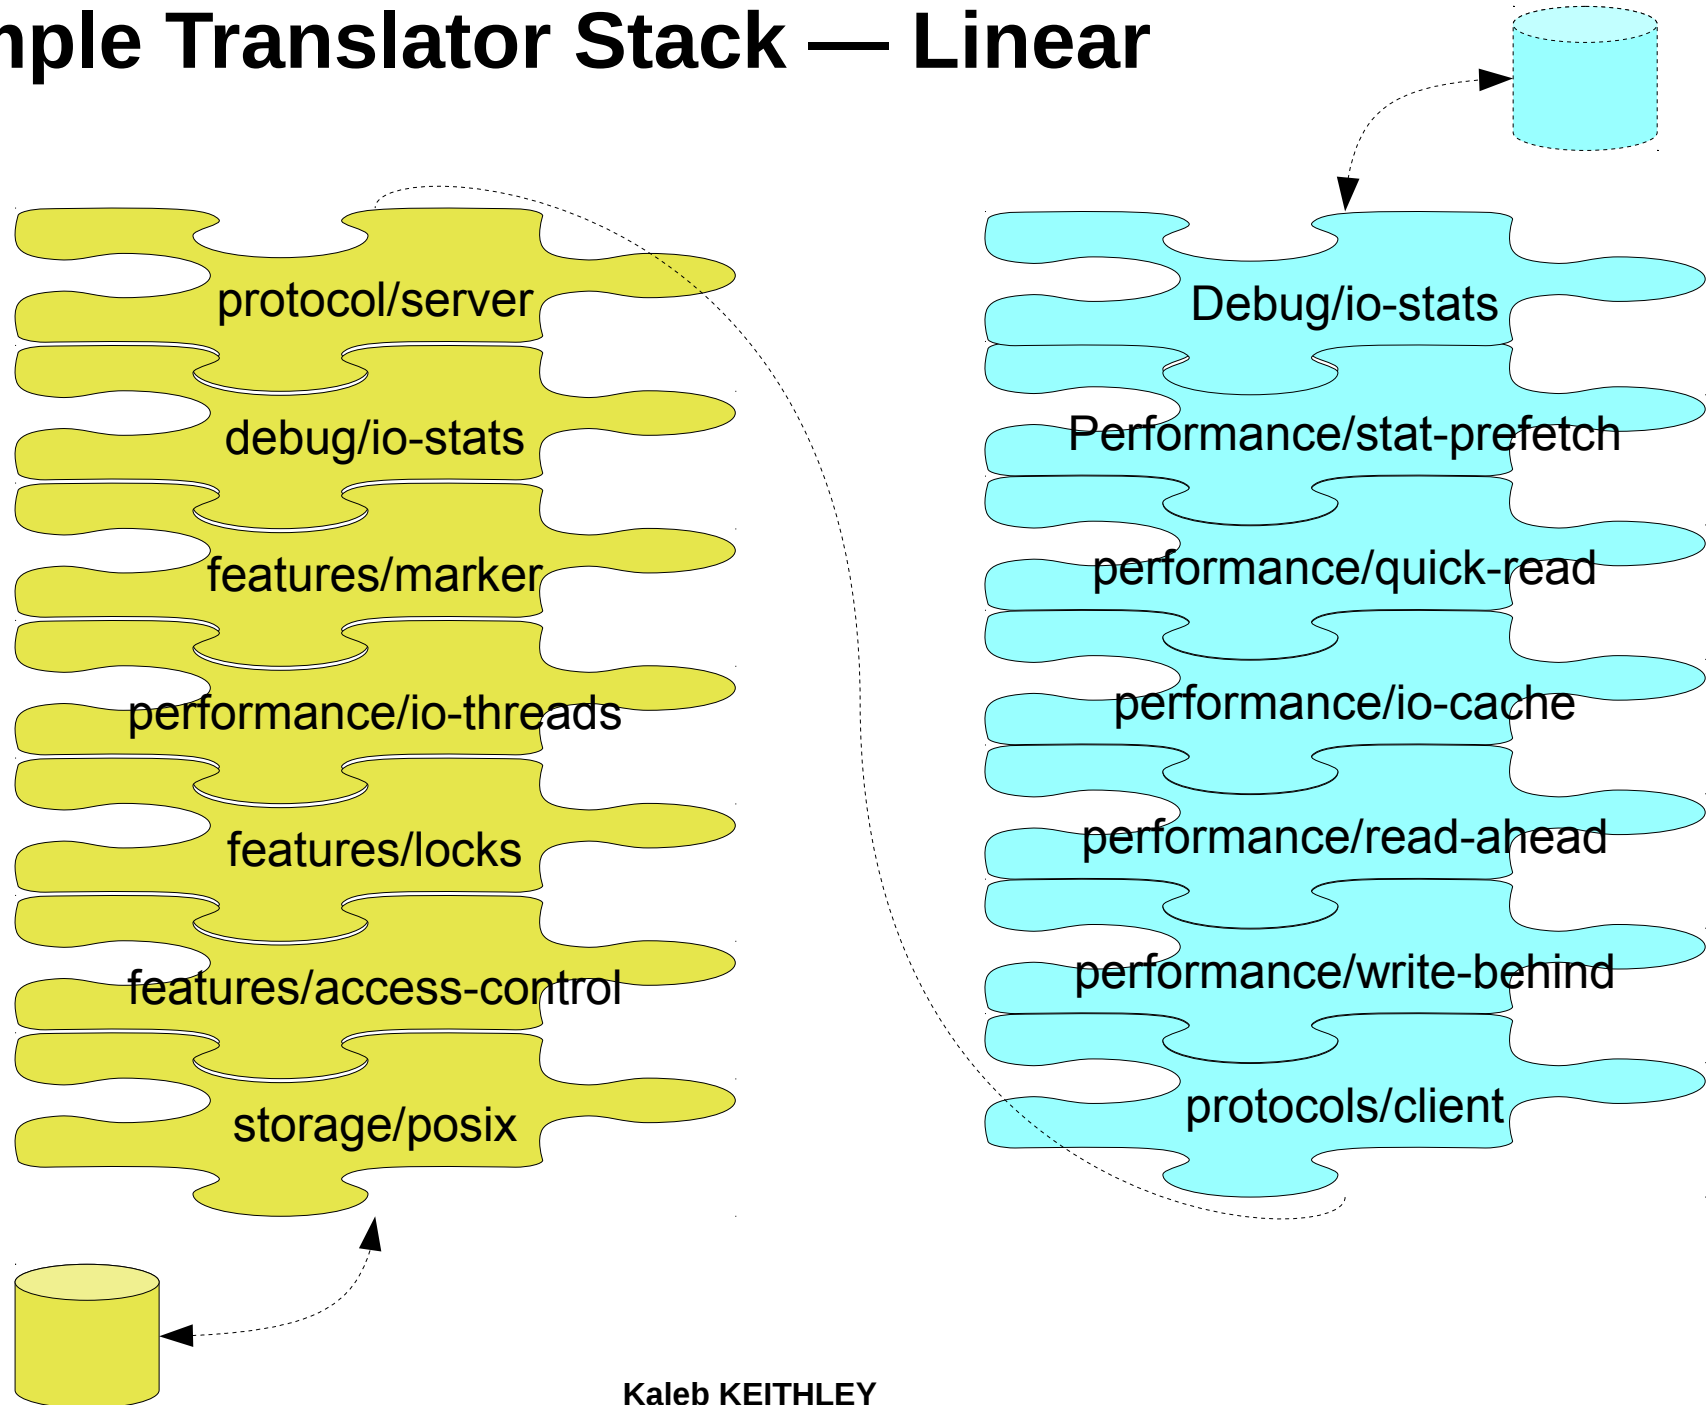
```
srv1% gluster volume start my_vol
```

```
...
```

```
client1% mount -t glusterfs srv1:my_vol  
/mnt
```

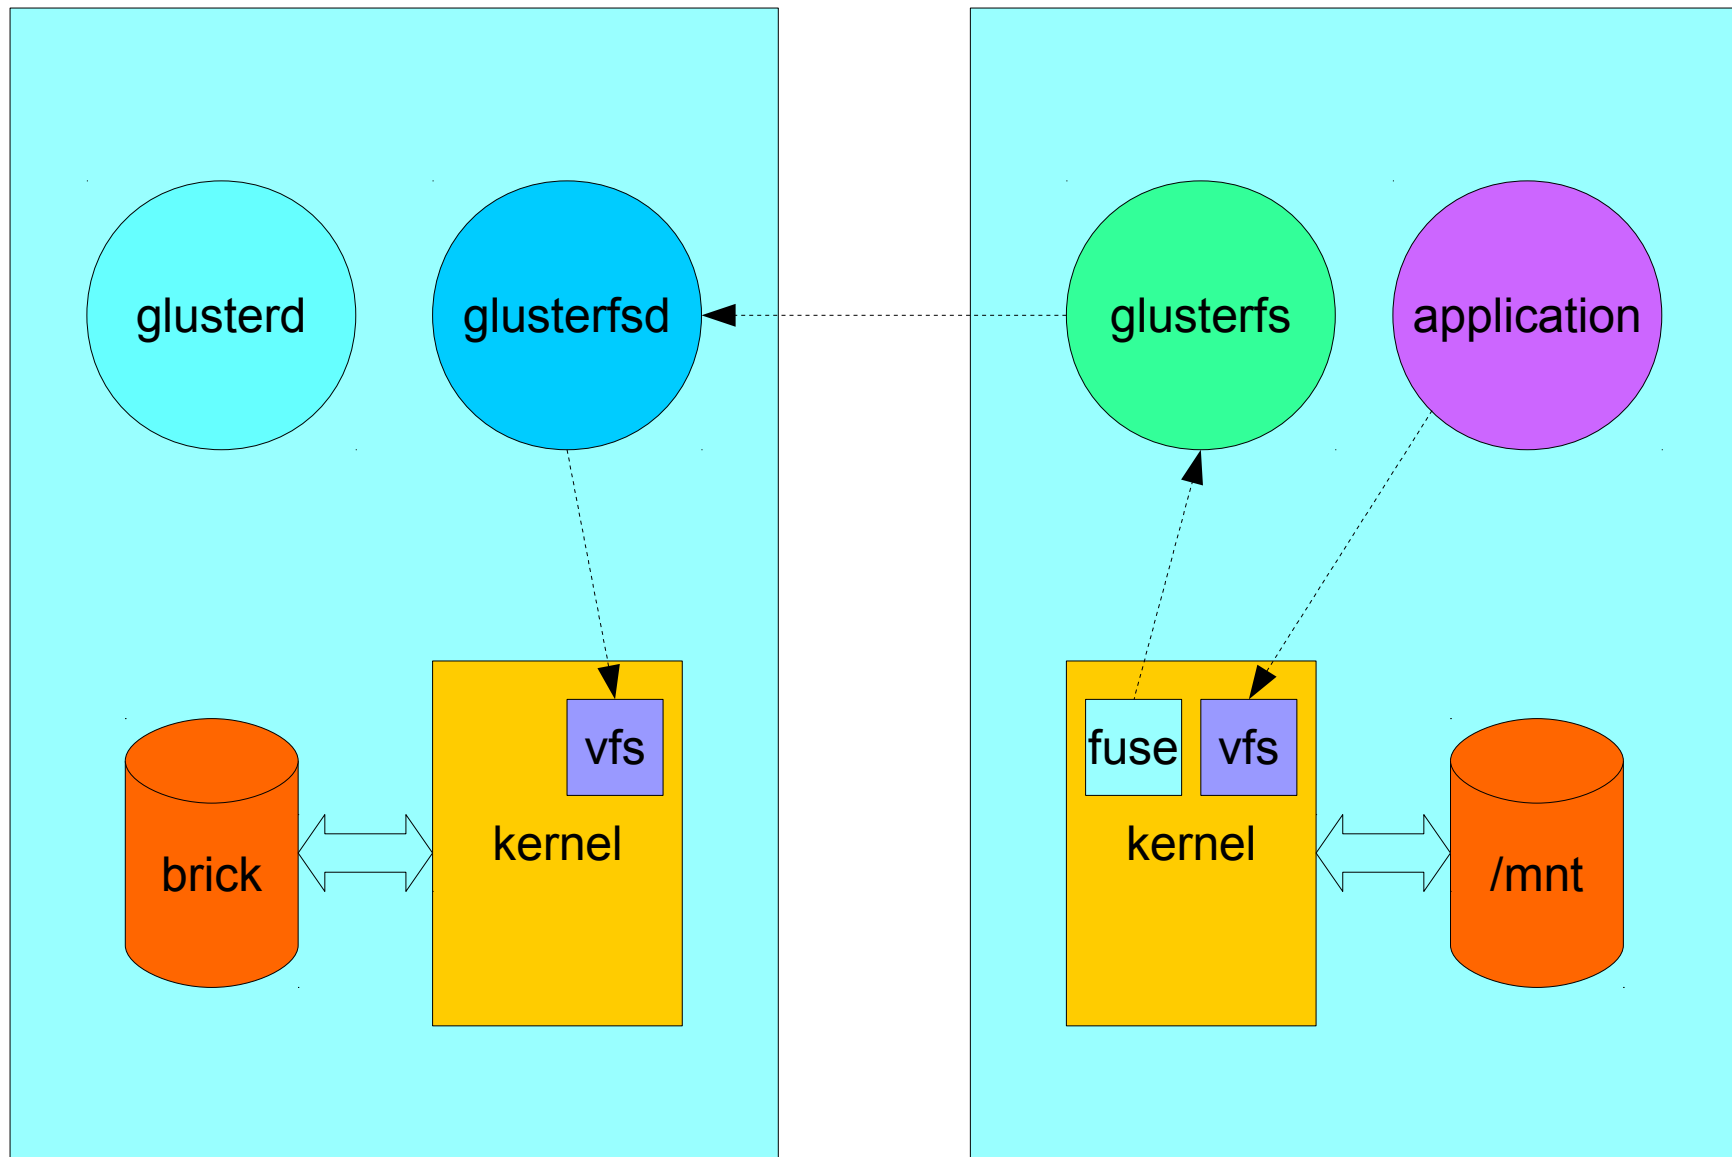


Simple Translator Stack — Linear



server

client



Let's get real — creating a distribute volume

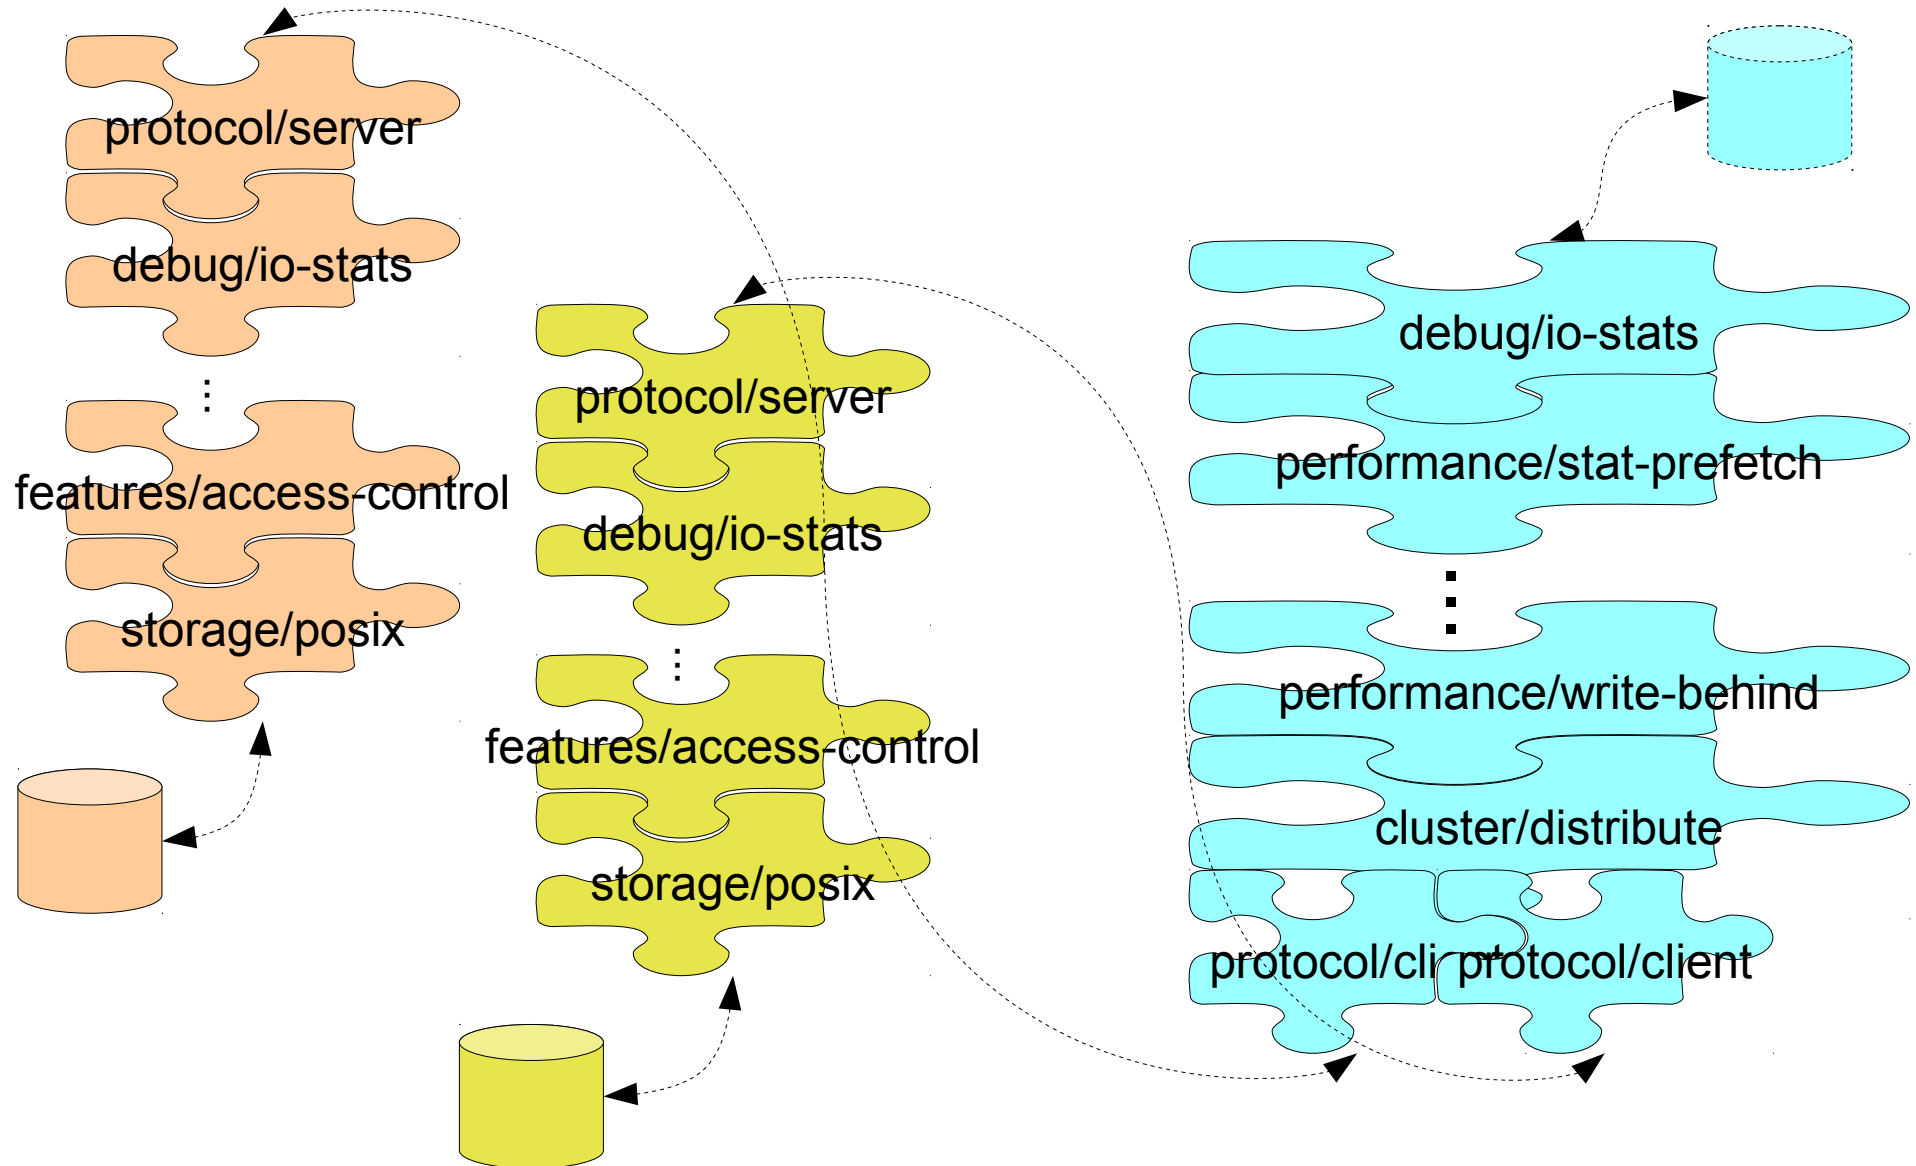
```
srv1% gluster volume create my_dist  
srv1:/bricks/my_vol srv2:/bricks/my_vol  
srv1% gluster volume start my_dist
```


...

```
client1% mount -t glusterfs srv1:my_dist /mnt
```

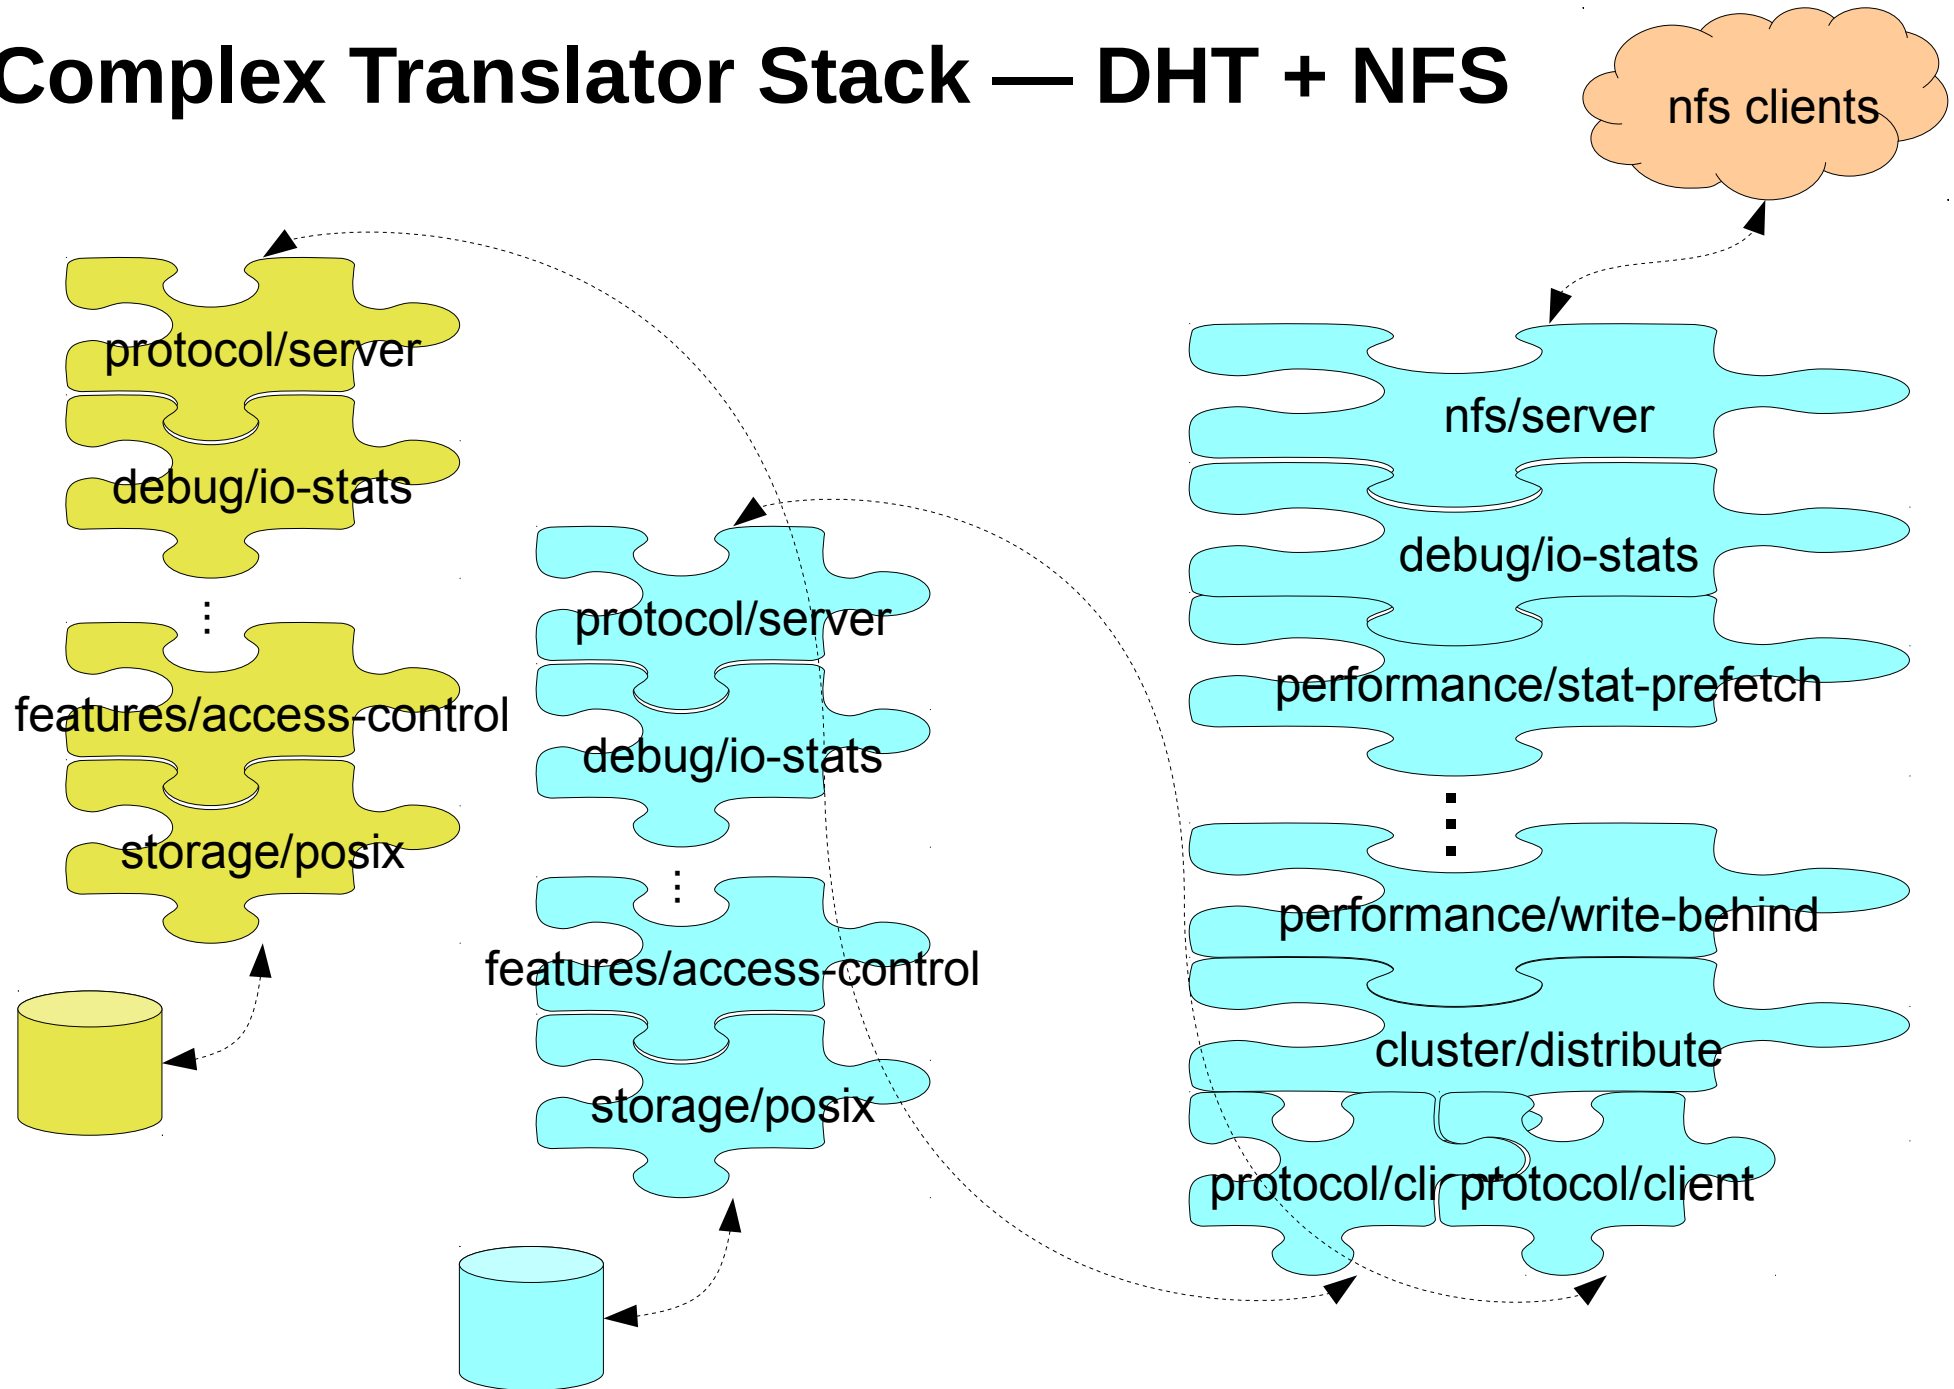


Complex Translator Stack — Distribute (DHT)





Complex Translator Stack — DHT + NFS



Let's get real — creating a replica volume

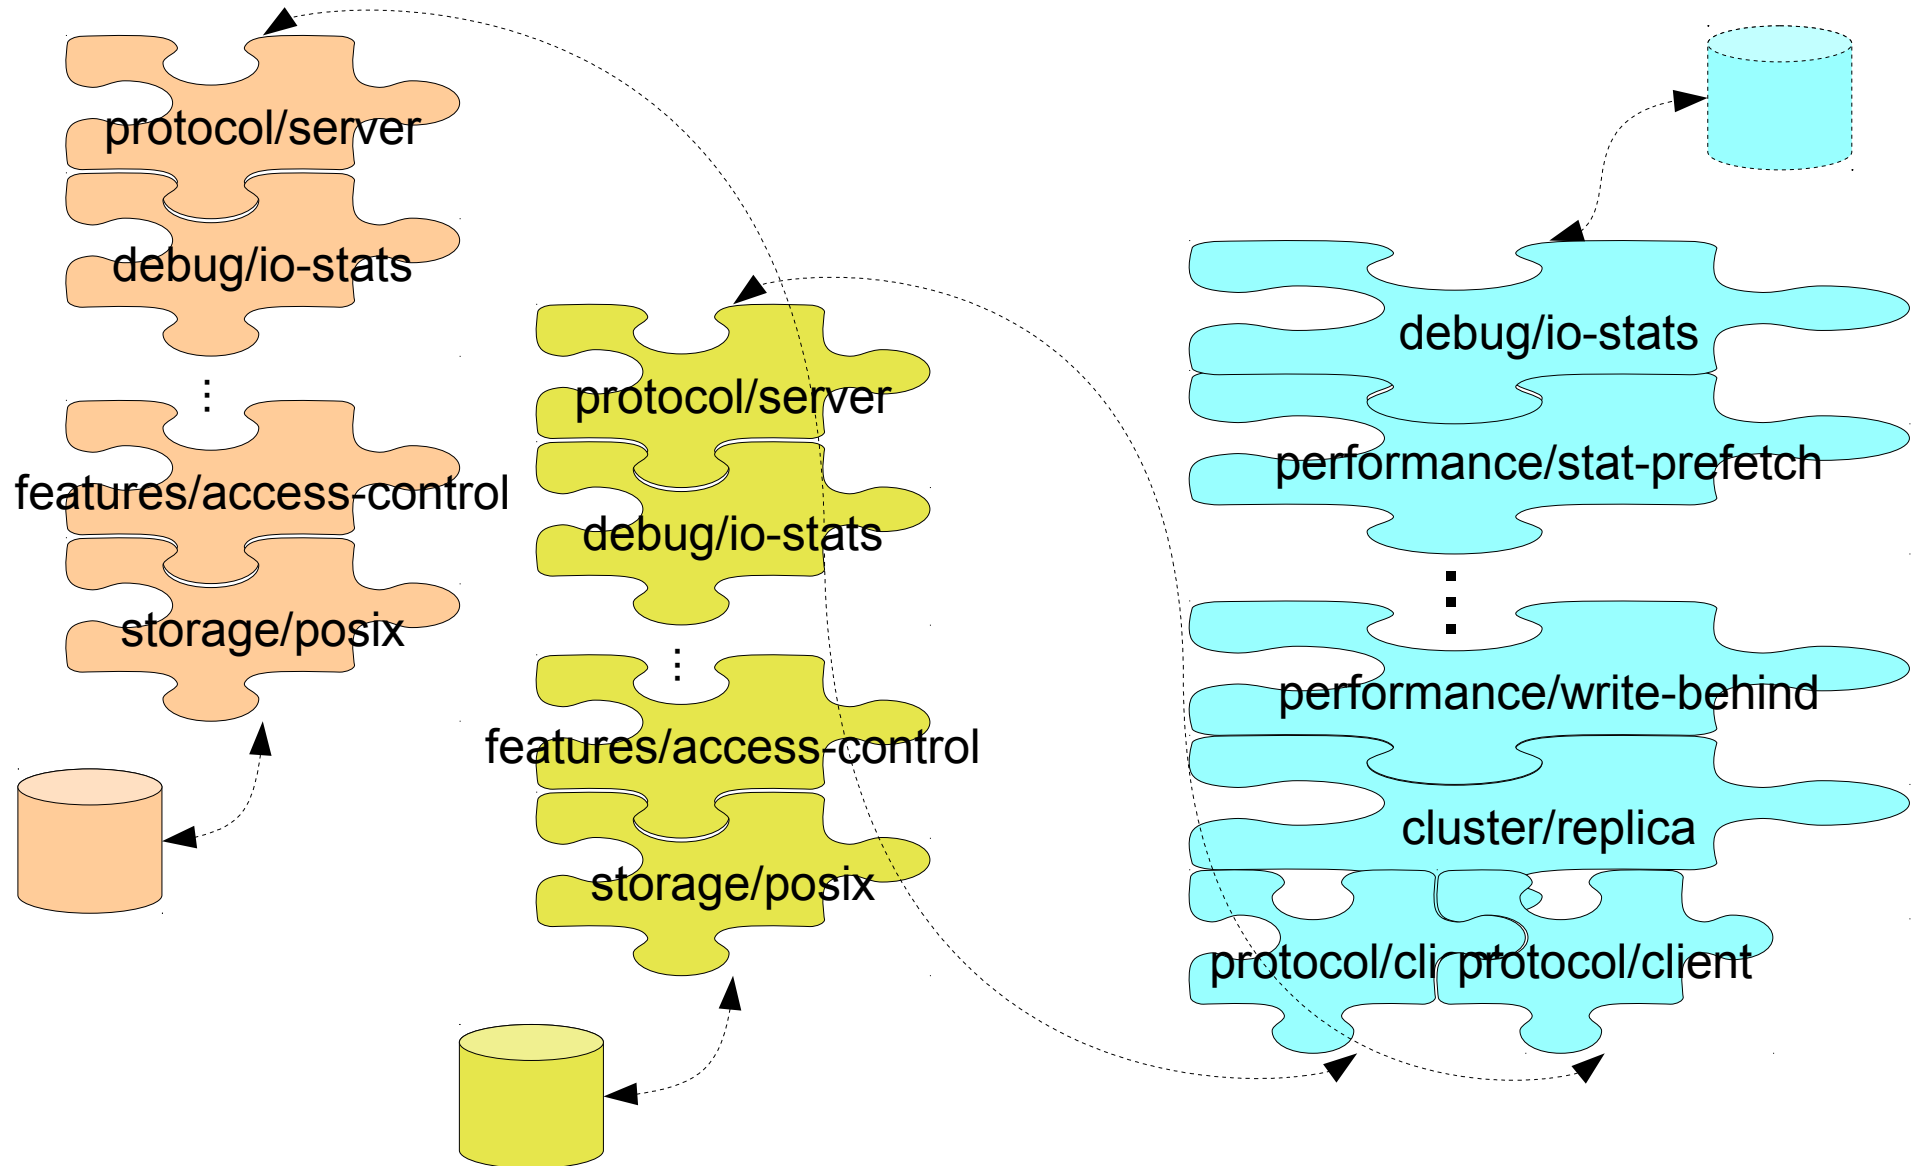
```
srv1% gluster volume create replica 2 my_repl  
srv1:/bricks/my_vol srv2:/bricks/my_vol  
  
srv1% gluster volume start my_repl
```

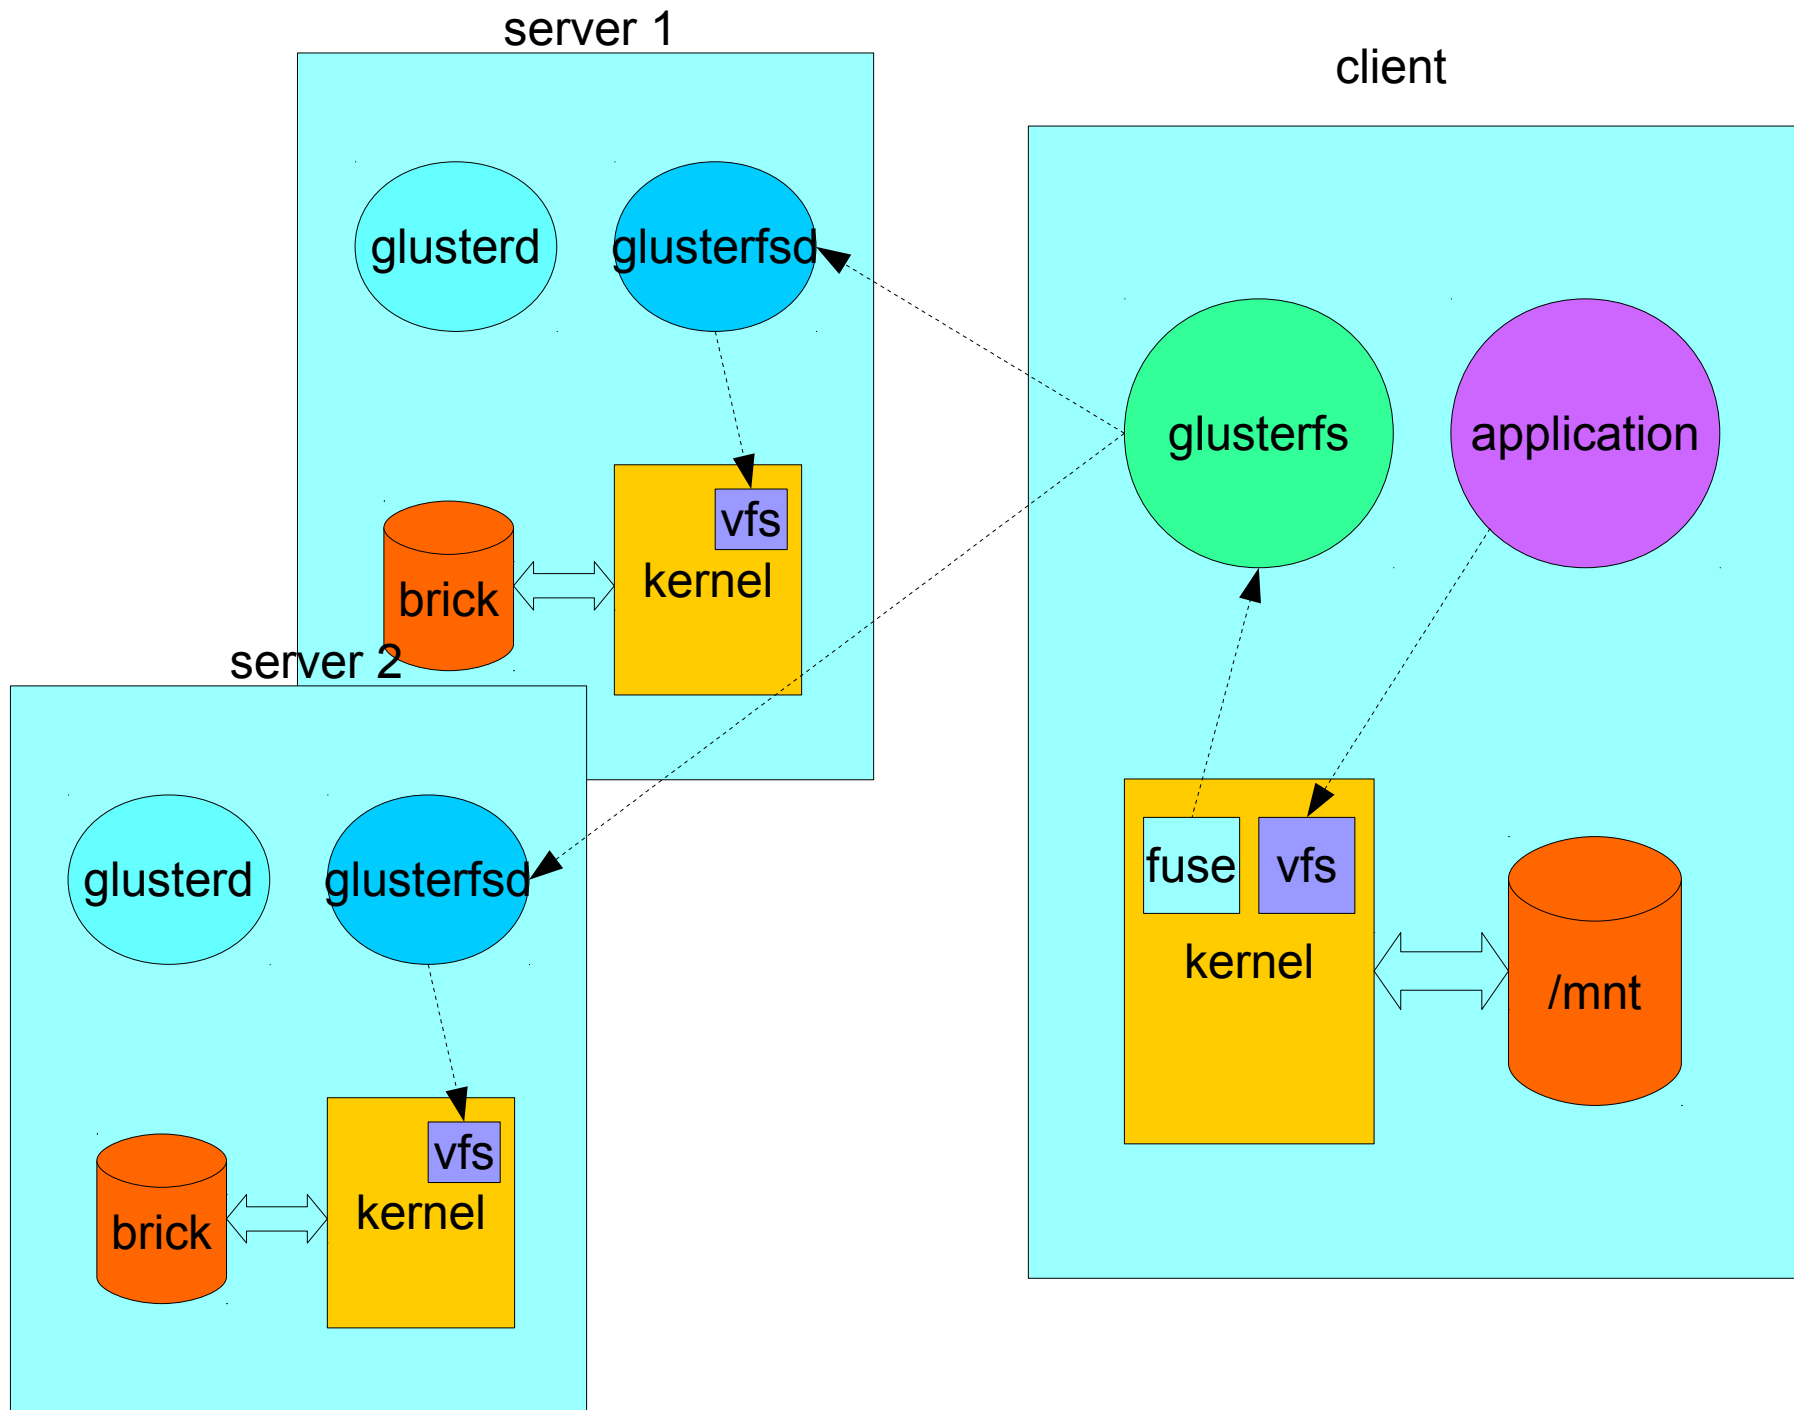
...

```
client1% mount -t glusterfs srv1:my_repl /mnt
```



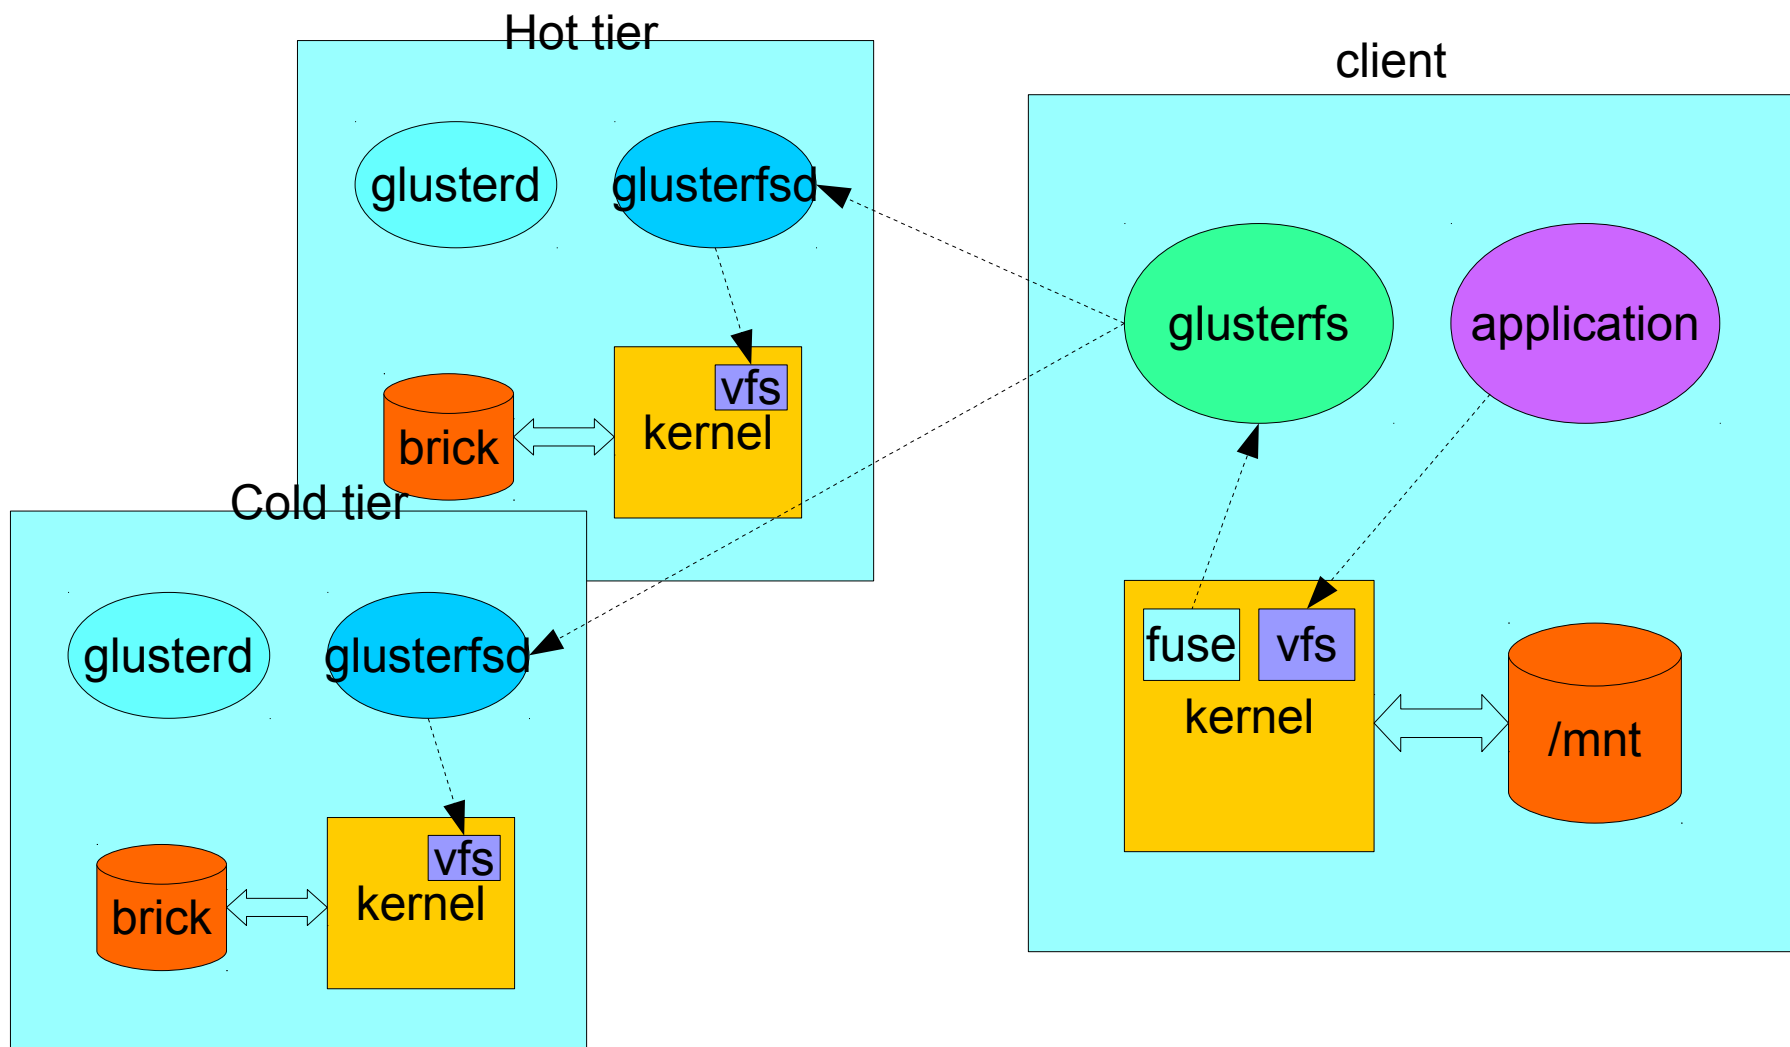
Complex Translator Stack — Replica (AFR)



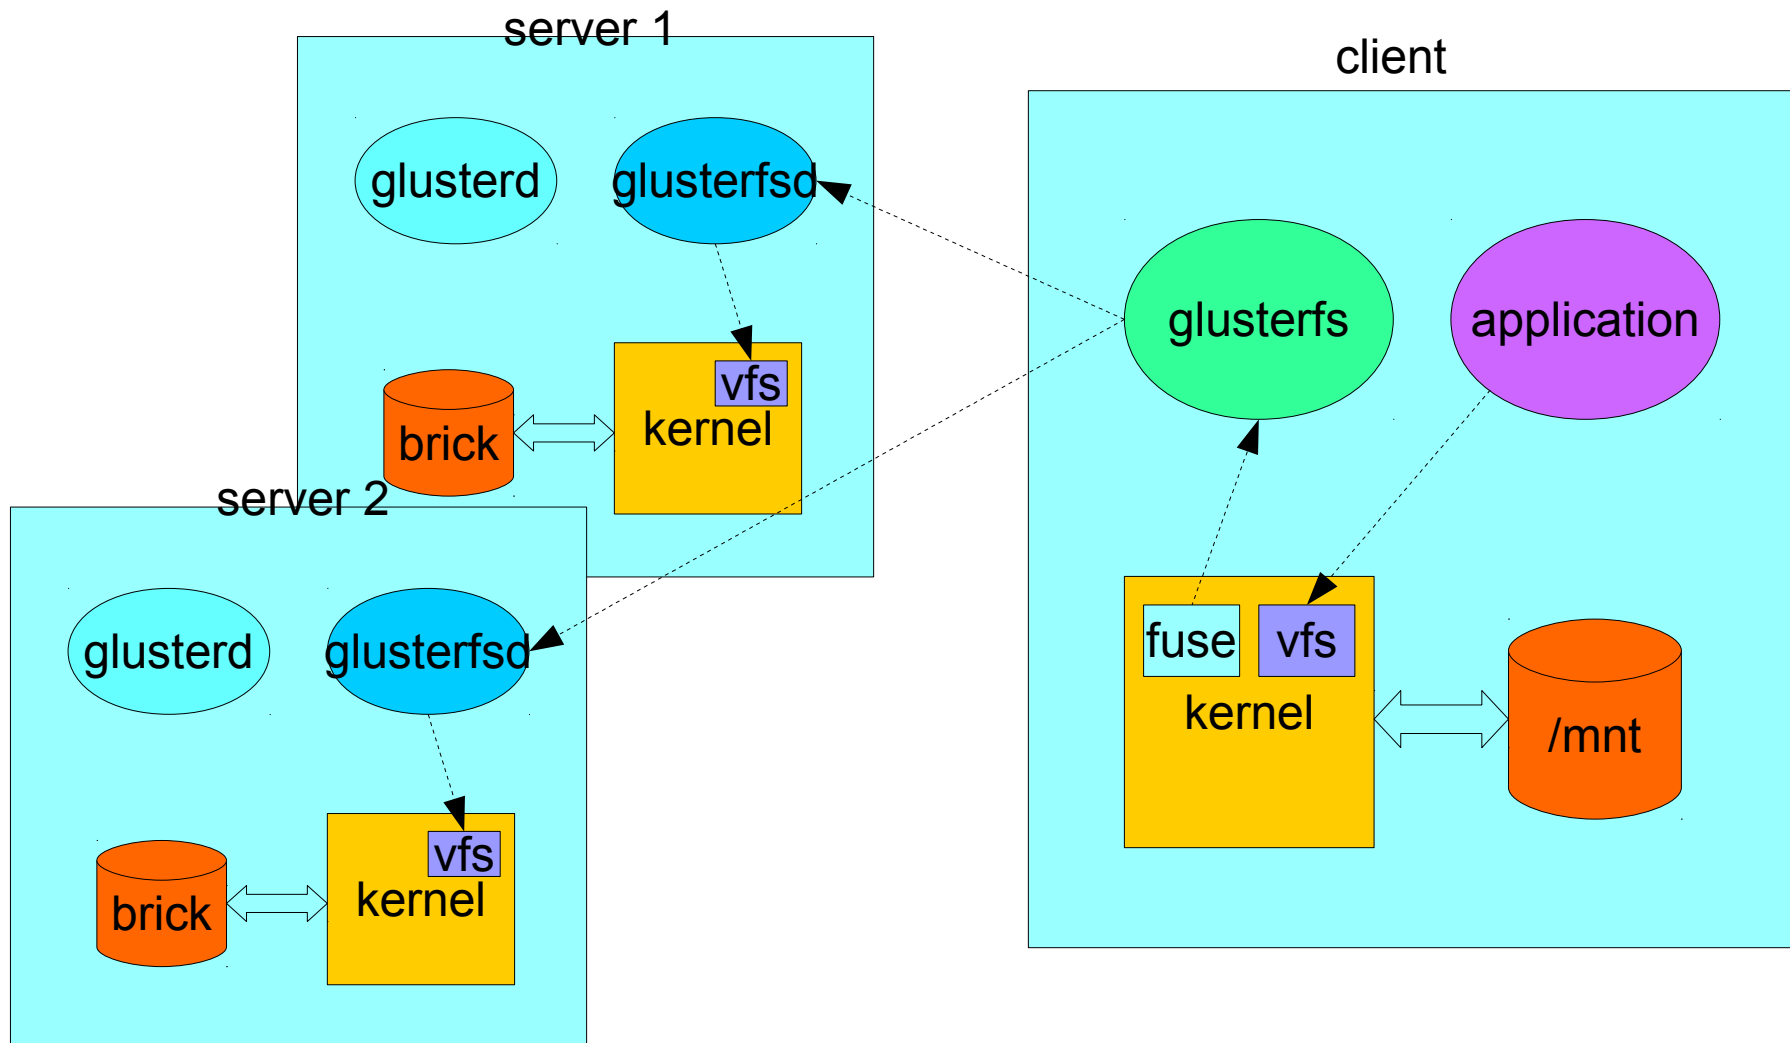



Tiering: a variation on distribution

- Hot and cold data




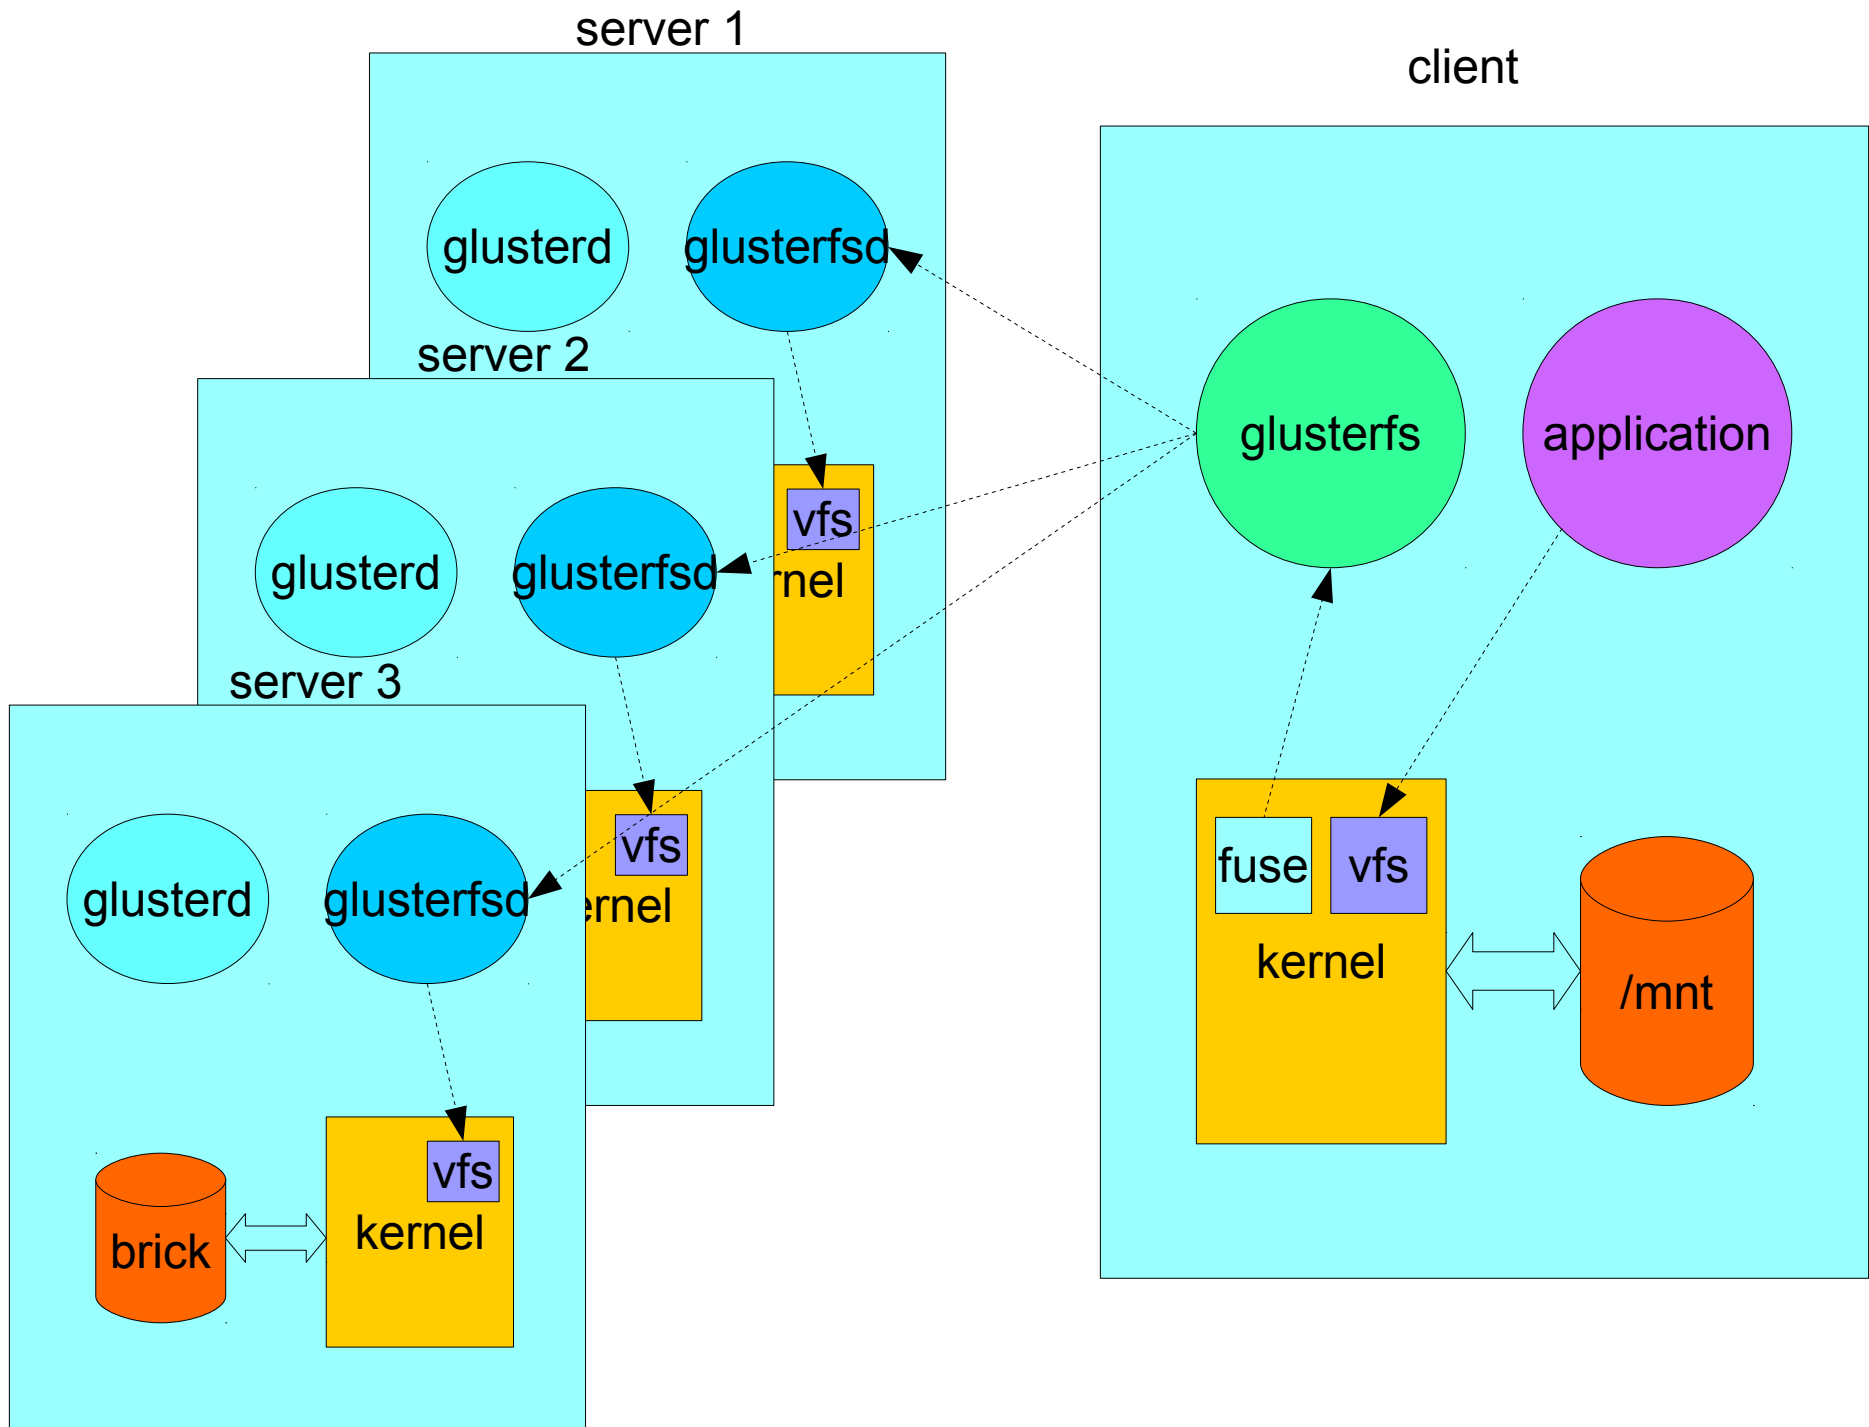
Adding a brick — distributed and rebalancing






Adding a brick — replica and healing







Isn't FUSE slow?




- Count the copies
- And context switches



gfapi to the rescue



- POSIX-like API
 - glfs_open()
 - glfs_close()
 - glfs_read(), glfs_write()
 - glfs_seek()
 - glfs_stat()
 - etc.



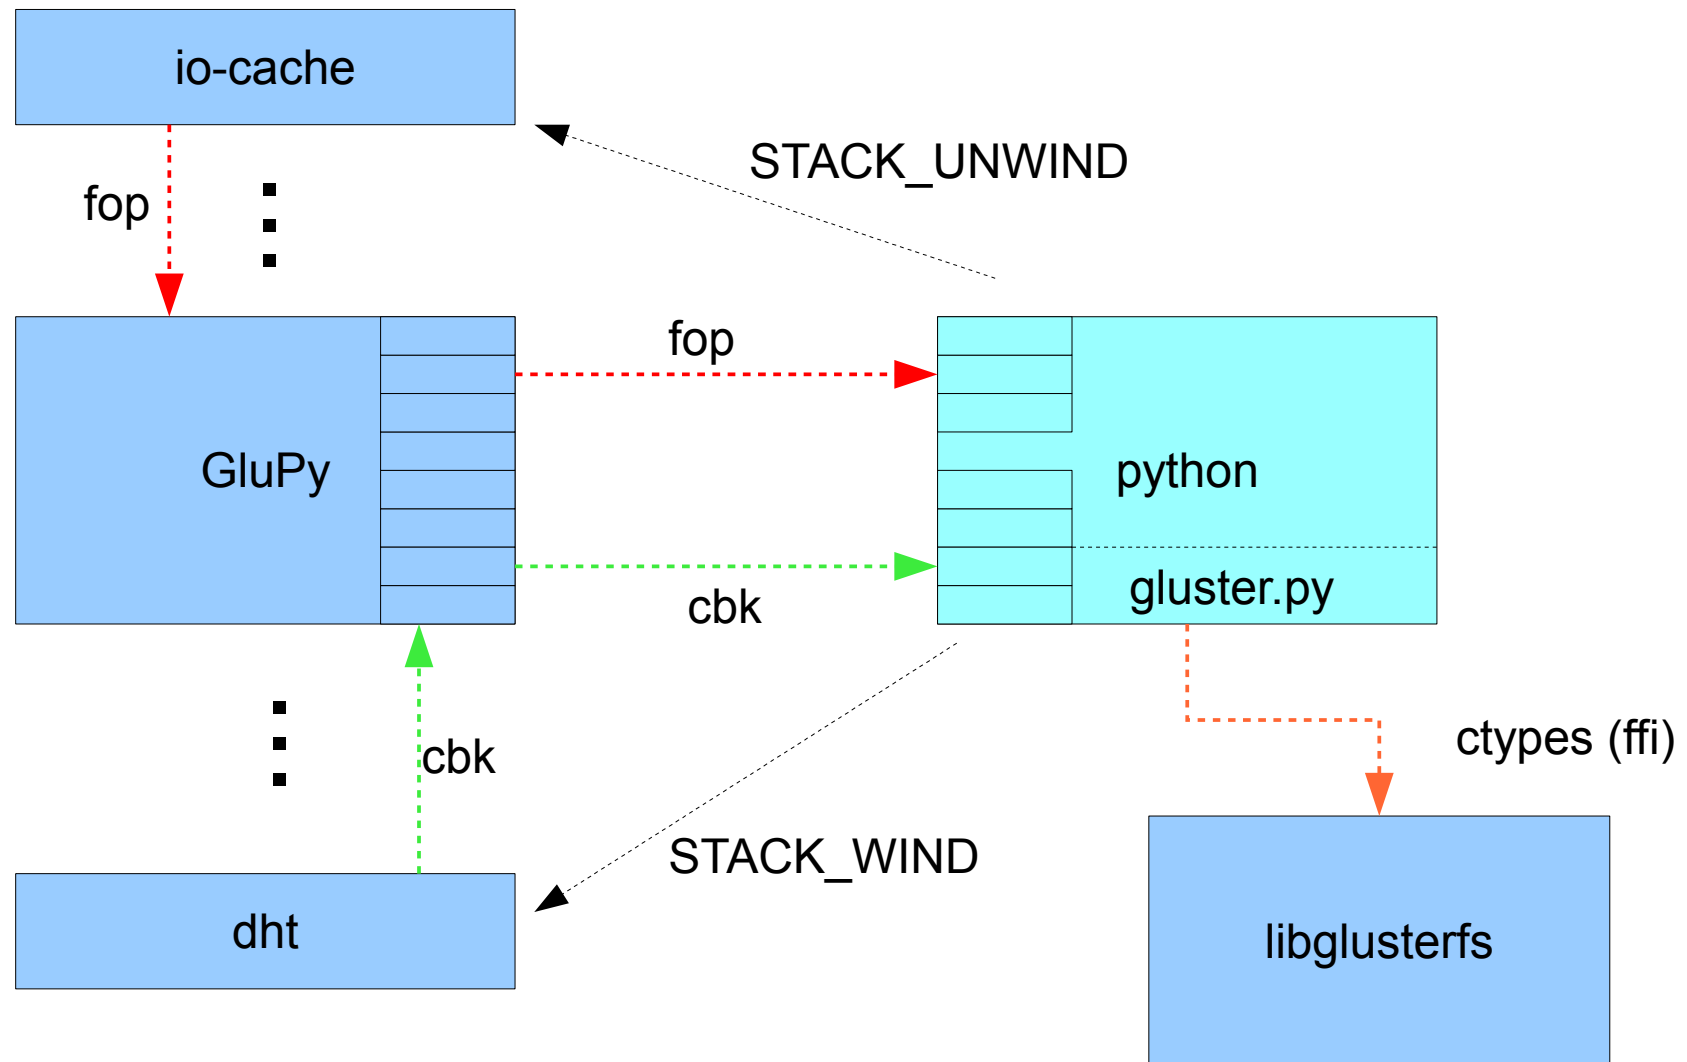
libgfapi — hellogluster.c

```
#include <gluster/api/glfs.h>

int main (int argc, char** argv)
{
    fs = glfs_new ("fsync");
    ret = glfs_set_volfile_server (fs, "tcp", "localhost",
    24007);
    /* or ret = glfs_set_volfile (fs, "/tmp/foo.vol"); */
    ret = glfs_init (fs);
    fd = glfs_creat (fs, filename, O_RDWR, 0644);
    ret = glfs_write (fd, "hello gluster", 14, 0);
    glfs_close (fd);
    return 0;
}
```



Writing a translator in Python with GluPy



SwiftOnFile — RESTful API, examples with curl

- Create a container: `curl -v -X PUT -H 'X-Auth-Token: $authtoken' https://$myhostname:443/v1/AUTH_$myvolname/$mycontainername -k`
- List containers: `curl -v -X GET -H 'X-Auth-Token: $authtoken' https://$myhostname:443/v1/AUTH_$myvolname -k`
- Copy a file into a container (upload): `curl -v -X PUT -T $filename -H 'X-Auth-Token: $authtoken' -H 'Content-Length: $filelen' https://$myhostname:443/v1/AUTH_$myvolname/$mycontainername/$filename -k`
- Copy a file from a container (download): `curl -v -X GET -H 'X-Auth-Token: $authtoken' https://$myhostname:443/v1/AUTH_$myvolname/$mycontainername/$filename -k > $filename`



Translator basics

- Translators are shared objects (shlibs)
 - Methods
 - `int32_t init(xlator_t *this);`
 - `void fini(xlator_t *this);`
 - Data
 - `struct xlator_fops fops { ... };`
 - `struct xlator_cbks cbks { };`
 - `struct volume_options options [] = { ... };`
- Client, Server, Client/Server
- Threads: write MT-SAFE
- Portability: GlusterFS != Linux only
- License: GPLv2 or LGPLv3+



Every method has a different signature

- Open fop method and callback

```
typedef int32_t (*fop_open_t) (call_frame_t *, xlator_t *,  
loc_t *, int32_t, fd_t *, dict_t *);
```

```
typedef int32_t (*fop_open_cbk_t) (call_frame_t *, void *,  
xlator_t *, int32_t, int32_t, fd_t *, dict_t *);
```

- Rename fop method and callback

```
typedef int32_t (*fop_rename_t) (call_frame_t *, xlator_t *,  
loc_t *, loc_t *, dict_t *);
```

```
typedef int32_t (*fop_rename_cbk_t) (call_frame_t , void *,  
xlator_t *, int32_t, int32_t, struct iatt *, struct iatt *,  
struct iatt *, struct iatt *, struct iatt *);
```



Data Types in Translators

- `call_frame_t` —
- `xlator_t` — translator context
- `inode_t` — represents a file on disk; ref-counted
- `fd_t` — represents an open file; ref-counted
- `iatt_t` — \sim struct stat
- `dict_t` — \sim Python dict (or C++ `std::map`)
- `client_t` — represents the connect client



fop methods and fop callbacks

```
uidmap_writev (...)
{
    ...
    STACK_WIND (frame, uidmap_writev_cbk,
                FIRST_CHILD (this),
                FIRST_CHILD (this)->fops->writev,
                fd, vector, count, offset, iobref);
    /* DANGER ZONE */
    return 0;
}
```

- Effectively lose control after STACK_WIND
 - Callback might have already happened
 - Or might be running right now
 - Or maybe it's not going to run 'til later



fop methods and fop callback methods, cont.

```
uidmap_writev_cbk (call_frame_t *frame, void *cookie, ...)
{
    ...
    STACK_UNWIND_STRICT (writev, frame
        op_ret, op_errno, prebuf, postbuf);
    return 0;
}
```

- The I/O is complete when the callback is called



STACK_WIND versus STACK_WIND_COOKIE

- Pass extra data to the cbk with STACK_WIND_COOKIE

```
quota_statfs (call_frame_t *frame,  
              xlator_t *this, loc_t *loc)  
{  
    inode_t *root_inode = loc->inode->table->root;  
    STACK_WIND_COOKIE (frame, quota_statfs_cbk,  
                       root_inode, FIRST_CHILD (this),  
                       FIRST_CHILD (this)->fops->statfs, loc, xdata);  
    return 0;  
}
```

- There is also frame->local
 - shared by all STACK_WIND callbacks



STACK_WIND, STACK_WIND_COOKIE, cont.

- Pass extra data to the cbk with
STACK_WIND_COOKIE

```
quota_statfs_cbk (call_frame_t *frame, void *cookie, ...)  
{  
    inode_t *root_inode = cookie;  
    ...  
}
```



STACK_UNWIND versus STACK_UNWIND_STRICT

- STACK_UNWIND_STRICT uses the correct type

```
/* return from function in a type-safe way */  
#define STACK_UNWIND (frame, params ...)   
    do {   
        ret_fn_t fn = frame->ret;   
        ...
```

versus

```
#define STACK_UNWIND_STRICT (op, frame, params ...)   
    do {   
        fop_##op##_cbk_t fn = (fop_##op##_cbk_t)frame->ret;   
        ...
```

- And why wouldn't you want strong typing?



Calling multiple children (fan out)

```
afr_writev_wind (...)  
{  
    ...  
    for (i = 0; i < priv->child_count; i++) {  
        if (local->transaction.pre_op[i]) {  
            STACK_WIND_COOKIE (frame, afr_writev_wind_cbk,  
                                (void *) (long) i,  
                                priv->children[i],  
                                priv->children[i]->fops->writev,  
                                local->fd, ...);  
        }  
    }  
    return 0;  
}
```



Calling multiple children, cont. (fan in)

```
afr_writev_wind_cbk (...)  
{  
    LOCK (&frame->lock);  
    callcnt = --local->call_count;  
    UNLOCK (&frame->lock);  
    if (callcnt == 0) /* we're done */  
        ...  
}
```

- failure by any one child means the whole transaction failed?
 - And needs to be handled accordingly



Dealing With Errors: I/O errors

```
uidmap_writev_cbk (call_frame_t *frame, void *cookie,  
    xlator_t *this, int32_t op_ret, int32_t op_errno, ...)  
{  
    ...  
    STACK_UNWIND_STRICT (writev, frame, -1, EIO, ...);  
    return 0;  
}
```

- op_ret: 0 or -1, success or failure
- op_errno: from <errno.h>
 - Use an op_errno that's valid and/or relevant for the fop



Dealing With Errors: FOP method errors

```
uidmap_writev (call_frame_t *frame, xlator_t *this, ...)
{
    ...
    if (horrible_logic_error_must_abort) {
        goto error; /* glusterfs idiom */
    }
    STACK_WIND(frame, uid_writev_cbk, ...);
    return 0;
error:
    STACK_UNWIND_STRICT (writev, frame, -1, EIO, NULL, NULL);
    return 0;
}
```



Call to action

- Go forth and write applications for GlusterFS!



£©€é¥¢βƆ

