



# Deploying pNFS solution over Distributed Storage

Jiffin Tony Thottan  
Niels de Vos

# Agenda

- pNFS protocol
- NFS-Ganesha
- GlusterFS
- Integration
- Challenges
- Configuring pNFS cluster
- Demo




# pNFS Protocol

## Overview

- pNFS is introduced as part of nfsv4.1 (RFC5661) in 2006
- clients access storage devices directly and in parallel
- Data and Metadata handled in two different paths



## Basic pNFS architecture



# pNFS terminologies

- MDS – Meta Data Server
  - NFSv4.1 server that supports the pNFS protocol. It provides access to the name space. Also handles I/O in case of failures
- Storage Devices
  - where actual data resides
- Storage Protocol
  - Used between client and Storage devices. It can be NFS, iSCSI, OSD etc
- Control Protocol
  - It maintains cluster coherence and it resides out side of scope standard NFS protocol.



## pNFS Layouts

Provides ability to access data for the clients, four types :

- File Layout (*mentioned in RFC5661*)
- Block Layout (*mentioned in RFC5663*)
- Object Layout (*mentioned in RFC5664*)
- Flexfile Layout (*<https://tools.ietf.org/html/draft-ietf-nfsv4-flex-files-07>*)



# pNFS Operations

Following operations are performed from client to MDS :

- GETDEVICEINFO(*device id*)
  - gets information about storage devices
- LAYOUTGET(*file handle, offset, length*)
  - fetch file information about data in the form of layout
- LAYOUTRETURN(*file handle, offset, length, stateid*)
  - releases the layout
- LAYOUTCOMMIT(*file handle, clientid, range, stateid*)
  - commits write using layout to the MDS



## **pNFS call back operation**

Following are notifications send from MDS to client :

- CB\_LAYOUTRECALL
  - recalls layout granted to a client
- CB\_RECALLABLE\_OBJ\_AVAIL
  - previously denied layout is available
- CB\_NOTIFY\_DEVICEID
  - informs client device id is invalid
- CB\_RECALL\_ANY
  - recalls delegations/layouts whose state can hold by the server





## **nfsv4.1 as Storage Protocol**

If storage devices is nfsv4.1 server(Data Server) ,  
following additional ops should be defined

- ds\_write
- ds\_read
- ds\_commit

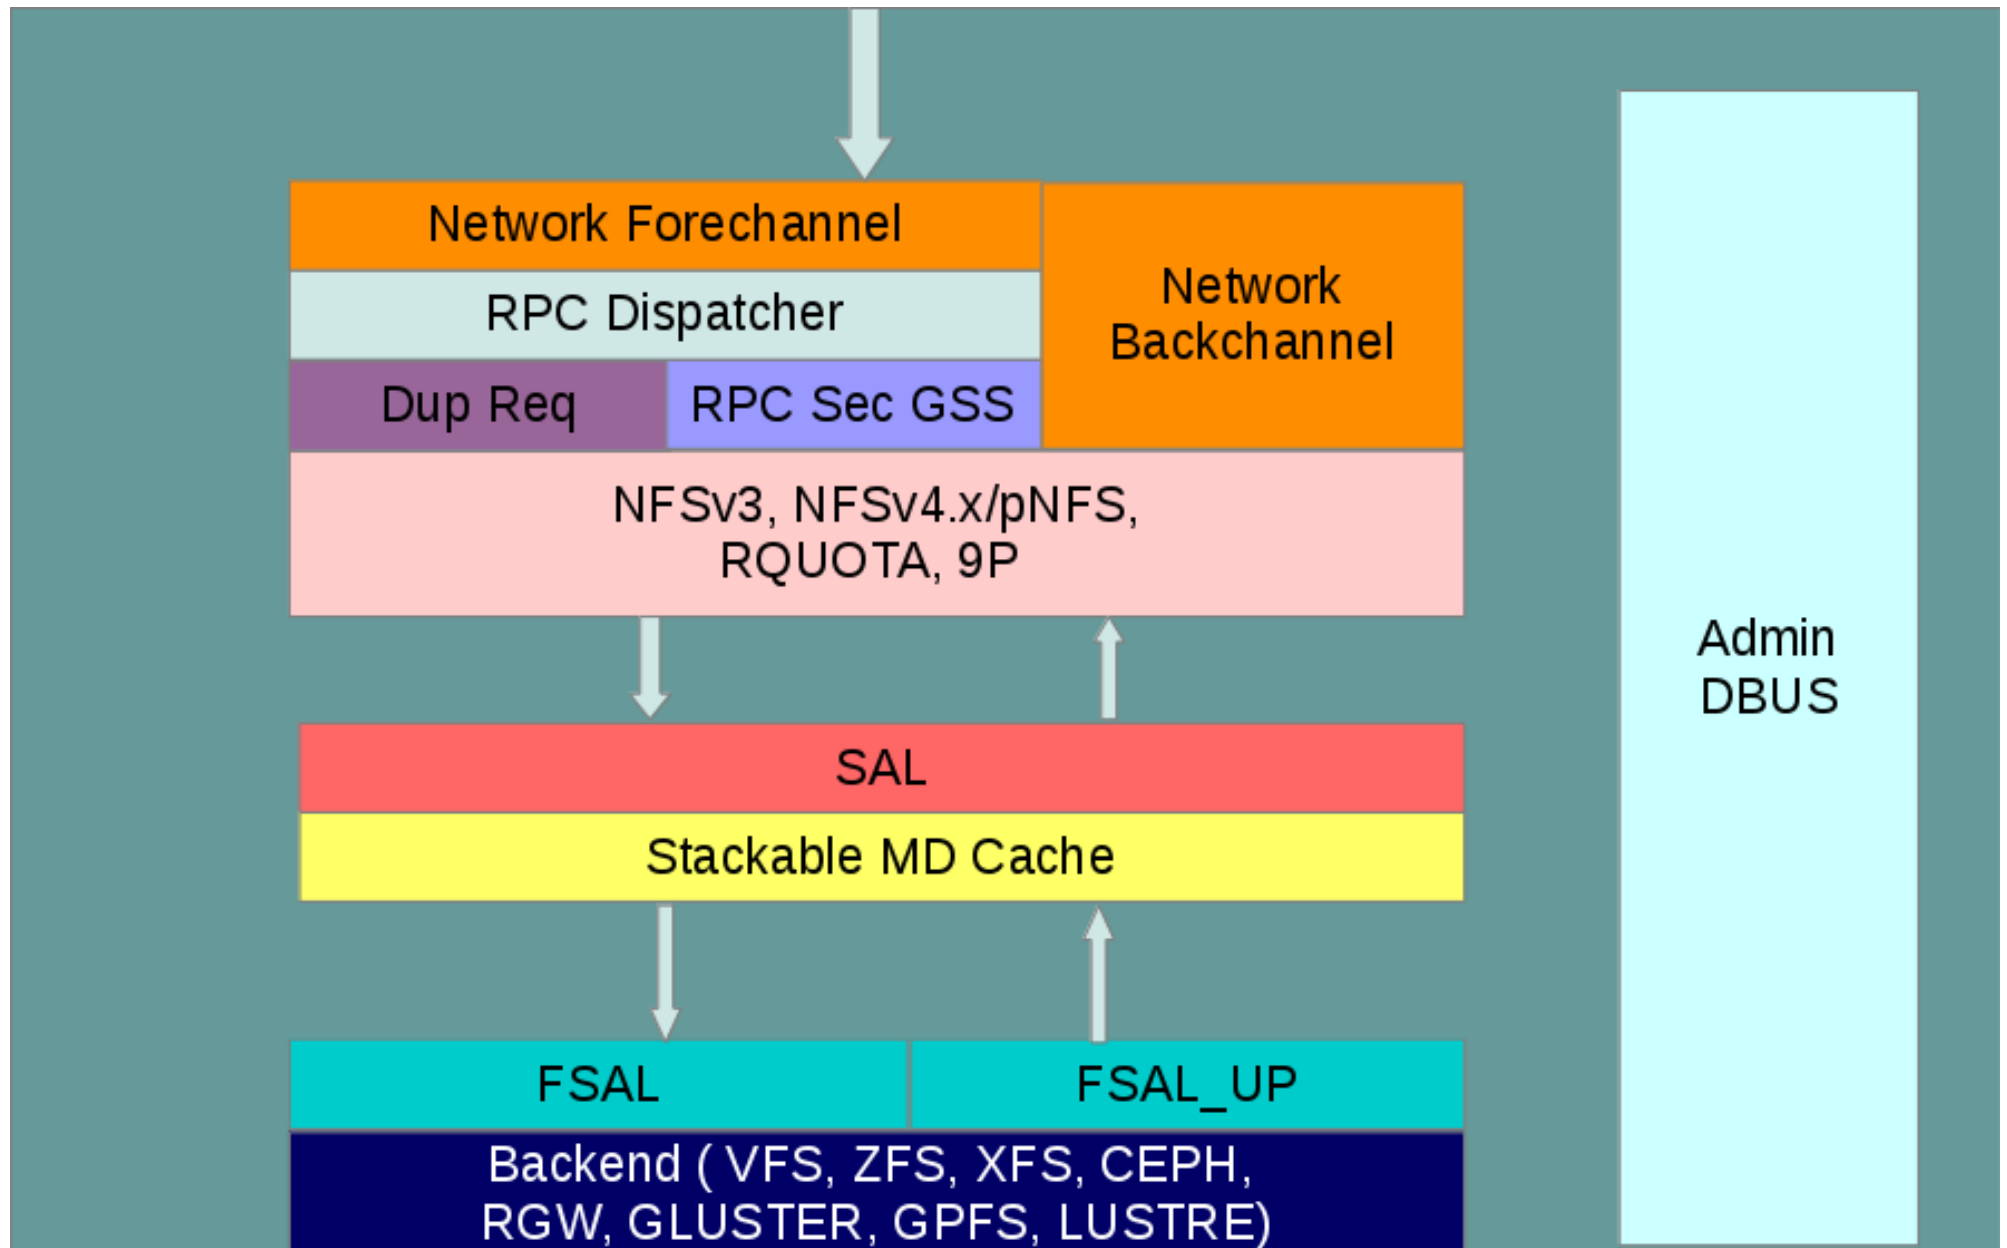


# NFS-Ganesha

- A user-space, protocol-complaint NFS server
- Supports NFS v3, 4.0, 4.1, pNFS and 9P from the Plan9 operating system.
- Provides a File System Abstraction Layer(FSAL) to plug in to any own storage mechanism
- Can provide simultaneous access to multiple file systems.
- Small but active and growing community ; CEA, Red Hat, IBM are active participants



# NFS-Ganesha architecture



## **Benefits of NFS-Ganesha**

- Can manage huge meta-data caches
- Dynamically export/unexport entries using D-Bus mechanism.
- Easy access to the services operating in the user-space (like Kerberos, NIS, LDAP)
- Provides better security and authentication mechanism for enterprise use
- Portable to any Unix-like file-systems

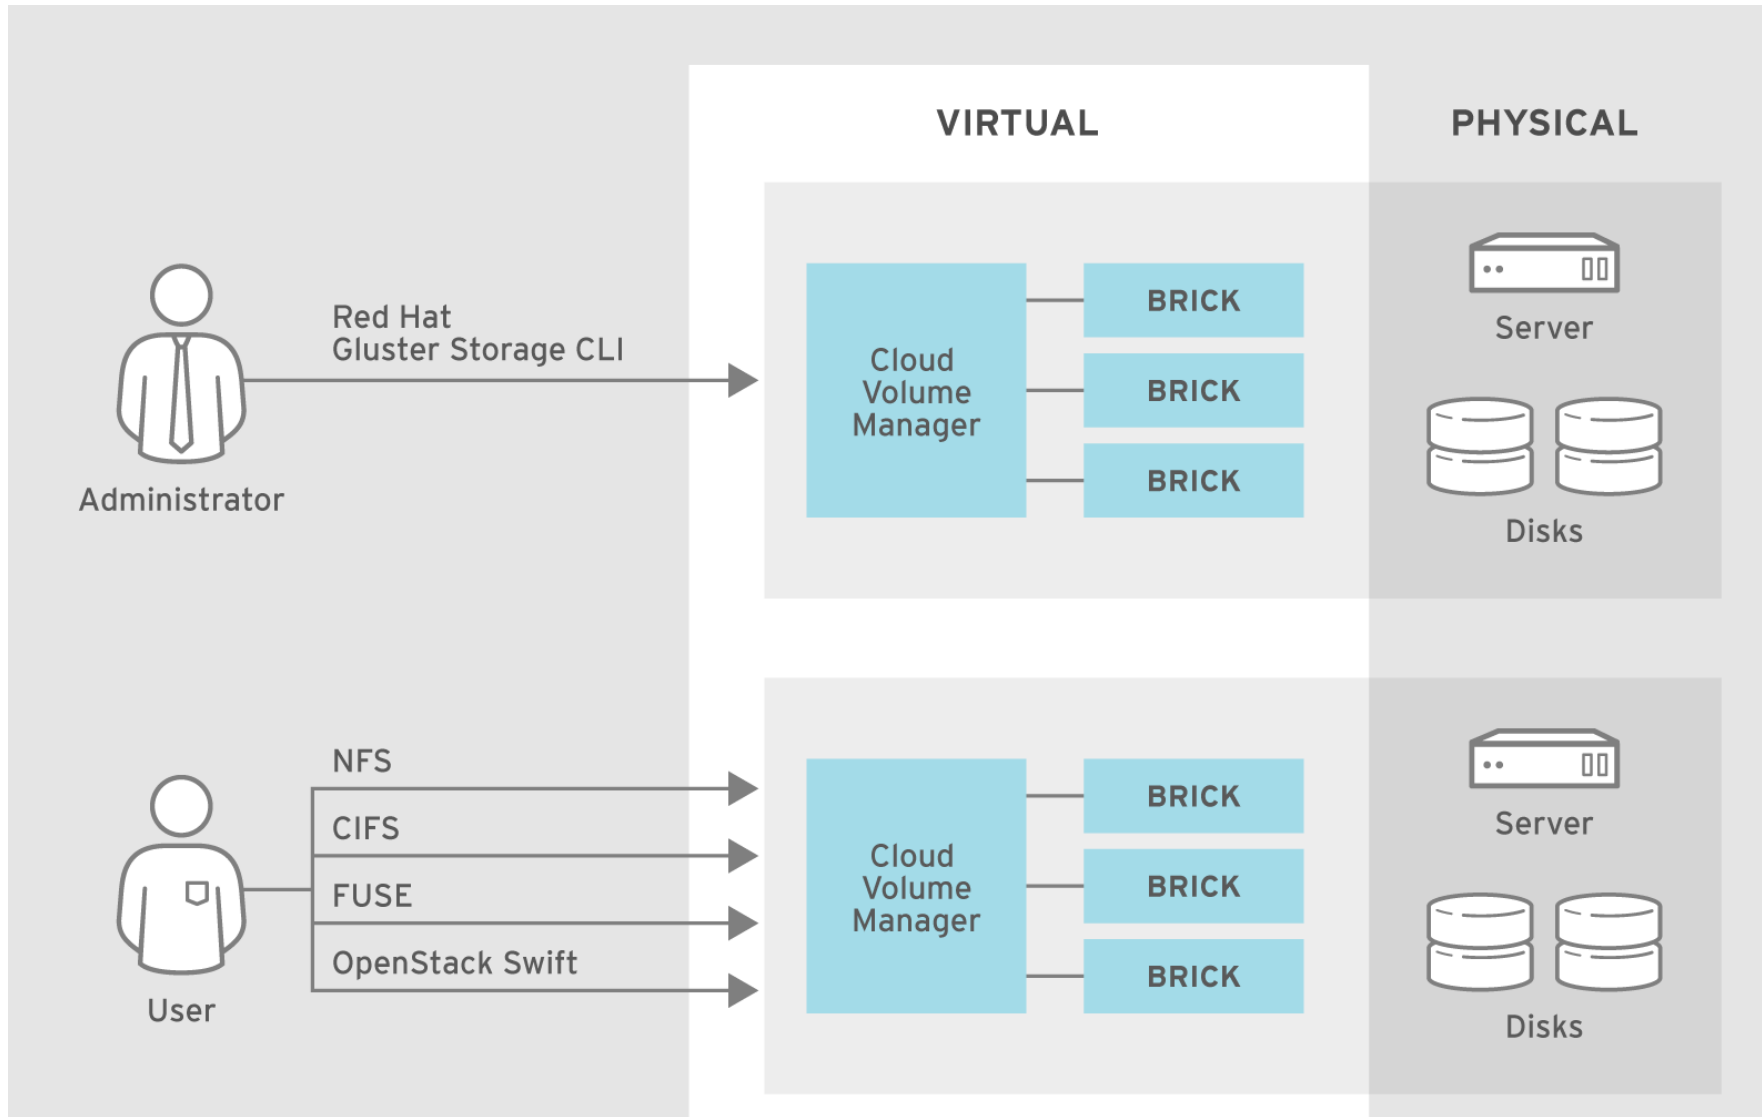


# GlusterFS

- An open source, scale-out distributed file system
- Software Only and operates in user-space
- Aggregates Storage into a single unified namespace
- No metadata server architecture
- Provides a modular, stackable design
- Runs on commodity hardware



# GlusterFS architecture



# GlusterFS Design

- Data is stored on disk using native formats (e.g. ext4, XFS)
- Has following components
  - Servers known as storage bricks (glusterfsd daemon), export local filesystem as volume
  - Clients (glusterfs process), creates composite virtual volumes from multiple remote servers using stackable translators
  - Management service (glusterd daemon) manages volumes and cluster membership
  - Gluster cli tool



# Integration = GlusterFS + NFS-Ganesha + pNFS

- Introduced in glusterfs 3.7, NFS-Ganesha 2.3
- Supports File Layout
- Entire file will present in a single node
- gfid passed with layout for the communications
- All symmetric architecture – ganessa process can act both as MDS and DS





## (conti..) Integration

- Commit through DS
- Only single MDS per volume
- Ganesha talks to glusterfs server using libgfapi
- Upcall used to sync between MDS and DS



# Libgfapi

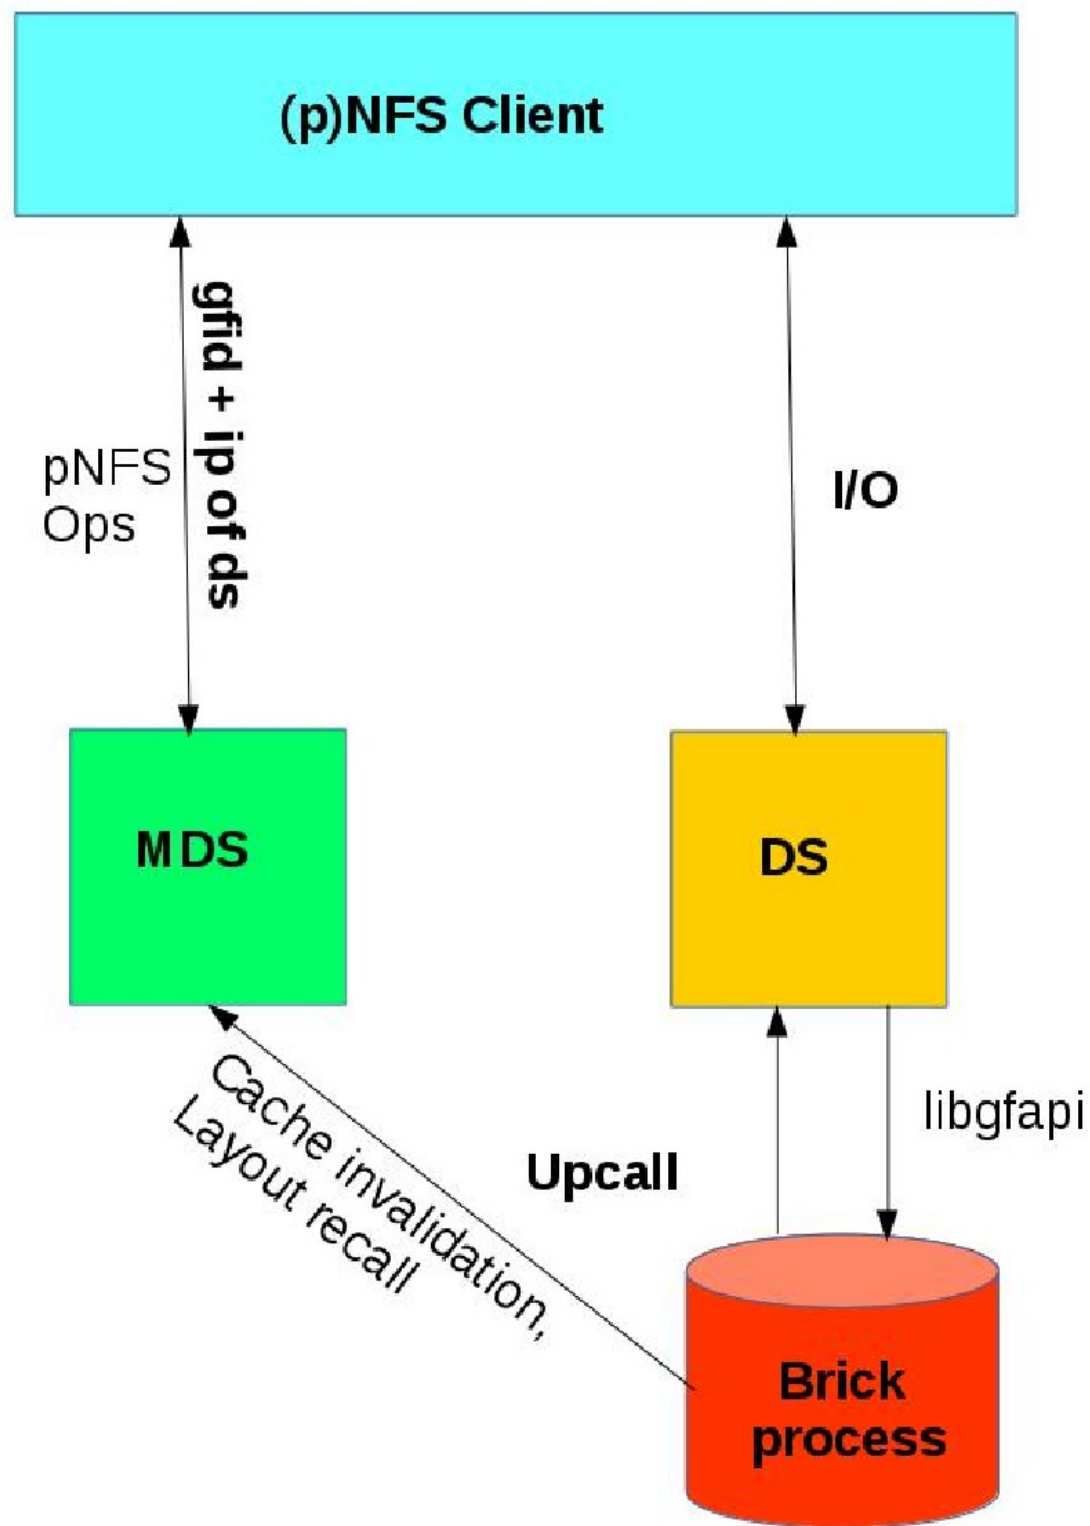
- A user-space library with APIs for accessing Gluster volumes.
- Reduces context switches.
- Many applications integrated with libgfapi (qemu, samba, NFS Ganesha).
- Both sync and async interfaces available.
- C and python bindings.
- Available via 'glusterfs-api\*' packages.



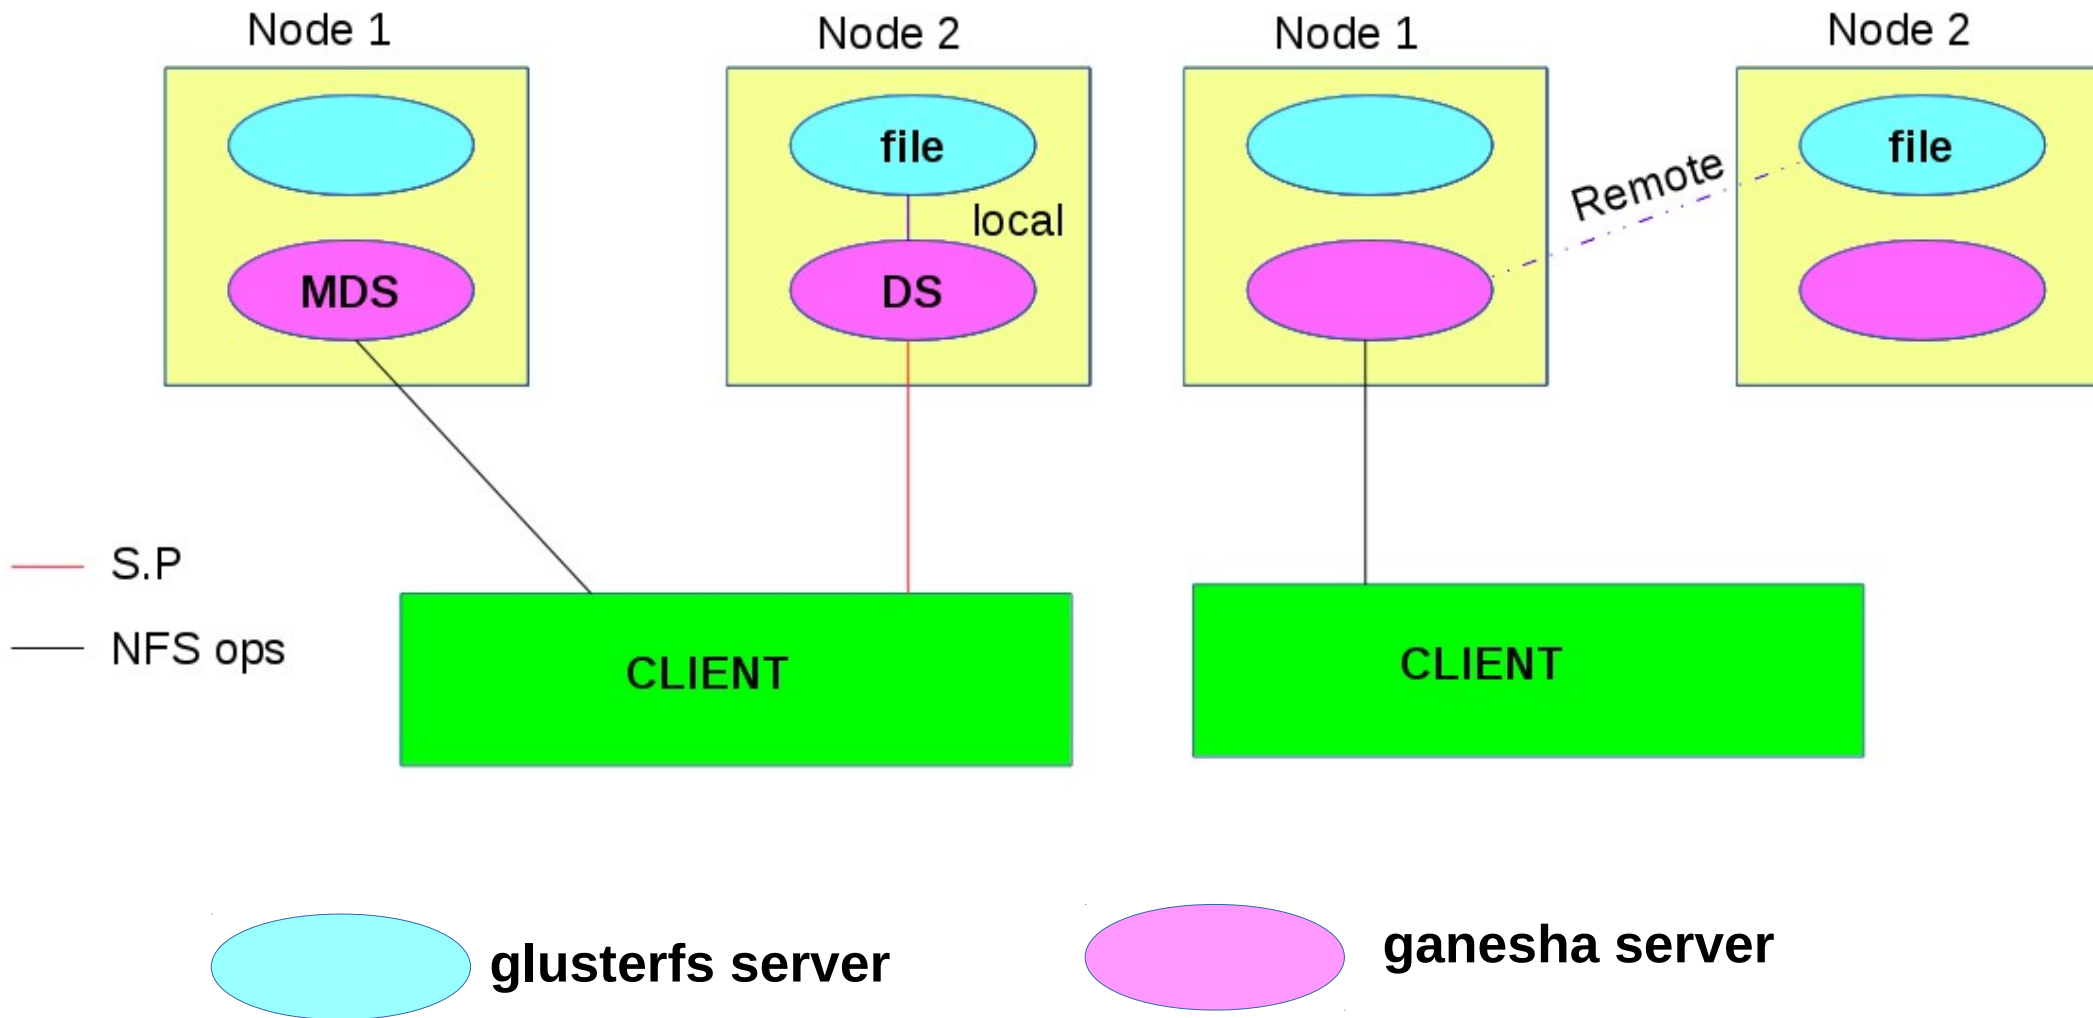
# Upcall Infrastructure

- A generic and extensible framework.
  - used to maintain states in the glusterfsd process for each of the files accessed
  - sends notifications to the respective glusterfs clients in case of any change in that state.
- Cache-Invalidation
  - Invalidate cache used by glusterfs client process
  - #gluster vol set <volname> \ features.cache-invalidation on/off





## pNFS v/s NFSv4



## Advantages

- Better bandwidth utilization
- Avoids additional network hops
- Requires no additional node to serve as MDS
- On different volume configurations

	Performance Improvement		Load balancing
	Read	Write	
<b>Distribute</b>	Yes	Yes	No
<b>Replicate</b>	Yes	No	Yes
<b>Disperse</b>	No	No	Yes




# Challenges

- Layout information
  - gfid + location + offset + iomode
- Perform I/O without open on DS
  - Similar to anonymous fd writes/reads
- Maintains cache coherency b/w MDS and DS
  - Using cache invalidation feature of upcall infra



- Load balancing between DS servers
  - If there are multiple DSes are available , MDS need to chose one among which guarantees local read or writes





- Store layout information as leases (in development)
  - Lease infrastructure provided by glusterfs server stores information about a layout. So when a conflict requests comes it can recall layout.



## Special Mention : Rick Macklem

- pNFS for Gluster NFS server
- GlusterFS has its own inbuilt nfsv3 server either known as gNFS or Gluster NFS will act as DS
- Flex File Layout support
- MDS – FreeBSD NFS server exporting GlusterFS volume



# Configuring pNFS

- Create and start a glusterfs volume
  - `gluster v create <volname> <options> <brick info>`
  - `gluster v start <volname>`
- Turn on cache-invalidation
  - `gluster v set <volname> cache-invalidation on`
- Adding configuration option for MDS in ganesha conf
  - `GLUSTER { PNFS_MDS = true; }`
- Start nfs-ganesha process on all storage nodes
  - `systemctl start nfs-ganesha`
- Mount the volume in the client
  - `mount -t nfs -o vers=4.1 \`  
`<ip of MDS>:/<volname> <mount point>`



## **Future Directions with GlusterFS Features**

- Lease Infrastructure
  - Multiple MDS support
- Storhaug
  - HA cluster for MDS
  - configure option for pNFS
- Adding support for sharded volume
- Integrating with DHT2



# Demo

- Creating cluster of nodes
- Creating a GlusterFS distribute volume (3x1)
- Exporting volumes via Ganesha server
- Mounting the volume via pNFS protocol
- Performing write operation on client.



# References

## ➤ Links (Home Page):

- <https://github.com/nfs-ganesha/nfs-ganesha/wiki>
- <http://www.gluster.org>

## ➤ References:

- <http://gluster.readthedocs.org>
- <http://blog.gluster.org>
- <https://tools.ietf.org/html/rfc5661>
- <http://events.linuxfoundation.org/sites/events/files/slides/pnfs.pdf>
- <https://people.freebsd.org/~rmacklem/pnfs-setup.txt>



# Contact

➤ Mailing lists:

- [nfs-ganesha-devel@lists.sourceforge.net](mailto:nfs-ganesha-devel@lists.sourceforge.net)
- [gluster-users@gluster.org](mailto:gluster-users@gluster.org)
- [gluster-devel@gluster.org](mailto:gluster-devel@gluster.org)

➤ IRC:

- #ganesha on freenode
- #gluster and #gluster-dev on freenode



# Q & A





# Thank You

