

# **An Introduction to the Implementation of ZFS**

Brought to you by

Dr. Marshall Kirk McKusick

European BSD Conference 2014  
27th September 2014

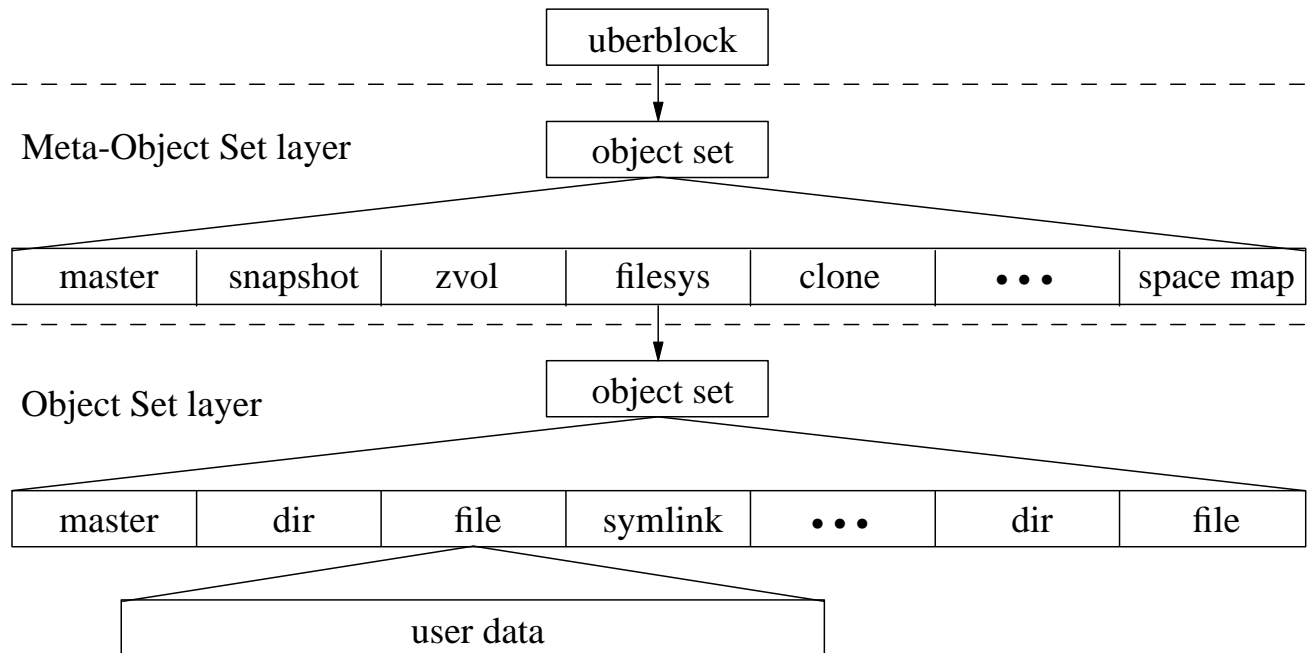
InterExpo Congress Center  
Sofia, Bulgaria

Copyright 2014 Marshall Kirk McKusick.  
All Rights Reserved.

# **Zettabyte Filesystem Overview**

- Never over-write an existing block
- Filesystem is always consistent
- State atomically advances at checkpoints
- Snapshots (read-only) and clones (read-write) are cheap and plentiful
- Metadata redundancy and data checksums
- Selective data compression and deduplication
- Pooled storage shared among filesystems
- Mirroring and single, double, and triple parity RAIDZ
- Space management with quotas and reservations
- Fast remote replication and backups

# Structural Organization



- Uberblock anchors the pool
- Meta-object-set (MOS) describes array of filesystems, clones, snapshots, and ZVOLs
- Each MOS object references an object-set that describes its objects
- Filesystem object sets describe an array of files, directories, etc.
- Each filesystem object describes an array of bytes

# ZFS Block Pointer

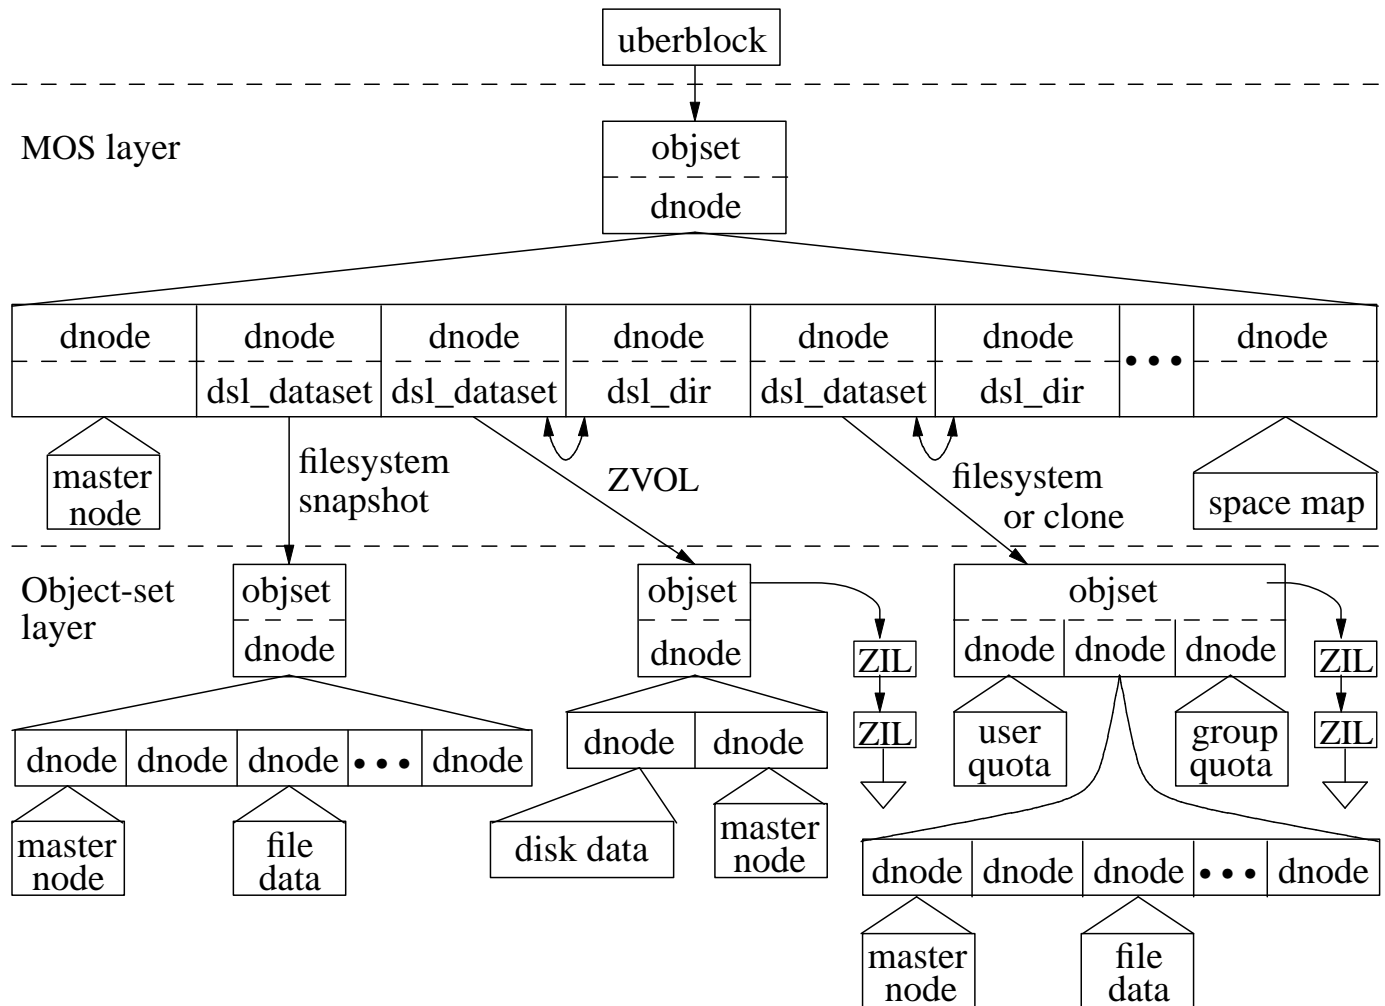
|   |                     |         |    |     |      |       |       |       |       |
|---|---------------------|---------|----|-----|------|-------|-------|-------|-------|
|   | 64                  | 56      | 48 | 40  | 32   | 24    | 16    | 8     | 0     |
| 0 | vdev1               |         |    |     |      | grid  | asize |       |       |
| 1 | G                   | offset1 |    |     |      |       |       |       |       |
| 2 | vdev2               |         |    |     |      | grid  | asize |       |       |
| 3 | G                   | offset2 |    |     |      |       |       |       |       |
| 4 | vdev3               |         |    |     |      | grid  | asize |       |       |
| 5 | G                   | offset3 |    |     |      |       |       |       |       |
| 6 | B                   | D       | X  | lvl | type | cksum | comp  | psize | lsize |
| 7 | spare               |         |    |     |      |       |       |       |       |
| 8 | spare               |         |    |     |      |       |       |       |       |
| 9 | physical birth time |         |    |     |      |       |       |       |       |
| A | logical birth time  |         |    |     |      |       |       |       |       |
| B | fill count          |         |    |     |      |       |       |       |       |
| C | checksum[0]         |         |    |     |      |       |       |       |       |
| D | checksum[1]         |         |    |     |      |       |       |       |       |
| E | checksum[2]         |         |    |     |      |       |       |       |       |
| F | checksum[3]         |         |    |     |      |       |       |       |       |

- Up to three levels of redundancy
- Checksum separate from data
- Birth time is the transaction-group number in which it was allocated
- Maintains allocated, physical (compressed), and logical sizes

# **ZFS Block Management**

- Disk blocks are kept in a pool
- Multiple filesystems and their snapshots are held in the pool
- Blocks from the pool are given to filesystems on demand and reclaimed to the pool when freed
- Space may be reserved to ensure future availability
- Quotas may be imposed to limit the space that may be used

# ZFS Structure




- MOS layer manages space and objects using that space
- Object-set layer manages filesystems, snapshots, clones, and ZVOLs

## ZFS Checkpoint

- Collect all updates in memory
- Periodically write all changes to an unused location to create a checkpoint
- Last step in checkpoint writes a new uberblock
- Entire pool is always consistent
- Checkpoint affects all filesystems, clones, snapshots, and ZVOL in the pool
- Need to log any changes between checkpoints that need to be persistent
- The **fsync** system call is implemented by forcing a log write not by doing a checkpoint
- Recovery starts from last checkpoint, rolls forward through log, then creates new checkpoint

# Flushing Dirty Data



Write modified data in this order:

- 1) new or modified user data
- 2) indirect blocks to new user data
- 3) new dnode block
- 4) indirect blocks to modified dnodes
- 5) object-set dnode for filesystem dnodes
- 6) filesystem dnode to reference objset dnode
- 7) indirect blocks to modified meta-objects
- 8) MOS object-set for meta-object dnode
- 9) new uberblock (plus its copies)



## **ZFS Strengths**

- High write throughput
- Fast RAIDZ reconstruction on pools with less than 40% utilization
- Avoids RAID "write hole"
- Blocks move between filesystems as needed
- Integration eases administration (mount points, exports, etc)

## ZFS Weaknesses

- Slowly written files scattered on disk
- Slow RAIDZ reconstruction on pools with greater than 40% utilization
- Block cache must fit in the kernel's address space, thus works well only on 64-bit systems
- Needs under 75% utilization for good write performance
- RAIDZ has high overhead for small block sizes such as 4 Kbyte blocks typically used by databases and ZVOLS.
- Blocks cached in memory are not part of the unified memory cache so inefficient for files and executables using **mmap** or when using **sendfile**

# Questions

More on ZFS:

- “The Design and Implementation of the FreeBSD Operating System, 2nd Edition”, Chapter 10
- Manual pages: `zfs(8)`, `zpool(8)`, `zdb(8)`

Marshall Kirk McKusick

`<mckusick@mckusick.com>`

<http://www.mckusick.com>