

OpenZFS

Why libshare? WHY!

...

George Wilson

gwilson@delphix.com

@zfsdude

- Utilize *sharenfs* from our applications to export NFS filesystems
- Properties are maintained with ZFS filesystem
- Allow for migration between platforms
- Utilize *libshare* functionality
 - built-in locking and caching
 - tightly coupled with NFS

- **Concurrency Issues**

- Multiple threads attempting to share/unshare a dataset
- Added file locking as a temporary solution

- **Scalability Issues**

- Constanting reading /proc/self/mounts
- Enabled caching of file to avoid excessive number of reads

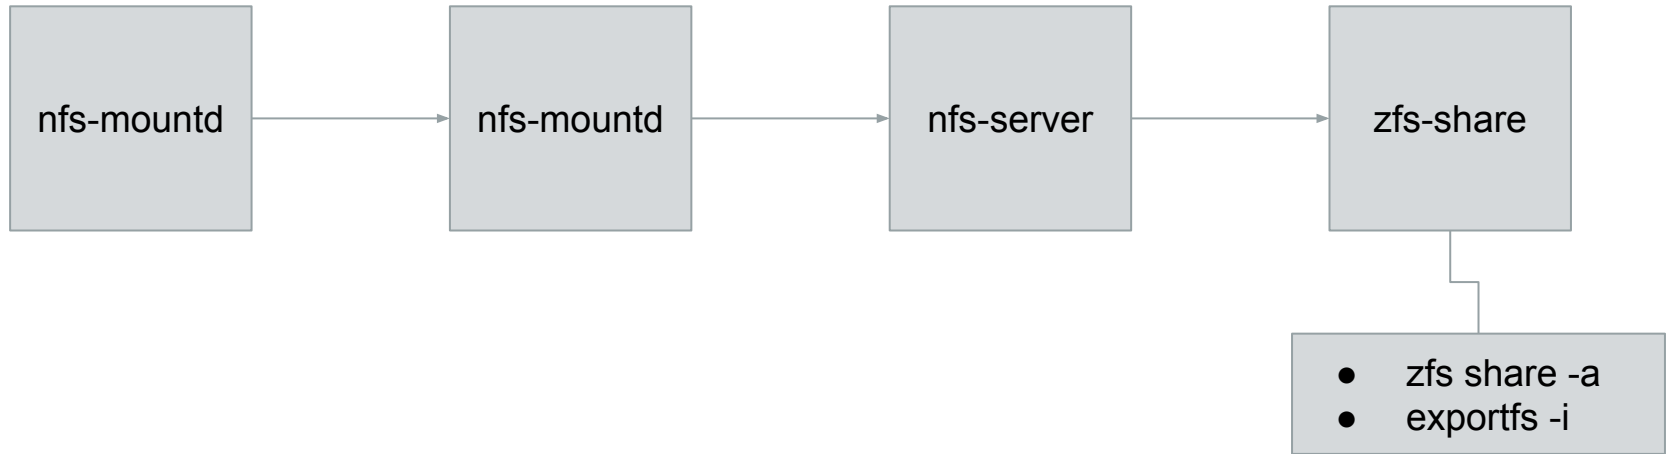
- **Stale File Handles**

- Rebooting/Restarting nfs server results in mount failures

NFS/ZFS Systemd Dependencies




Open**ZFS**



Problem with dependencies



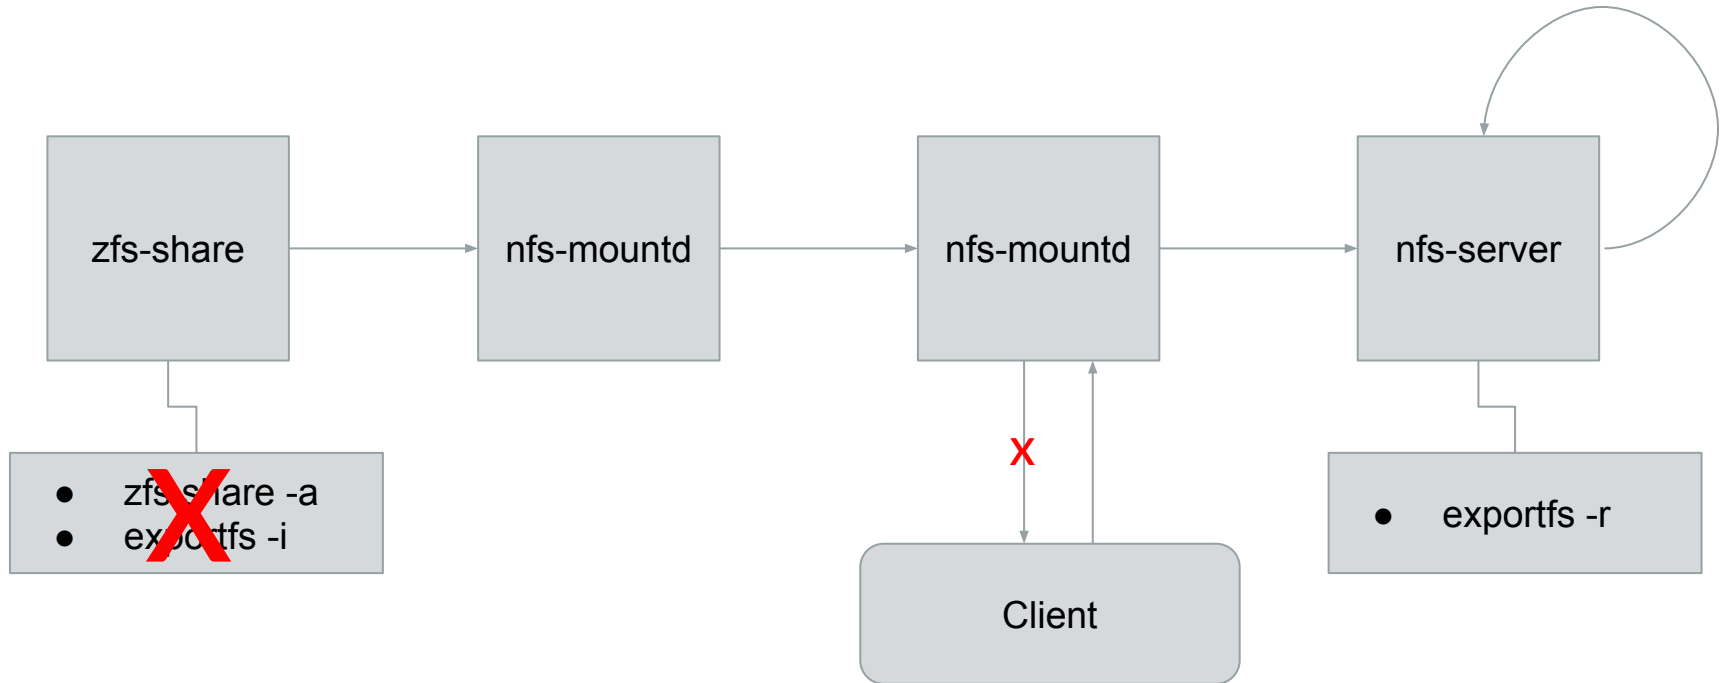
Open**ZFS**



Reorder Dependencies



Open**ZFS**



REQUIREMENTS

- ZFS sharenfs logic needs to happen before mountd/nfsd
- Ideally should have close ties to nfs service

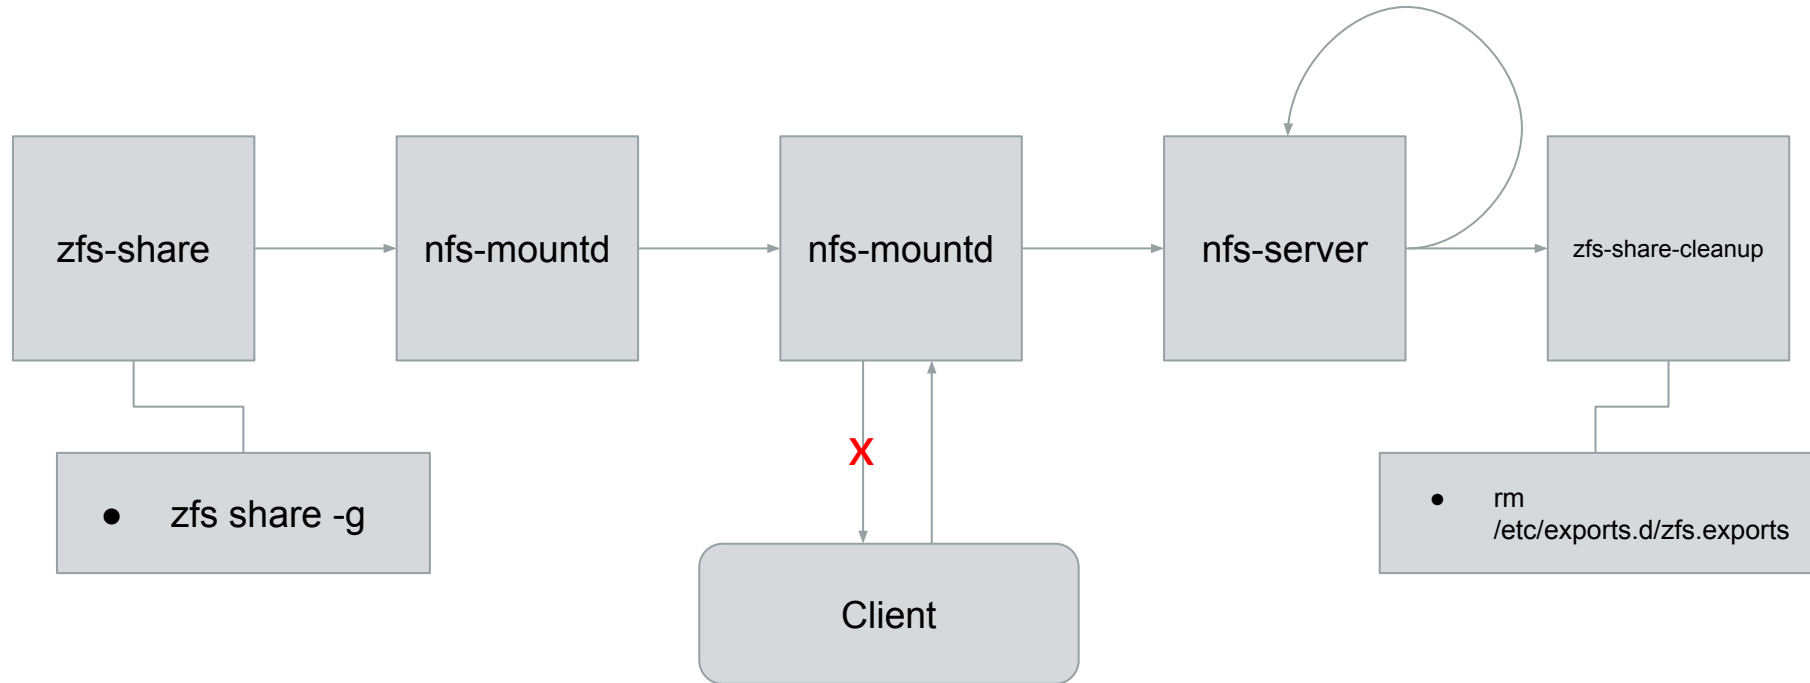
SOLUTION

- Utilize `/etc/exports`
 - Automatically consumed by nfs-service
- Need a way to generate all exports
 - Replace “zfs share -a” with “zfs share -g”
 - Generates nfs-service consumable output
 - Much faster than “zfs share -a”
- Introduce cleanup service

New Dependencies



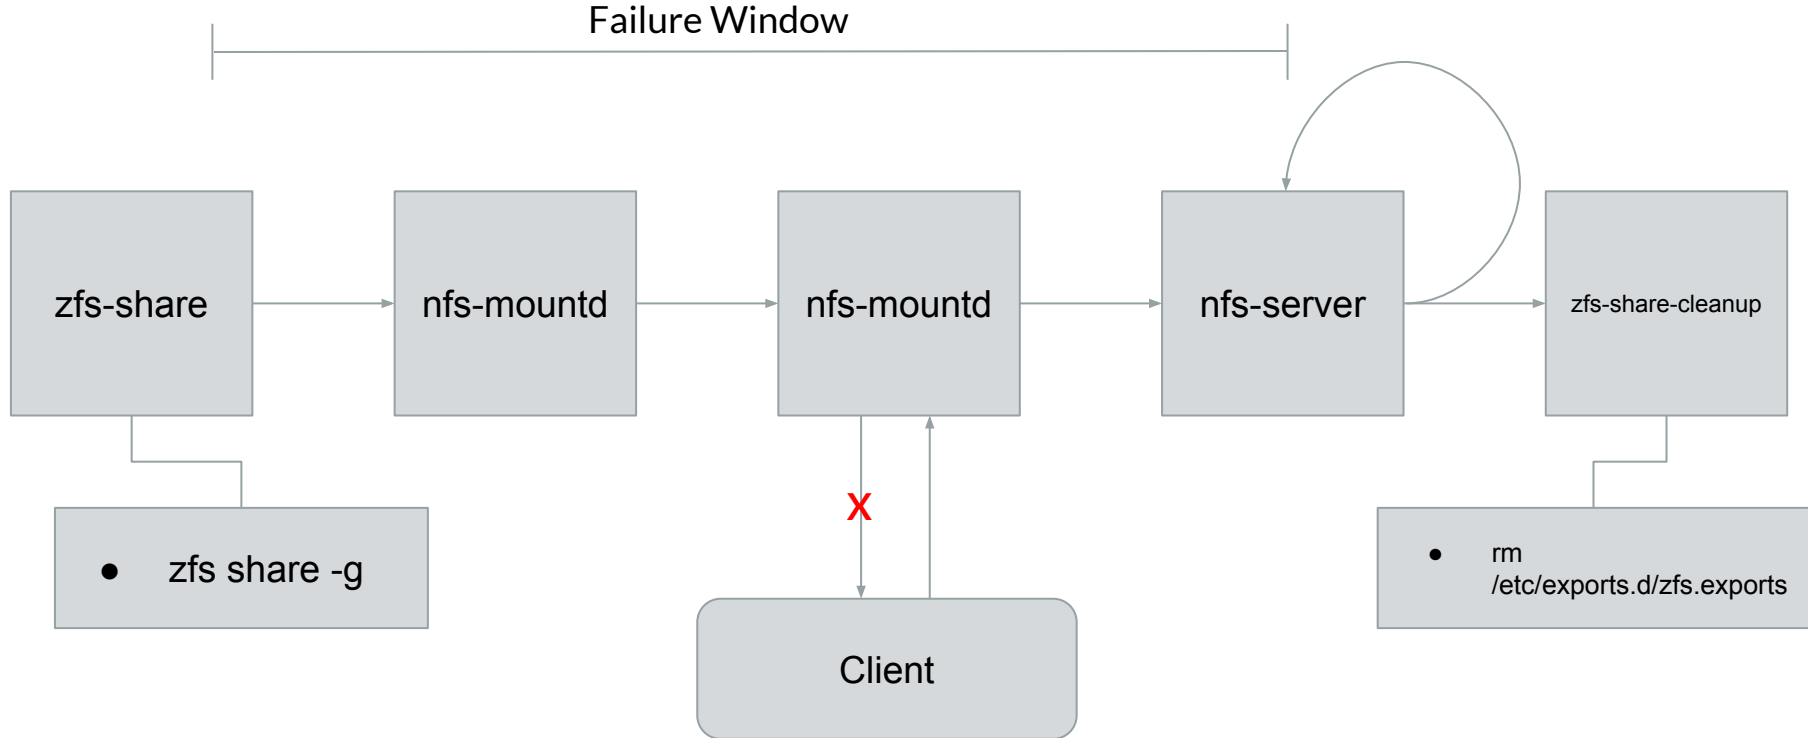
Open**ZFS**



Failure Window

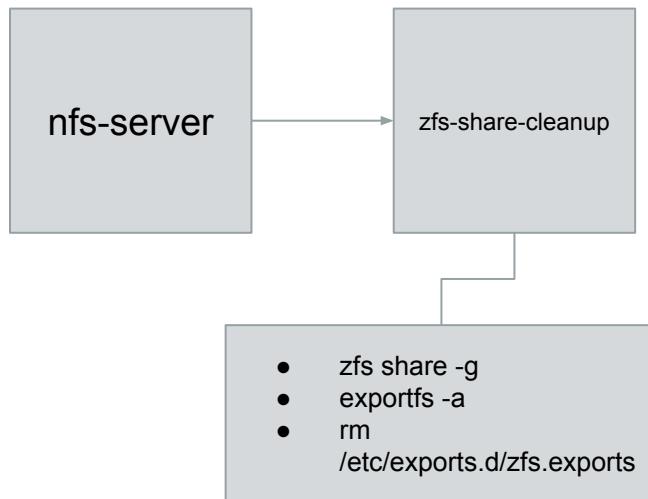


Open**ZFS**



Closing the window (mostly)

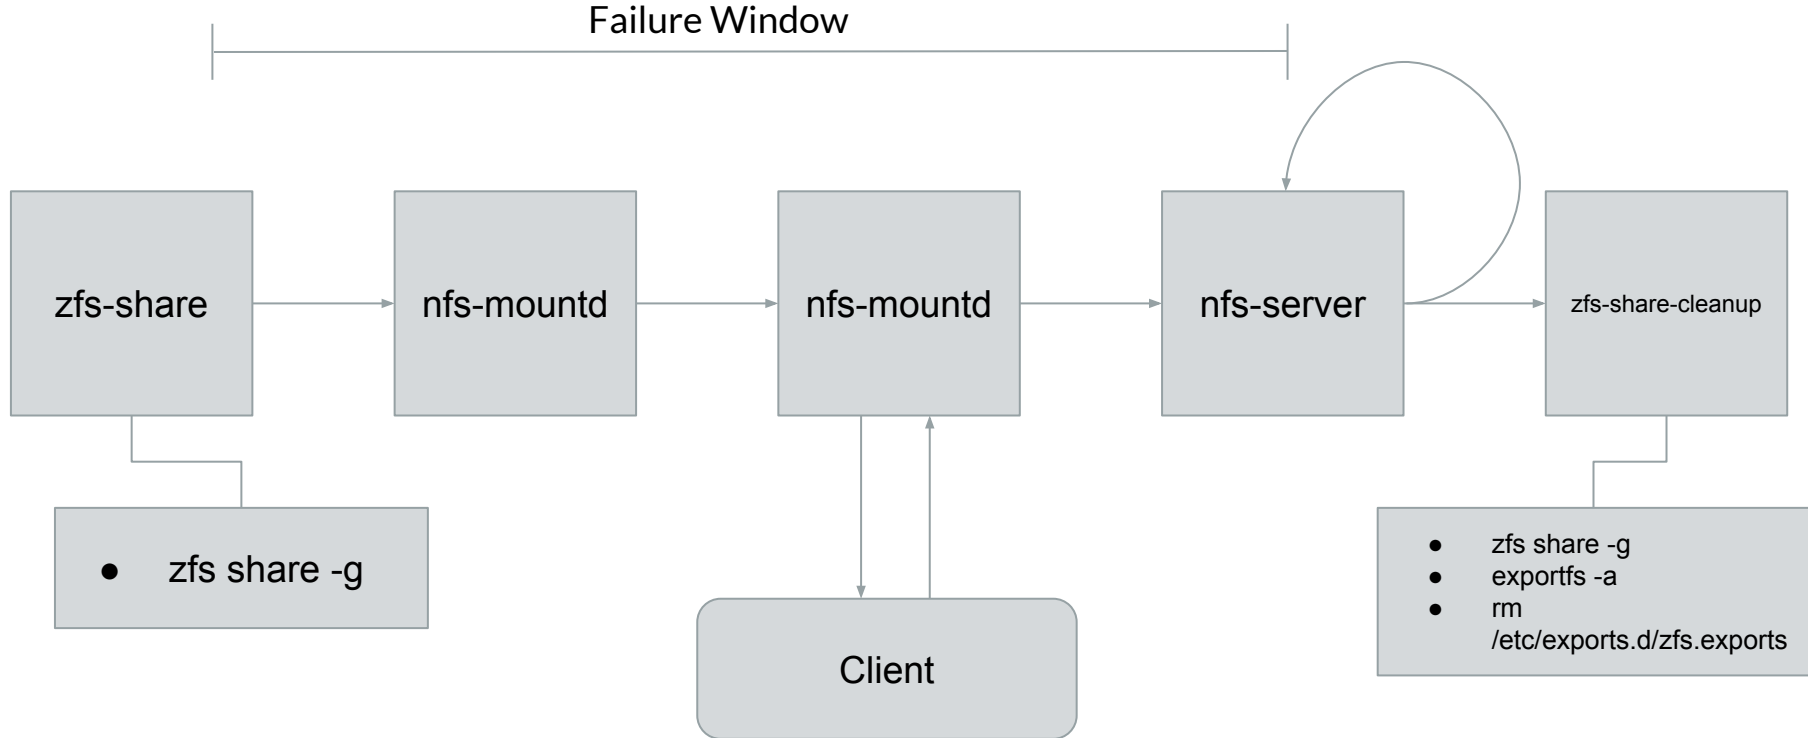
- Extend cleanup service
- Regenerate any exports file
- Re-export all filesystems



Temporary Solution



Open**ZFS**



Performance Impact

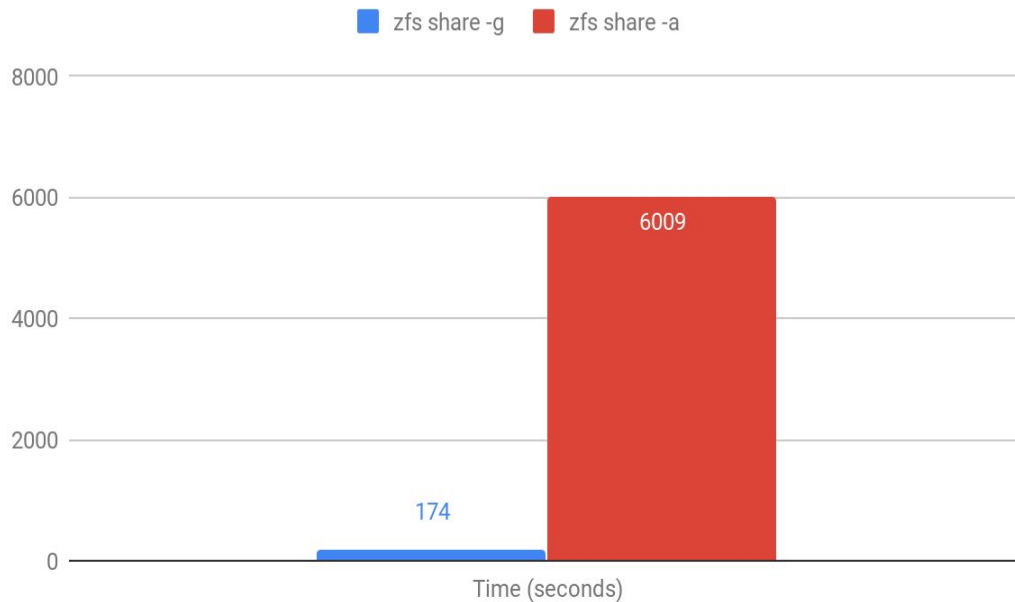


OpenZFS

- VM with 64GB of RAM
- 60K+ filesystem
- 8 vCPUs

**Over 3000%
Improvement**

zfs share -g and zfs share -a





- Redefine libshare
 - Define simplified API for loosely coupled platforms
- Rework Linux libshare
 - Update `/etc/exports` or `/etc/export.d` directly
 - Interact with `systemd` as required
- First phase, focus only on NFS sharing
 - SMB shares need more investigation

Questions?