



# OpenZFS


## Persistent L2ARC



George Amanakis  
gamanakis@gmail.com



- L2ARC caches buffers from ARC
- Rotary implementation
  1. Writes sequentially on device
  2. Evicts previously written buffers
  3. Writes new ones
  4. Repeat
- Loops from the beginning if cache device is full



# How can we make the L2ARC persistent?



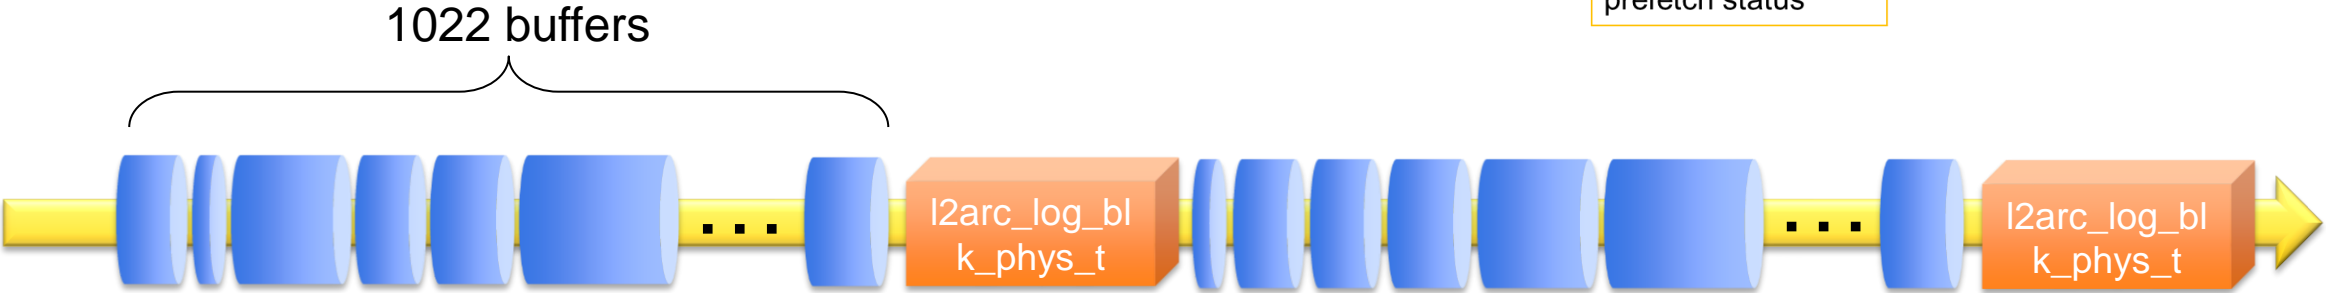
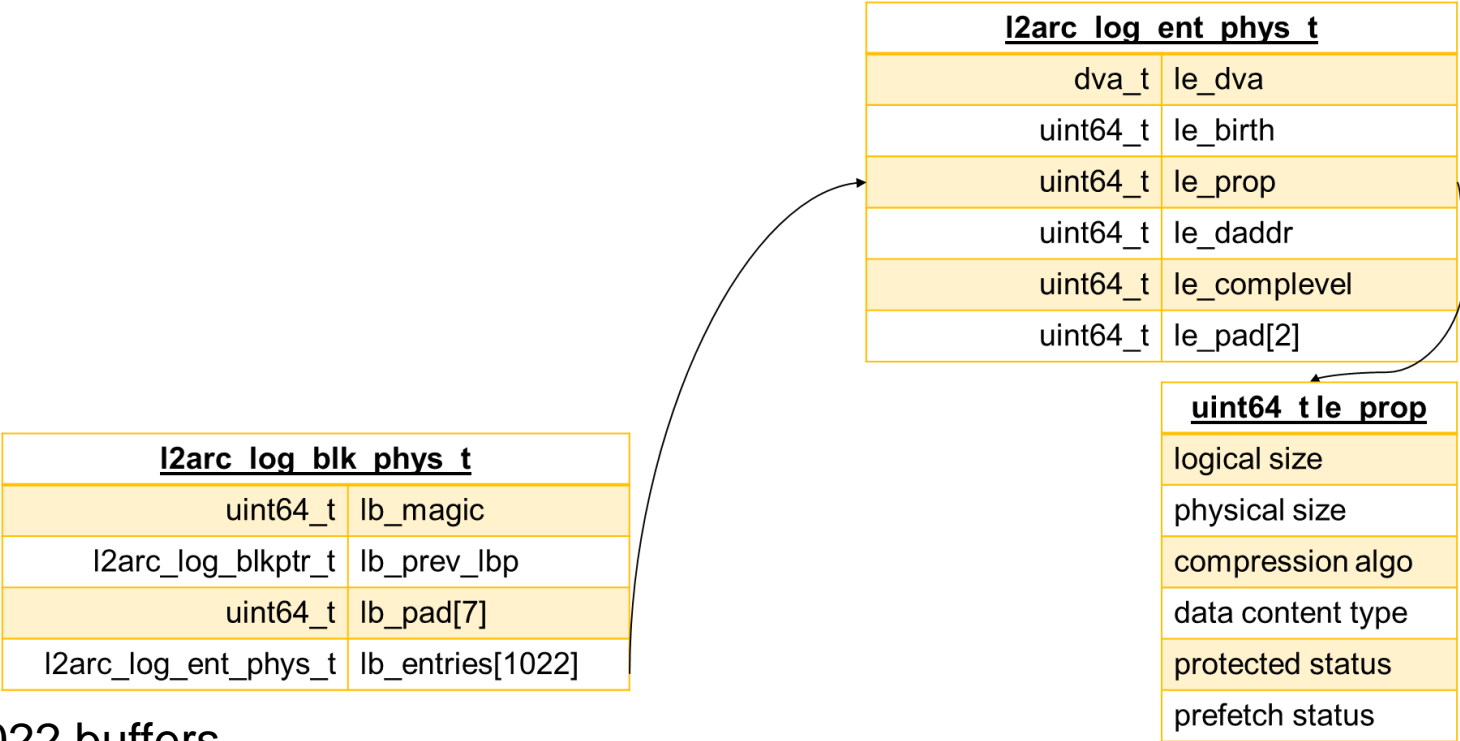
- Enabling persistence means we need to restore the header entries of L2ARC buffers in ARC
- **L2ARC log blocks:** L2ARC metadata containing buffer header entries
  - Written on disk every 1022 L2ARC buffers

<u>l2arc_log_blk_phys_t</u>	
uint64_t	lb_magic
l2arc_log_blkptr_t	lb_prev_lbp
uint64_t	lb_pad[7]
l2arc_log_ent_phys_t	lb_entries[1022]

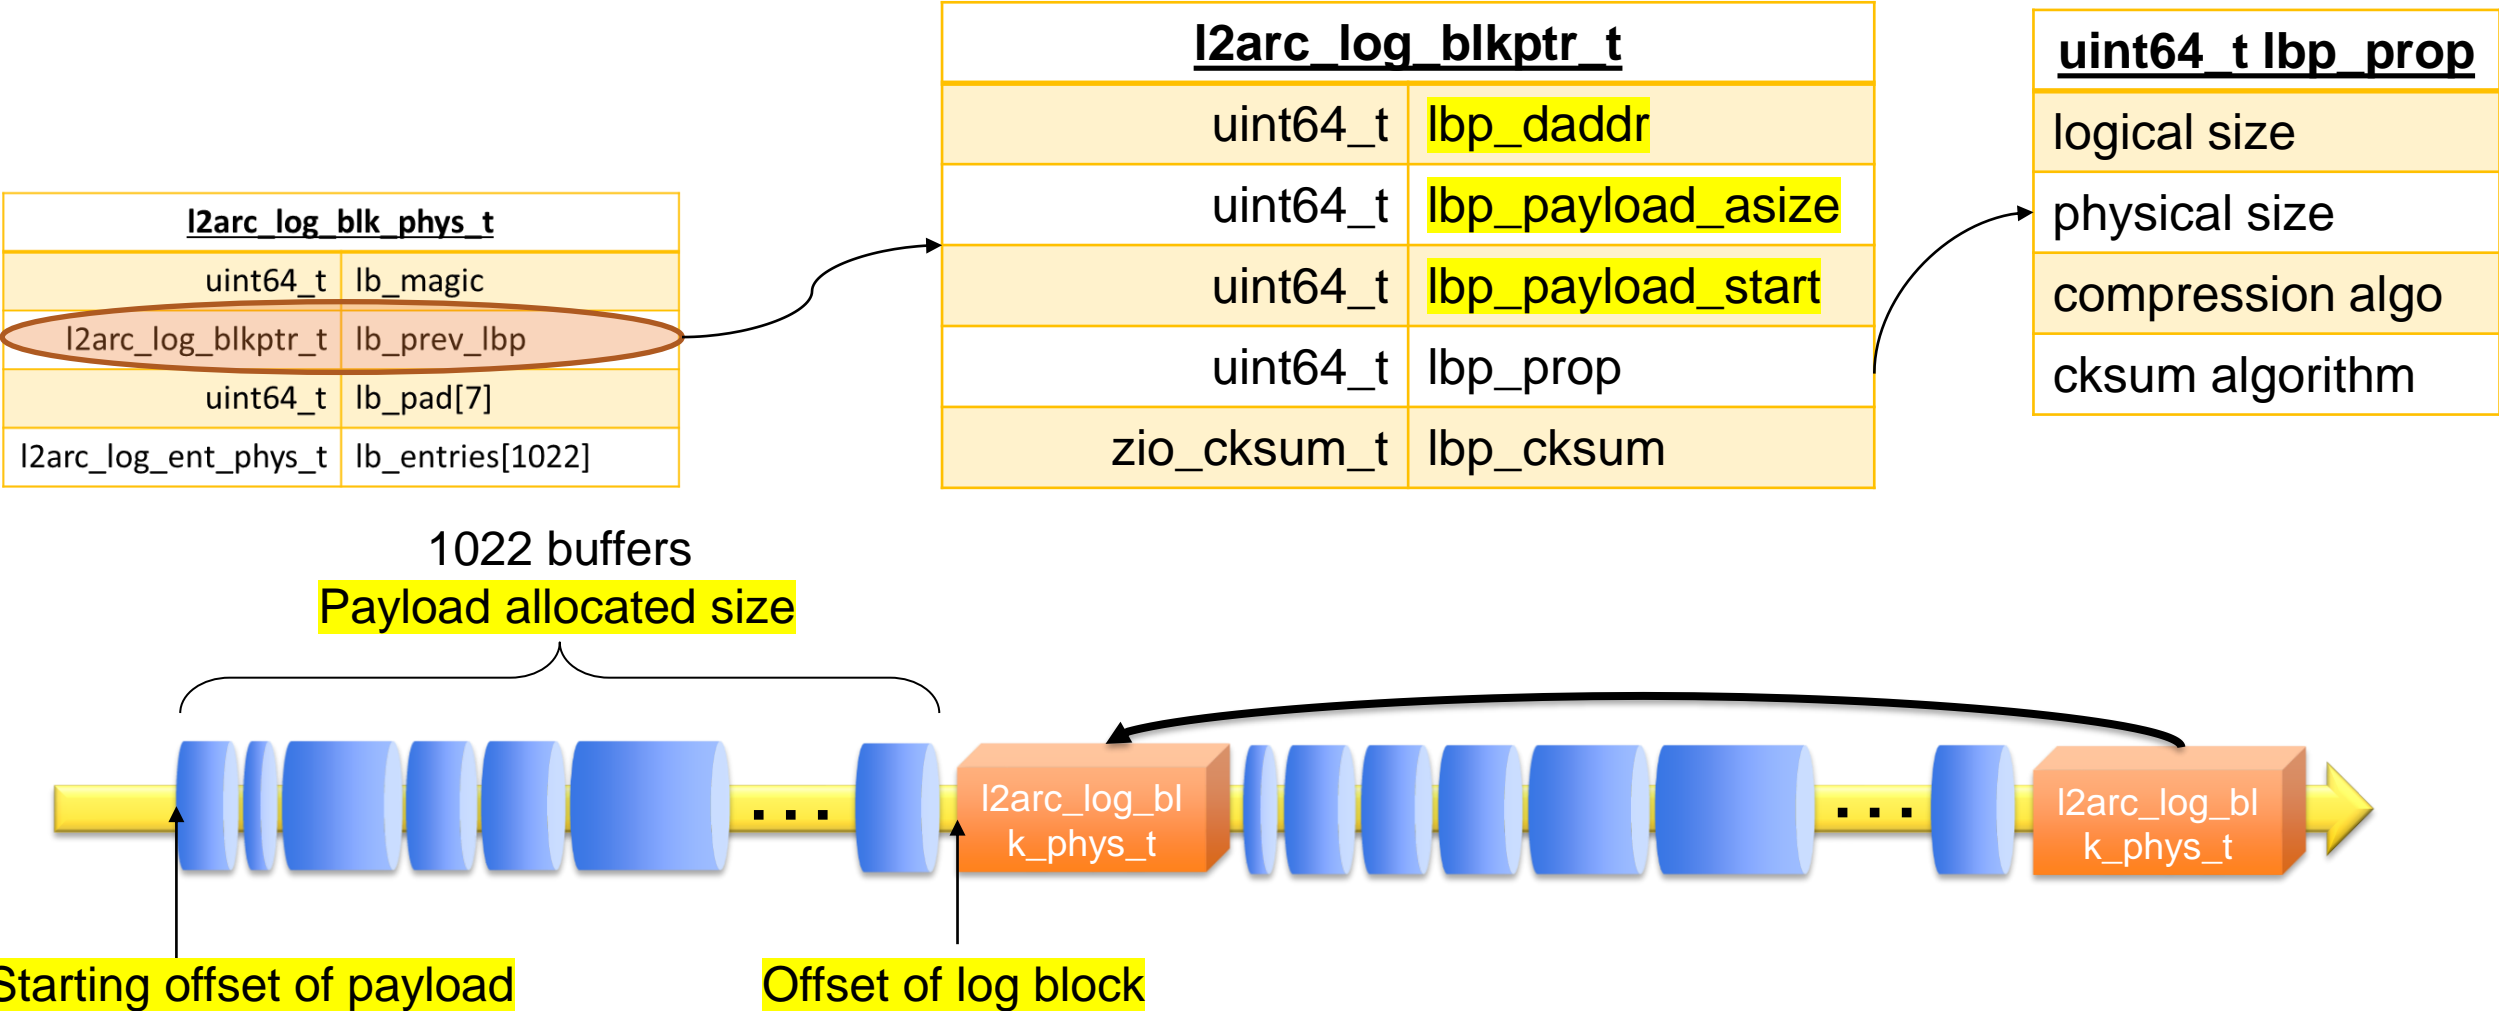
<u>l2arc_log_ent_phys_t</u>	
dva_t	le_dva
uint64_t	le_birth
uint64_t	le_prop
uint64_t	le_daddr
uint64_t	le_complevel
uint64_t	le_pad[2]

<u>uint64_t le_prop</u>
logical size
physical size
compression algo
data content type
protected status
prefetch status


# Storing buffer header entries: L2ARC log blocks



# Keeping track of L2ARC log blocks: L2ARC log block pointers



# Performance optimization: two interleaved chains of log blocks



## Rebuild process

- Issue a sync read to read log block (4)
- Issue an async read to read the log block (3)
- Decompress and restore log block (4)
- Check if log block (3) has been read...

## Performance

- consumer grade SATA SSD
- Intel Xeon E5-2667v2
- 64GB L2ARC device, ~100GB logical size
  - 2.8 sec with 1 chain
  - 1.9 sec with 2 interleaved chains (~32% faster)

- L2ARC rebuild is done asynchronously with respect to pool import
- No buffers are written to the cache device until the rebuild has been completed

# How does the L2ARC rebuild start?

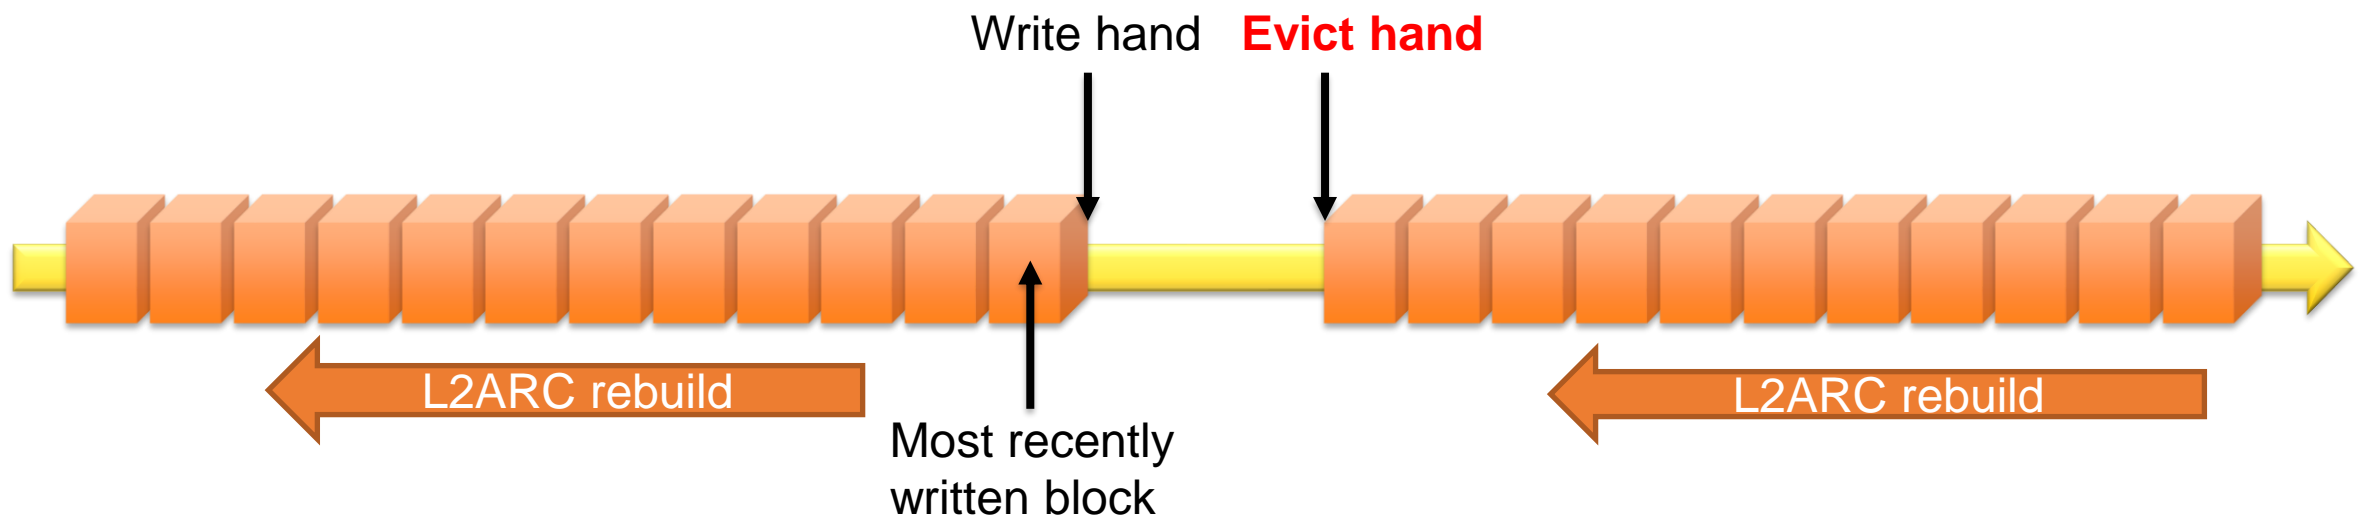
- **L2ARC device header**

- Is updated each time a log block is written to the cache device
- Contains pointers to the two **most recently written** log blocks

<u>l2arc_dev_hdr_phys_t</u>	
uint64_t	dh_magic
uint64_t	dh_version
uint64_t	dh_spa_guid
uint64_t	dh_vdev_guid
uint64_t	dh_log_entries
uint64_t	dh_evict
uint64_t	dh_flags (first pass)
uint64_t	dh_lb_asize
uint64_t	dh_lb_count
l2arc_log_blkptr_t	dh_start_lbps[2]
...	
zio_eck_t	dh_tail

# When does the L2ARC rebuild stop?

- Stop if the rebuild if we reach the **evict hand**
- The evict hand is stored in the device header
- The range between the write and the evict hand may have been zeroed-out if L2ARC TRIM is enabled







- *l2arc\_rebuild\_enabled*
  - Controls whether ZFS will attempt to rebuild the L2ARC
  - Log blocks are still written on the device
  - Defaults to 1 (true)
- *l2arc\_rebuild\_blocks\_min\_l2size*
  - Disables the writing of L2ARC log blocks on the device
  - Beneficial for small devices
  - Defaults to 1GB



## # zdb -ll cachedevice

### ----- L2ARC device header -----

```

magic: 6504978260106102853
version: 1
pool_guid: 2721483723720346099
flags: 1
start_lbps[0]: 310026240
start_lbps[1]: 298739200
log_blk_ent: 87
start: 4194816
end: 1467482112
evict: 4194816
lb_asize_refcount: 50176
lb_count_refcount: 28
trim_action_time: 0
trim_state: 0

```

```

log_blk_count:    28 with valid cksum
                  0 with invalid cksum
log_blk_asize:    50176

```

### ----- L2ARC device log blocks -----

```

lb[  1]  magic: 5498692020116080708
|
|      daddr: 310026240
|
|      payload_asize: 11285504
|
|      payload_start: 298740736
|
|      lsize: 65536
|
|      asize: 1536
|
|      compralgo: 15
|
|      cksumalgo: 7
|

```

← LZ4

← FLETCHER4

```

lb[  1]  le[  1]  DVA asize: 131072, vdev: 0, offset: 348140032
|
|      birth: 18
|
|      lsize: 131072
|
|      psize: 131072
|
|      compr: 2
|
|      complevel: 0
|
|      type: 1
|
|      protected: 0
|
|      prefetch: 0
|
|      address: 298740736
|
|      ARC state: 1

```

← Off

← Buffer content: Data



```
...
l2_size          4      320079872
l2_asize         4      305926656
l2_hdr_size      4      230592
l2_log_blk_writes 4      28
l2_log_blk_avg_asize 4    2215
l2_log_blk_asize  4    50176
l2_log_blk_count  4     28
l2_data_to_meta_ratio 4    3762
l2_rebuild_success 4     1
l2_rebuild_unsupported 4    0
l2_rebuild_io_errors 4    0
l2_rebuild_dh_errors 4    0
l2_rebuild_cksum_lb_errors 4  0
l2_rebuild_lowmem  4    0
l2_rebuild_size    4    318228480
l2_rebuild_asize   4    305782784
l2_rebuild_bufs    4    2436
l2_rebuild_bufs_precached 4  30
l2_rebuild_log_blks 4    28
...
```

## # zpool history -i pool

```
2020-09-17.16:53:18 [txg:7775] L2ARC rebuild
successful, restored 28 blocks
```

```
2020-09-17.16:53:18 lt-zpool import tst3 -d
/home/user/vdevs
```

Questions?



- Status:
  - Done!
  - In master branch and upcoming OpenZFS 2.0!

- Acknowledgments:
  - Saso Kiselkov (Nexenta)
  - Yuxuan Shui
  - Jorgen Lundman

Everybody who reviewed code!

- Matt Ahrens
- Brian Behlendorf
- Ryan Moeller
- Kjeld Schouten-Lebbing
- George Wilson